

@ λ : Conservatively Extending a Type System From Within

Cyrus Omar Jonathan Aldrich

Carnegie Mellon University
{comar, aldrich}@cs.cmu.edu

Abstract

Researchers often extend an existing language with new type and operator constructors. But because this is not generally possible from within the language, researchers must implement extensions as distinct forks of the language. This limits their utility: building applications where different components are written in languages with different type systems is unsafe and unnatural. An internally-extensible type system could help address these issues, but designing an extension mechanism that is both expressive and reliable is a challenge: extensions must not be permitted to compromise key metatheoretic properties of the language, nor should they be allowed to interfere with one another, in any combination. This latter property, *conservativity*, is not guaranteed by prior mechanisms.

We introduce @ λ , a simply-typed lambda calculus with simply-kinded type-level computation where the semantics of new type and operator constructors can be defined in a functional style from within. Its semantics are structured as an extensible variant of the Harper-Stone elaboration semantics used to specify Standard ML. We incorporate typed compilation techniques and a form of type abstraction that together allow us to prove type safety, decidability and conservativity. We intend @ λ to serve as a minimal foundation for future work on extension mechanisms for typed programming languages, but it is already quite expressive. We begin by showing how a more natural concrete syntax can be introduced by a type-directed desugaring to @ λ . Then, we show how a number of abstractions that must be built into other languages can be implemented orthogonally, as libraries, within @ λ . We conclude with a discussion of its present limitations.

1. Introduction

Typed programming languages are typically organized around a collection of indexed type and operator constructors. The simply-typed lambda calculus (STLC), for example, provides a single type constructor, \rightarrow , indexed by a pair of types, and two operator constructors, λ , indexed by a type, and *ap*, which can be thought of as being indexed trivially. We can extend it with universally quantified and recursive types by adding the type constructors \forall and μ , each indexed by a type binder (a type-level term that may refer to a free type variable), and their associated operator constructors. It is tempting to stop here: the language is now universal (i.e. Turing

complete) and the dynamic semantics of operators associated with other useful types, including sums and products, can be defined by a method analogous to Church encodings in the untyped lambda calculus [?]. This is of clear theoretical and pedagogical interest. But this calculus is plainly impractical as a programming language. Reynolds, striking a cautionary tone reminiscent of Perlis in his description of the “Turing tarpit” [?], put it succinctly [?]:

To say that any reasonable function can be expressed by some program is not to say that it can be expressed by the most reasonable program. It is clear that the language requires a novel programming style.

A simple language suffices when only the dynamic semantics of a program are being considered but continuing to add type and operator constructors (collectively, *constructs*) to a language can make reasoning about programs simpler and more precise and support a more natural programming style. Consistent with this view, typed programming languages generally expose more constructs externally than their compilers work with internally. A functional language like ML, for example, might expose n -ary product types (tuples), labeled product types (records), recursive labeled sum types (datatypes) and forms of type abstraction (polymorphism and abstract data types). A compiler for the language might then translate these to an intermediate language (IL) with only simple products and sums for further compilation to machine code. Correctly implementing the dynamic semantics of the language is the primary concern during this phase, so a simple IL that decreases the number of cases that must be considered when implementing and verifying the compiler is a wise choice.

It is again tempting to stop here. Indeed, languages like ML are often referred to as “general-purpose languages” and have been used successfully for a range of computing tasks. The constructs they provide certainly occupy “sweet spots” in the design space, and many abstractions can be encoded using these constructs in a completely satisfactory manner. However, they admit no new “universality” theorems stronger than those for the simple calculus described above, and situations continue to arise in both research and practice where new constructs are desirable:

1. General-purpose constructs continue to evolve. There are many variations on record types, for example: structurally-typed records, or records with forms of dynamic dispatch. We will show an example of records with prototypic dispatch in Sec. ??, and argue that they cannot be safely and naturally defined in terms of standard functional constructs. Sum types also admit variants, as we will discuss in the next section.
2. Specialized type systems that enforce stronger invariants than general-purpose abstraction mechanisms can enforce are often proposed by researchers. One need not take more than a cursory glance through the literature to discover type systems for parallel programming [? ?], concurrent programming [? ? ?], distributed programming [?], dataflow programming [?], authen-

right
termi-
nol-
ogy?

ticated data structures [?], information flow [? ?], databases [?], aliasing [?], effects [?], network protocols [?], regular expressions [?], units of measure [?] and many others.

- Interoperability layers that allow programmers to safely and naturally interact with libraries written in a different language require enforcing the type system of that language within the calling language. Using library-based interoperability layers today can lead to safety issues when this is not possible.

Today, complete solutions to these problems generally take the form of a *fork* of an existing general-purpose language (discussed further in the next section). This is unsatisfying because then these solutions are *non-orthogonal*: a programmer can choose either a language supporting a reasonable approach to parallel programming or one that builds in support for working with and statically reasoning about regular expressions, but there is not necessarily a language that supports both. Even if a separation of such concerns can be achieved, such that each component of an application can be written using a specialized language (as advocated by proponents of the *language-oriented approach* to software development [?]), there will be an *interoperability* problem at the interface between components, as described in item 3 above.

everything below is todo
Rather than specifying a complete language, a researcher might simply specify the new type and operator constructors and

An extensible programming language could address these problems by providing a language-integrated mechanism for introducing new type and operator constructors and implementing their associated static and dynamic semantics directly. But, as mentioned in Section ??, some significant challenges must be addressed before such a mechanism can be relied upon. The desire for expressiveness must be balanced against concerns about maintaining various safety properties in the presence of arbitrary combinations of user-defined extensions to the language's core semantics. The mechanism must ensure that desirable *metatheoretic properties* (e.g. type safety, decidability) of the language are maintained by extensions. Because multiple independently developed extensions might be used within one program, the mechanism must further guarantee their *non-interference*. These are the issues we seek to address in this work.

2. From Extensible Compilers to Extensible Languages

To understand the genesis of our internal extension mechanism, it is helpful to begin by considering why most implementations of programming languages cannot even be externally extended. Let us consider, as a simple example, an implementation of Gödel's T, a typed lambda calculus with recursion on primitive natural numbers (see Appendix). A compiler for this language written using a functional language will invariably represent the primitive type families and operators using closed inductive datatypes. For example, a simple implementation in Standard ML may be based around these datatypes:

```
datatype Type = Nat | Arrow of Type * Type
datatype Exp = Var of var
              | Lam of var * Type * Exp | Ap of Exp * Exp
              | Z | S of Exp | Natrec of Exp * Exp * Exp
```

The logic governing typechecking and translation to a suitable intermediate language (for subsequent optimization and compilation by some back-end) will proceed by exhaustive case analysis over the constructors of Exp.

In an object-oriented implementation of Gödel's T, we might instead encode types and operators as subclasses of abstract classes Type and Exp. Typechecking and translation will proceed by the ubiquitous *visitor pattern* by dispatching against a fixed collection of known subclasses of Exp.

programs	ρ	$::=$	tycon TYCON of $\kappa_{\text{idx}} \{ \text{schema } \tau_{\text{rep}}; \theta \}; \rho$ def $t : \kappa = \tau; \rho \mid e$
primitive ops	θ	$::=$	opcon <i>op</i> of $\kappa_{\text{idx}} (\tau) \mid \theta; \theta$
external terms	e	$::=$	$x \mid \lambda x : \tau. e \mid \text{FAM.op}(\tau_1)(e_1; \dots; e_n)$
type-level terms	τ	$::=$	$t \mid \lambda t : \kappa. \tau \mid \tau_1 \tau_2 \mid \bar{z} \mid \tau_1 \oplus \tau_2 \mid \text{"str"}$ $\llbracket \kappa \mid \tau_1 :: \tau_2 \mid \text{fold}(\tau_1; \tau_2; \mathbf{h}, \mathbf{t}, \mathbf{r}. \tau_3) \mid$ $() \mid (\tau_1, \tau_2) \mid \text{fst}(\tau) \mid \text{snd}(\tau)$ equality types if $\tau_1 \equiv \tau_2$ then τ_3 else τ_4 TYCON(τ) denotations case τ of TYCON(\mathbf{x}) $\Rightarrow \tau_1$ ow τ_2 $\llbracket \tau_{\text{term}} \text{ as } \tau_{\text{type}} \rrbracket \mid \text{error}$ case τ of $\llbracket \mathbf{x} \text{ as } \mathbf{t} \rrbracket \Rightarrow \tau_1$ ow τ_2 reified IL $\nabla(\gamma) \mid \blacktriangledown(\sigma)$
kinds	κ	$::=$	$\kappa_1 \rightarrow \kappa_2 \mid \text{list}[\kappa] \mid \mathbb{Z} \mid \text{Str} \mid 1 \mid \kappa_1 \times \kappa_2$ $\star \mid \text{Elab} \mid \text{ITy} \mid \text{ITm}$
internal terms	γ	$::=$	$x \mid \lambda x : \sigma. \gamma \mid \gamma_1 \gamma_2 \mid \text{fix } f : \sigma \text{ is } \gamma$ $(\gamma_1, \gamma_2) \mid \text{fst}(\gamma) \mid \text{snd}(\gamma)$ $\bar{z} \mid \gamma_1 \oplus \gamma_2 \mid \text{if } \gamma_1 \equiv \gamma_2 \text{ then } \gamma_3 \text{ else } \gamma_4$ abs($\llbracket \tau_1 \text{ as } \tau_2 \rrbracket$) $\mid \Delta(\tau)$
internal types	σ	$::=$	$\sigma_1 \rightarrow \sigma_2 \mid \mathbb{Z} \mid \sigma_1 \times \sigma_2 \mid \text{rep}(\tau) \mid \blacktriangle(\tau)$

Figure 1. Syntax of @ λ . Variables x are used in expressions and internal terms and are distinct from type-level variables, t . Names FAM are type family names (we assume that globally unique type family names can be generated by some external mechanism) and *op* are operator family names. “str” denotes string literals, \bar{z} denotes integer literals and \oplus stands for binary operations over integers.

In either case, we encounter the same basic issue: there is no way to modularly add new primitive type families and operators and implement their associated typechecking and translation logic.

A number of language mechanisms have been proposed that allow new cases to be added to datatypes and the functions that operate over them in a modular manner. In functional languages, we might use *open datatypes*. For example, if we wish to extend Gödel's T with product types and we have written our compiler in a language supporting open inductive datatypes, it might be possible to add new cases like this:

```
newcase Prod of Type * Type extends Type
newcase Pair of Exp * Exp extends Exp      (* Intro *)
newcase PrL of Exp extends Exp              (* Elim Left *)
newcase PrR of Exp extends Exp              (* Elim Right *)
```

The logic for functionality like typechecking and translation could then be implemented for only these new cases. For example, the typeof function that assigns a type to an expression could be extended like so:

```
typeof PrL(e) = case typeof e of
  Prod(t1, _) => t1
  _ => raise TypeError("<appropriate error message>")
```

If we allowed users to define new modules containing definitions like these and link them into our compiler, we will have succeeded in creating an externally-extensible compiler, albeit one where safety is not guaranteed (we will return to this point shortly). We have not, however, created an extensible programming language, for two reasons. First, compiler extensions are distributed and activated separately from libraries, so dependencies become more difficult to manage. Second, other compilers for the same language will not necessarily support the same extensions. If our

```

tycon NAT of 1 {schema  $\nabla(\mathbb{Z})$ ; (1)
  opcon z of 1 (empty a  $\llbracket \nabla(0) \text{ as } \text{NAT}(\langle \rangle) \rrbracket$ ); (2)
  opcon s of 1 (pop_final a  $\lambda x: \text{!Tm}.\lambda t: \star.$  (3)
    check_type t  $\text{NAT}(\langle \rangle) \llbracket \nabla(\Delta(x) + 1) \text{ as } \text{NAT}(\langle \rangle) \rrbracket$ ); (4)
  opcon rec of 1 (pop a  $\lambda x1: \text{!Tm}.\lambda t1: \star.\lambda a: \text{list}[\text{Elab}].$  (5)
    pop a  $\lambda x2: \text{!Tm}.\lambda t2: \star.\lambda a: \text{list}[\text{Elab}].$  (6)
    pop_final a  $\lambda x3: \text{!Tm}.\lambda t3: \star.$  (7)
    check_type t1  $\text{NAT}(\langle \rangle)$  ( (8)
      check_type t3  $\text{ARROW}(\langle \text{NAT}(\langle \rangle), \text{ARROW}(\langle t2, t2 \rangle))$  (9)
       $\llbracket \nabla((\text{fix } f: \mathbb{Z} \rightarrow \text{rep}(t2) \text{ is } \lambda x: \mathbb{Z}.$  (10)
        if  $x \equiv_{\mathbb{Z}} 0$  then  $\Delta(x2)$  else  $\Delta(x3) (x - 1) (f (x - 1))) \Delta(x1) \text{ as } t2 \rrbracket$ ) (11)
    ); (12)
  def nat :  $\star = \text{NAT}(\langle \rangle)$ ; (13)
  ( $\lambda \text{plus}: \text{ARROW}(\langle \text{nat}, \text{ARROW}(\langle \text{nat}, \text{nat} \rangle)) . \lambda \text{two}: \text{nat}.$  (14)
     $\text{ARROW}.\text{ap}(\langle \rangle)(\text{ARROW}.\text{ap}(\langle \rangle)(\text{plus}; \text{two}); \text{two})$  (15)
    ( $\lambda x: \text{nat}.\lambda y: \text{nat}.\text{NAT}.\text{rec}(\langle \rangle)(x; y; \lambda p: \text{nat}.\lambda r: \text{nat}.\text{NAT}.\text{s}(\langle \rangle)(r))$  (16)
     $\text{NAT}.\text{s}(\langle \rangle)(\text{NAT}.\text{s}(\langle \rangle)(\text{NAT}.\text{z}(\langle \rangle)())$  (17)

```

Figure 2. Gödel’s T in @λ, used to calculate 2+2. Helper functions for working with lists (**empty**, **pop**, **pop_final**) and types (**check_type**), described below, are given in the appendix. We will refer to this program as ρ_{nat} in the text.

newly-introduced constructs are exposed at a library’s interface boundary, clients using different compilers face the same problems with interoperability that those using different languages face. That is, extending a language by extending a single compiler for it is morally equivalent to creating a new language. Several prominent language ecosystems today are in a state where a prominent compiler has introduced or enabled the introduction of extensions that many libraries have come to rely on, including the Glasgow Haskell Compiler, SML/NJ and the GNU compilers for C and C++.

A more appropriate and useful place for extensions like this is directly within libraries, alongside abstractions that can be adequately implemented in terms of existing primitive abstractions. To enable this, the language must allow for the introduction new primitive type families, like *Prod*, operators, like *Pair*, *PrL* and *PrR*, and associated typechecking and translation logic. When encountering these new operators in expressions, the compiler must effectively hand control over typechecking and translation to the appropriate user-defined logic. Because this mechanism is language-internal, all compilers must support it to satisfy the language specification.

Statically-typed languages typically make a distinction between *expressions*, which describe run-time computations, and type-level constructs like types, type aliases and datatype declarations. The design described above suggests we may now need to add another layer to our language, an extension language, where extensions can be declared and implemented. In fact, we will show that **the most natural place for type system extensions is within the type-level language**. The intuition is that extensions to a statically-typed language’s semantics will need to manipulate types as values at compile-time. Many languages already allow users to write type-level functions for various reasons, effectively supporting this notion of types as values at compile-time (see Sec. 6 for examples). The type-level language is often constrained by its own type system (where the types of type-level values are called *kinds* for clarity) that prevents type-level functions from causing problems

during compilation. This is precisely the structure that a distinct extension layer would have, and so it is quite natural to unify the two, as we will show in this work.

3. @λ

In this section, we will develop a core calculus, called @λ for the “actively-typed lambda calculus”, by way of a semantics and a simple example, and discuss how it addresses the safety concerns that arise.

3.1 Overview

The grammar of @λ is shown in Figure 1. A program, ρ , consists of a series of declarations followed by an expression. Declarations can be either bindings of type-level terms to type-level variables using **def** or a primitive type family declared using **family**, which can contain implementations of one or more operator families, θ . Expressions, e , can be either variables, lambdas, or applications of operators.

The language is structured as a simply-typed lambda calculus with simply-kinded type-level computation. Kinds, κ , classify type-level terms, τ . Types are type-level values of kind \star (following System F_{ω}) and classify expressions. The type-level language also includes other kinds of terms: type-level functions, lists (required by our mechanism), integers, strings and products for the sake of our examples (see Sec. 4 for a discussion on other acceptable kinds of type-level data) and constructs for developing extensions – denotations and reified internal terms and types – which we will discuss in the sections below.

All expressions are given meaning by translation to a typed internal language. This language has been chosen, for simplicity, to be a variant of Plotkin’s PCF with primitive integers and products, but in practice would include other constructs consistent with its role as a high-level intermediate language. The grammar of internal terms, γ , and internal types, σ , also includes special forms containing type-level terms; these are used for developing extensions and during compilation and will be erased before compilation ends.

3.2 Example: Gödel’s T as an Active Type Family

To make our explanation of each of the constructs in the calculus concrete, we will work through an example showing how to introduce primitive natural numbers with bounded recursion in the style of Gödel’s T [8]. These will be implemented internally as integers (that is, internal terms of internal type \mathbb{Z}). Figure 2 shows how to define the indexed type family *NAT*. This family contains only one type, written $\text{NAT}(\langle \rangle)$, which we alias on line 13 by defining the type-level variable **nat**. We define the typechecking and translation logic for the operators associated with this family (**z**, **s** and **rec**) on lines 2-11 and use these to define a *plus* function on line 16 and compute 2+2. We will refer back to this figure as we describe each construct below.

3.3 Indexed Type Families and Types

The syntactic form **tycon** *TYCON* of κ_{idx} {schema τ_{rep} ; θ } declares a new primitive type family named *FAM* indexed by type-level values of kind κ_{idx} with representation schema $\mathbf{i}.\tau_{\text{rep}}$ and operators θ . The purpose of the representation schema and of associating of operators directly with types will be explained below.

Declaring a type family in this way is a language-internal analog to adding a new constructor to the compiler-internal datatype *Type*, as suggested in Sec. 2. The index represents the data associated with this constructor. A type (that is, a type-level term of kind \star) is constructed by naming a family in scope and providing a type-level term of the appropriate kind as an index. A base type like **nat** can be thought of as being the only type in the family *NAT*

trivially indexed by the unit value, of kind 1, while families like `NTUPLE` might be indexed by a list of types, having kind `list[*]`. For example, `NTUPLE(nat :: nat :: [*])` might be the type of a pair of natural numbers. Given a type, its family can be case analyzed to extract the value of its index. It is important that type equality be decidable, so only kinds for which equivalence coincides with syntactic equality can be used as type family indices. The main consequence of this restriction is that indices cannot contain type-level functions.

3.4 Representations and Representation Schemas

As we will discuss further below, it is important that all expressions classified by a type compile to consistently-typed internal terms. For this reason, we require that every type have a single internal type associated with it, called its *representation*. This is computed by substituting the type index for the bound variable `i` in the term τ_{rep} , called the *representation schema* of the type family, and evaluating to a value representing an internal type. Internal types, σ , are reified as type-level terms of kind `ITy` using the introductory form $\nabla(\sigma)$.

3.5 Indexed Operator Families and Denotations

Type families are also equipped with a collection of primitive operator families. An operator family named `op` is declared using the form `opcon op of κ_{idx} (τ_{op})`. Like type families, operator families are indexed by values of some kind, κ_{idx} , but because operators are not first-class type-level values in our calculus, there are no equality restrictions. In the example in Fig. 2, all the operators are trivially indexed by the kind 1, so each family only contains one operator. However, a type family like `NTUPLE` would be equipped with a family of projection operators, `pr`, indexed by a position (e.g. an integer in our calculus). A family implementing record types or object types might have a similar operator indexed by a type-level string representing the field being accessed (see Sec. 5).

To apply an operator, the grammar provides a uniform form of expression: `FAM.op(τ_{idx})($e_1; \dots; e_n$)`, where $n \geq 0$. The type-checking and translation of an expression of this form is controlled by the term τ_{op} in the operator’s declaration. This term must evaluate to a *denotation*, which is a type-level value of kind `Elab`, when given the operator index, τ_{idx} , and a list constructed from the denotations recursively assigned to each argument, e_1 through e_n .

There are two forms of denotations that an expression can be assigned. A *valid denotation* has the form $\llbracket \tau_{\text{term}} \text{ as } \tau_{\text{type}} \rrbracket$, where τ_{term} is the *translation* of the expression to an internal term and τ_{type} is the type it has been assigned. Internal terms are represented as type-level terms using the form $\nabla(\gamma)$, and have kind `ITm` (similar to `ITy`). If a type error is detected by an operator, the *error denotation*, `error`, is returned instead of a valid denotation. In a practical implementation, a specialized error message and other diagnostic information would be provided when returning error, but we omit such details for simplicity. Terms of kind `Elab` can be case analyzed to determine if they are valid or errors, and if valid, to extract the translation and type.

In the example in Fig. 2, the operator `z` checks that no arguments were passed in using the simple helper function `empty`. If so, it returns a valid denotation by pairing the translation $\nabla(0)$ with the type `NAT()`, as expected¹. If an argument was provided, the helper function returns error. The successor operator, `s`, takes one argument, so it pops a denotation off the argument list, making sure there are no more, and binds its translation and type to `x` and `t` respectively, all using the helper function `pop.final`. It then checks that the argument’s type is also `NAT()`, returning a denotation pairing the translation $\nabla(\Delta(x) + 1)$ with the type `NAT()` if so.

¹ Actually, there is no theoretical barrier to a different “zero” being used!

$$\frac{\text{Kind Checking} \quad \frac{\emptyset \vdash_{\Sigma_0} \rho \text{ prog}}{\rho \rightarrow \gamma} \quad \text{Active Typechecking and Translation} \quad \frac{\vdash_{\Phi_0} \rho \Rightarrow \gamma}{\rho \rightarrow \gamma}}{\rho \rightarrow \gamma}$$

Figure 3. Central Compilation Judgement of `@λ`.

The form $\Delta(\tau)$ is used to “un-reify” reified internal terms, of kind `ITm` (thus serving as the left-inverse of $\nabla(\gamma)$). In this case, `x` is a type-level variable of kind `ITm` representing the translation of the argument to the successor operator, so we simply need to add one to it. Because we have checked that the denotation’s type was `NAT()`, and the representation schema will guarantee that expressions of this type always translate to integers, we know that it is safe to do so. If any of these steps fail, the various helper functions we use simply return error (in practice, it would be prudent to equip each failure condition with a different error message). Compilation will also fail if we accidentally violate the representation schema, as we will see in Sec. 3.8.

3.6 Functions, Variables and Arrow Types

The recursor operator, `rec`, proceeds similarly, extracting the translations and types of each of its three arguments. Of note, however, is how it handles the third argument, which binds two variables (the predecessor and the result of recursing on it). In `@λ`, the built-in `λ` operator serves as the sole mechanism for introducing bound variables. In most calculi, lambda terms have types of the form $\tau_1 \rightarrow \tau_2$. In `@λ`, \rightarrow corresponds to the built-in type family, `ARROW`. This family is indexed by a pair of types, $\star \times \star$, and is always in scope. Its representation schema simply maps to the corresponding internal arrow type, `i.▼(rep(fst(i)) → rep(snd(i)))`. The special form `rep(τ)` is used to refer, abstractly, to the representation of the type τ (we also see this used on line 10). Although the `λ` operator is built-in, because it needs to bind variables, application is just an operator associated with the `ARROW` family, `ap`, as is seen on line 15. The details are straightforward and given in the appendix.

3.7 Compilation

The *central compilation judgement*, shown in Figure 3, captures the two phases of compilation: kind checking and active typechecking and translation. The first phase ensures that all type-level terms in the program are well-kinded and that all expressions and internal terms are closed. The kinding rules for programs are given in Figure 4, and they rely on the kinding rules for type-level terms given in Figure 6. These rules use contexts Σ and Θ to track family and operator signatures, but beyond that, the rules are largely consistent with those of a simply-typed lambda calculus shifted into the type-level, with the addition of the handful of special forms constrained as described in the previous sections. The reader is encouraged to verify that the example in Fig. 2 is well-kinded. The second phase of compilation involves invoking the logic implemented by user-defined operator families to typecheck and translate the program. The rules for this phase are given in Fig. 5. The context Φ tracks the operator definitions and representation schemas associated with families in scope. Because the logic inside operators must be invoked during this phase, we need an evaluation semantics for type-level terms, given in Fig. 7. This is again a largely unsurprising collection of rules with the exception of `DENCASE-EVAL-VALID`, which we will explain below.

3.8 Abstract Representations and Abstract Internal Types

To typecheck and translate an expression, `e`, the rule `ATT-EXP` first assigns a type, τ , and an *abstract translation*, γ_{abs} , to it, as deter-

$$\begin{array}{c}
\boxed{\Delta \vdash_{\Sigma} \rho \text{ prog}} \quad \Delta ::= \emptyset \mid \Delta, \mathbf{t} : \kappa \quad \Sigma ::= \Sigma_0 \mid \Sigma, \text{FAM}[\kappa_{\text{idx}}, \Theta] \\
\\
\text{FAMILY-KINDING} \\
\frac{\text{FAM} \notin \text{dom}(\Sigma) \quad \kappa_{\text{idx}} \text{ eq} \quad \Delta \vdash_{\Sigma, \text{FAM}[\kappa_{\text{idx}}, \Theta]} \theta : \Theta \quad \Delta, \mathbf{i} : \kappa_{\text{idx}} \vdash_{\Sigma, \text{FAM}[\kappa_{\text{idx}}, \Theta]} \tau : \text{!Ty} \quad \Delta \vdash_{\Sigma, \text{FAM}[\kappa_{\text{idx}}, \Theta]} \rho \text{ prog}}{\Delta \vdash_{\Sigma} \text{tycon TYCON of } \kappa_{\text{idx}} \{ \text{schema } \tau_{\text{rep}}; \theta \}; \rho \text{ prog}} \\
\\
\text{DEF-KINDING} \quad \frac{\Delta \vdash_{\Sigma} \tau : \kappa \quad \Delta, \mathbf{t} : \kappa \vdash_{\Sigma} \rho \text{ prog}}{\Delta \vdash_{\Sigma} \text{def } \mathbf{t} : \kappa = \tau; \rho \text{ prog}} \quad \text{EXP-KINDING} \quad \frac{\Delta \emptyset \vdash_{\Sigma} e \text{ expr}}{\Delta \vdash_{\Sigma} e \text{ prog}} \\
\\
\boxed{\Delta \vdash_{\Sigma} \theta : \Theta} \quad \Theta ::= \text{op}[\kappa_{\text{idx}}] \mid \Theta, \Theta \\
\\
\text{OP-KINDING} \\
\frac{\Delta, \mathbf{i} : \kappa_i, \mathbf{a} : \text{list}[\text{Elab}] \vdash_{\Sigma} \tau : \text{Elab}}{\Delta \vdash_{\Sigma} \text{opcon op of } \kappa_{\text{idx}} (\tau) : \text{op}[\kappa_{\text{idx}}]} \\
\\
\text{OPS-KINDING} \\
\frac{\Delta \vdash_{\Sigma} \theta_1 : \Theta_1 \quad \Delta \vdash_{\Sigma} \theta_2 : \Theta_2 \quad \text{dom}(\theta_1) \cap \text{dom}(\theta_2) = \emptyset}{\Delta \vdash_{\Sigma} \theta_1; \theta_2 : \Theta_1, \Theta_2} \\
\\
\boxed{\Delta \Omega \vdash_{\Sigma} e \text{ expr}} \quad \Omega ::= \emptyset \mid \Omega, x \\
\\
\text{E-VAR-KINDING} \quad \frac{}{\Delta \Omega, x \vdash_{\Sigma} x \text{ expr}} \quad \text{E-LAM-KINDING} \quad \frac{\Delta \vdash_{\Sigma} \tau : \star \quad \Delta \Omega, x \vdash_{\Sigma} e \text{ expr}}{\Delta \Omega \vdash_{\Sigma} \lambda x : \tau. e \text{ expr}} \\
\\
\text{E-OP-KINDING} \\
\frac{\text{FAM}[\kappa_{\text{idx}}, \Theta] \in \Sigma \quad \text{op}[\kappa_i] \in \Theta \quad \Delta \vdash_{\Sigma} \tau_i : \kappa_i \quad \Delta \Omega \vdash_{\Sigma} e_1 \text{ expr} \quad \dots \quad \Delta \Omega \vdash_{\Sigma} e_n \text{ expr}}{\Delta \Omega \vdash_{\Sigma} \text{FAM.op}(\tau_i)(e_1; \dots; e_n) \text{ expr}}
\end{array}$$

Figure 4. Kinding for programs. Variable contexts Δ and Ω obey standard structural properties. Kinding rules for type-level terms are given in Figure 6.

mined by the judgement $\Gamma \vdash_{\Phi} e : \tau \Rightarrow \gamma_{\text{abs}}$. It then *deabstracts* this abstract translation to complete the compilation process, as determined by the judgement $\vdash_{\Phi} \gamma_{\text{abs}} \not\leq \gamma$. The purpose of the abstract translation phase is to ensure that the implementation details of type families are not exposed to other families, so that invariants that they rely on are preserved when families are composed. For example, our implementation of natural numbers as integers in Fig. 2 maintains the invariant that the translation is non-negative. If the knowledge that natural numbers are implemented as integers was externally visible, a different extension could introduce an operator like `opcon badnat of 1 (const a [∇(−1) as NAT⟨()⟩])`, breaking this invariant (and thus the guarantee that the recursor always terminates!)

We preclude such operations by a mechanism similar to the abstract type mechanism supported by ML-style module systems [8]. By keeping a type family’s representation schema private to the operators associated with it, they maintain full control over representation invariants. So, because it cannot be shown given the knowledge available in the type family containing `badnum` that the internal type associated with `NAT⟨()⟩` is \mathbb{Z} , the translation `−1` will not be permitted according to the rules we will give below. The only way to produce a term of type `NAT⟨()⟩` as the result of applying an operator not associated with the family `NAT` is if that operator extracts the term from an argument (e.g. when projecting it out of an n -tuple), in which case the necessary invariants are inductively maintained (see Sec. 4). Translations extracted from denotations via case analysis are tracked during this phase as *abstract internal terms*, `abs([τ1 as τ2])` (see rule `DENCASE-EVAL-VALID`).

Let us first review how expressions are assigned types and abstract translations. Variables (ATT-VAR) translate directly to variables in the internal language, with the type determined by the typing context, Γ . Lambda terms (ATT-LAM) also translate to

$$\begin{array}{c}
\boxed{\vdash_{\Phi} \rho \Rightarrow \gamma} \quad \Phi ::= \Phi_0 \mid \Phi, \text{FAM}[\theta, \mathbf{i}. \tau_{\text{rep}}] \\
\\
\text{ATT-FAM} \\
\frac{\vdash_{\Phi, \text{FAM}[\theta, \mathbf{i}. \tau_{\text{rep}}]} \rho \Rightarrow \gamma}{\vdash_{\Phi} \text{tycon TYCON of } \kappa_{\text{idx}} \{ \text{schema } \tau_{\text{rep}}; \theta \}; \rho \Rightarrow \gamma} \\
\\
\text{ATT-DEF} \\
\frac{\tau \Downarrow \tau' \quad \vdash_{\Phi} [\tau'/\mathbf{t}] \rho \Rightarrow \gamma}{\vdash_{\Phi} \text{def } \mathbf{t} : \kappa = \tau; \rho \Rightarrow \gamma} \\
\\
\text{ATT-EXP} \\
\frac{\emptyset \vdash_{\Phi} e : \tau \Rightarrow \gamma_{\text{abs}} \quad \vdash_{\Phi} \gamma_{\text{abs}} \not\leq \gamma}{\vdash_{\Phi} e \Rightarrow \gamma} \\
\\
\boxed{\Gamma \vdash_{\Phi} e : \tau \Rightarrow \gamma_{\text{abs}}} \quad \Gamma ::= \emptyset \mid \Gamma, x : \tau \\
\\
\text{ATT-VAR} \\
\frac{}{\Gamma, x : \tau \vdash_{\Phi} x : \tau \Rightarrow \gamma} \\
\\
\text{ATT-LAM} \\
\frac{\tau_1 \Downarrow \tau'_1 \quad \Gamma, x : \tau'_1 \vdash_{\Phi} e : \tau_2 \Rightarrow \gamma_{\text{abs}}}{\Gamma \vdash_{\Phi} \lambda x : \tau_1. e : \text{ARROW}(\langle \tau'_1, \tau_2 \rangle) \Rightarrow \lambda x : \text{rep}(\tau'_1). \gamma_{\text{abs}}} \\
\\
\text{ATT-OP} \\
\frac{\text{FAM}[\theta, \mathbf{i}. \tau_{\text{rep}}] \in \Phi \quad \text{opcon op of } \kappa_{\text{idx}} (\tau_{\text{op}}) \in \Theta \quad \tau_{\text{idx}} \Downarrow \tau'_{\text{idx}} \quad \Gamma \vdash_{\Phi} e_1 : \tau_1 \Rightarrow \gamma_1 \quad \dots \quad \Gamma \vdash_{\Phi} e_n : \tau_n \Rightarrow \gamma_n}{\left[\llbracket \nabla(\gamma_1) \text{ as } \tau_1 \rrbracket :: \dots :: \llbracket \nabla(\gamma_n) \text{ as } \tau_n \rrbracket :: \llbracket \nabla(\gamma_{\text{abs}}) \text{ as } \tau \rrbracket \right] \tau_{\text{op}} \Downarrow \llbracket \nabla(\gamma_{\text{abs}}) \text{ as } \tau \rrbracket} \\
\frac{\vdash_{\Phi}^{\Xi_0, \text{FAM}} \text{rep}(\tau) \rightsquigarrow \sigma_{\text{abs}} \quad \vdash_{\Phi}^{\Xi_0, \text{FAM}} \gamma_{\text{abs}} \sim \sigma_{\text{abs}}}{\Gamma \vdash_{\Phi} \text{FAM.op}(\tau_{\text{idx}})(e_1; \dots; e_n) : \tau \Rightarrow \gamma_{\text{abs}}}
\end{array}$$

Figure 5. Active typechecking and translation of programs. Evaluation semantics for type-level terms are given in Fig. 7.

lambda terms in the internal language and are given the appropriate `ARROW` type by extending the typing context and proceeding recursively as is usual. The internal type of the argument, however, is left as `rep(τ1)`, which is called the *abstract representation* of `τ1`. The actual representation of this type is only available from within operators in its family.

Now we will consider the important operator application rule (ATT-OP). This rule operates by extracting the appropriate operator definition, evaluating the operator index to a value, `τidx`, and recursively assigning a type and abstract translation to each argument. From these, it constructs a list of denotations and passes this list, along with the fully evaluated operator index, to the operator implementation, `τop`. If this results in a valid denotation, compilation can proceed, but if an error is produced, compilation will stop because no other rule will apply (in practice, we would display a type error at this point).

Before a valid denotation produced in this way can be used, however, we check that it is *representationally consistent*, meaning that the abstract translation, `γabs`, is of an abstract internal type consistent with the abstraction representation of the type it is paired with, `τ`. The premise $\vdash_{\Phi}^{\Xi_0, \text{FAM}} \text{rep}(\tau) \rightsquigarrow \sigma_{\text{abs}}$ determines the abstract internal type and the premise $\Psi \vdash_{\Phi}^{\Xi_0, \text{FAM}} \gamma_{\text{abs}} \sim \sigma_{\text{abs}}$ performs the internal type checking. To do so, the variable context, Γ , which associates variables with types, must be converted to a context associating those variables with their corresponding abstract internal types, Ψ . The relevant judgements are defined in Fig. 8. The schema context Ξ is used to track which representation schemas are visible. In our calculus, $\Xi_0 = \text{ARROW}$ is always visible alongside the schema of the family associated with the operator being considered. This could be expanded to support module-scoped visibility (we do not include modules in our calculus for simplicity),

mutually-defined type families or type families that intentionally expose their representation schemas publicly (as `ARROW` does).

To clarify this fundamental mechanism let us examine how the expression $\text{NAT.s}(\langle \rangle)(\text{NAT.z}(\langle \rangle)(\langle \rangle))$ will be processed during this phase. The inner application of z will produce the abstract translation 0 paired with the type $\text{NAT}(\langle \rangle)$. Because the representation schema for NAT is available, we have that $\vdash_{\Phi}^{\Xi_0, \text{NAT}} \text{rep}(\text{NAT}(\langle \rangle)) \rightsquigarrow \mathbb{Z}$ by `SHOW-REP`, as needed. In contrast, if *badnat* were used inside, the representation schema of NAT would not be available, so we can only apply `HIDE-REP`, which does not give us enough information to give -1 a type.

When the compiler next considers the outer application of s , it will pass the denotation $\llbracket \nabla(0) \text{ as } \text{NAT}(\langle \rangle) \rrbracket$ into its definition. There, it binds the translation to the variable x (within the `pop_final` helper function). However, x is not simply $\nabla(0)$. Instead, its provenance is tracked by using the form $\text{abs}(\llbracket \nabla(0) \text{ as } \text{NAT}(\langle \rangle) \rrbracket)$. When attempting to check that $\Delta(x) + 1$ is valid, the `SHOW-TRANS` rule will reveal that it is an integer because, again, the appropriate representation schema is available. If the schema were not available, the most that could be derived is that $\Psi \vdash_{\Phi}^{\Xi} \Delta(x) \sim \text{rep}(\text{NAT}(\langle \rangle))$. The fact that natural numbers are represented using integers is not derivable (though it is true), so the addition operation would fail to type. This fact would, however, be sufficient for implementing families, like `NTUPLE`, that do not require knowledge about a type's representation. We encourage the reader to derive these judgements for the logic in the *rec* operator to strengthen their understanding of this fundamental mechanism.

We cannot leave abstract representations and abstract internal terms in the result of compilation, so a final deabstraction phase erases these, replacing them with their underlying internal types and terms in all contexts. Just as with abstract types in modules or existential types, there is no run-time overhead to this mechanism – programs run at full speed. We give the deabstraction rules in the appendix due to their simplicity.

4. Safety of $@\lambda$

In giving users of a language direct influence on the typechecking and translation of expressions, it is essential to consider safety properties. It must not be possible for well-typed programs to fail to finish compiling or go wrong at run-time because of a buggy extension. It is also considered desirable for typechecking to be decidable. And for extensions to be reliable, they must not be allowed to interfere with one another under any circumstance. By carefully designing our extension mechanism, we believe we have achieved each of these goals.

4.1 Representational Consistency and Type Safety

We do not need to give a semantics to internal terms and internal types that do not survive deabstraction. The remaining terms and types form a variant of PCF with well-understood primitives (here, integers and binary products) known to be type safe [8]. If compilation can be shown to always result in a well-typed term of this language, then type safety follows by composing these two facts. Fortunately, this fact is a corollary of our representational consistency mechanism, described in Sec. 3.8. The proof is by straightforward (though deeply nested!) rule induction.

4.2 Decidability

To show that compilation is decidable, we need to show that both kind checking and active typechecking and compilation are decidable. This hinges on showing that the type-level language is type-safe and that evaluation of type-level terms always terminates.

The proof is straightforward because the type-level language is based on a conventional simply-typed lambda calculus with only

bounded recursion over lists, so standard techniques (e.g. logical relations) can be used directly. The only new constructs are the types, denotations and reified internal language forms, but because none of these can contain type-level functions by the statics, there is no risk of introducing self-reference by their inclusion. If the type-level language included constructs like inductive datatypes (without a strict positivity condition) or general recursion, this theorem would be weakened.

4.3 Non-Interference

Our abstraction mechanism guarantees that the representation invariants collectively maintained by the operators associated with a type family cannot be violated by operators in other type families, by ensuring that introductory forms cannot be defined outside the family. One way to state this is as follows:

The proof relies on the fact that the show/hide rules are mutually exclusive and so there are two cases to consider for each form of expression, and weakening properties of the family contexts. This powerful theorem allows us to compose type families arbitrarily without needing to handle cases in our implementation that are ruled out based on a local analysis. The details will be provided in the appendix.

5. Use Cases

Aldrich argues that dynamic dispatch has proven useful for building large software frameworks [?]. Object-oriented languages support dynamic dispatch directly but encodings of dynamic dispatch in terms of standard functional primitives can be awkward and difficult to reason about (e.g. [8] Ch. X).

Indexed type families and indexed operator families are a powerful tool for programming language semantics. In *Practical Foundations for Programming Languages*, for example, nearly every chapter defines a new collection of indexed types and operators and gives their semantics in isolation from the constructs in the other chapters [8]. For example, n -ary sums and products are families indexed by a list of types. Labeled variants of these simply pair the labels with the types as static strings. Even nested pattern matching, such as that found in modern functional languages, can be understood as an operator family indexed by a series of patterns, which can be represented as type-level data. In the appendix, we show how to implement sums and products in $@\lambda$.

The book also shows constructs as varied as dynamic types as well as a number of constructs for parallelism and concurrency, such as futures, promises and actors. With a sufficiently capable internal language (exposing basic concurrency primitives, for example), each of these could be fully implemented as libraries using our mechanism. Few languages include primitive support for several such abstractions, even though they are often useful together. In a preliminary implementation of this calculus, Ace, which is beyond the scope of this paper, we have implemented the entirety of the OpenCL programming language as a library, along with primitives supporting partitioned global address spaces.

Object systems too could be implemented using this mechanism, with indices serving to capture the inheritance data and the signatures. Operator families for reading and writing fields and/or sending messages would be parameterized by strings naming them. The operator definition would simply search the signatures going up the inheritance hierarchy, and implement objects in a conventional way (e.g. using a v-table together with a pointer to a structure containing field data). A variety of object systems could coexist within the same language. Another related example would be to implement a safe and efficient interoperability layer with an existing OO language by capturing its type system as an extension, to be used only at language boundaries.

Finally, a variety of domain-specific type systems, capturing complex rule systems like the system of scientific units of measure (a family indexed by the measure being used and the type which the measure is being applied to), regular expressions [?] (a family indexed by the number of captured groups) or XML [?] (indexed by a document schema) would be possible. Recent work examined a specialized type system capturing the semantics of the widely-used jQuery javascript library [?]. Today, these all require *ad hoc* language-external solutions, or the development of new languages.

6. Related Work

Our representation schema abstraction mechanism relates closely to abstract and existential types [8?]. Our calculus enforces the abstraction barriers in a purely syntactic manner, as in previous work on syntactic type abstraction [?]. While this work is all focused on abstracting away the identity of a particular type outside of the “principal” it is associated with (e.g. a module), we focus on abstracting away the knowledge of how a primitive type family is implemented outside of a limited scope.

The representational consistency mechanism brings into the language work on typed compilation, especially work done on Standard ML in the TILT project [?]. Indeed, the specification of Standard ML is structured around a typed internal language and a judgement that assigns a type and an internal term to each expression [?]. Representational consistency is related to the notion of type-directed compilation in this work.

Type-level computation is supported in some form by a growing number of languages. For example, Haskell supports a simple form of it [2]). Ur uses type-level records and names to support type-safe metaprogramming, with applications to web programming [4]. Omega adds algebraic data types at the type-level, using these to increase the expressive power of algebraic data types at the expression level [12]. Dependently-typed languages blur the traditional phase separation between types and expressions, so type-level computation is often implicitly used (though not always in its most general form, e.g. Deputy [3], ATS [5]). We show how to integrate language extensions into the type-level language, drawing on ideas about [13].

Many languages and tools allow developers to rewrite expressions according to custom rules. These can broadly be classified as *term rewriting systems*. Macro systems, such as those characteristic of the LISP family of languages [10], are the most prominent example. Some compile-time metaprogramming systems also allow users to manipulate syntax trees (e.g. MetaML [11]), and external rewrite systems also exist for many languages. These facilities differ from AT&T in that they involve direct manipulation of terms, while AT&T involves extending typechecking and translation logic directly. We also draw a distinction between the type-level, used to specify types and compile-time logic, the expression grammar, used to describe run-time behavior, and the internal language, used to implement this behavior. By doing so, each component can be structured and constrained as appropriate for its distinct role, as we have shown.

Previous work on extensible languages has suffered from problems with either expressiveness or safety. For example, a number of projects, such as SugarJ [6], allow for user-defined desugarings (and indeed, our system would clearly benefit from integration with such a mechanism), but this does not allow the semantics to be fundamentally extended nor for implementation details to be hidden. Recent variants of this work has investigated introducing new typing rules, but only if they are admissible by the base type system [?]. Our work allows for entirely new logic to be added to the system, requiring only that the implementation of this logic respect the internal type system. A number of other extensible languages (e.g. Xroma [?]) and compilers do allow new static semantics to

be introduced, but in an unconstrained manner based on global pattern matching and rewriting, leading to significant problems with interference and safety. Our work aims to provide a sound and safe underpinning, based around well-understood concepts of type and operator families, to language extensibility.

In the future, we will investigate mechanisms to enable type systems with specialized binding and scoping rules, as well as integration of dependent kinds to support mechanized verification of key properties like representational consistency and adequacy against a declarative specification at kind-checking time. We will also investigate mechanisms that enable a more natural syntax (e.g. Wyvern’s type-directed syntax [?]).

References

- [1] H. Chafi, A. K. Sujeeth, K. J. Brown, H. Lee, A. R. Atreya, and K. Olukotun. A domain-specific approach to heterogeneous parallelism. In *Proceedings of the 16th ACM symposium on Principles and practice of parallel programming*, PPoPP ’11, pages 35–46, New York, NY, USA, 2011. ACM.
- [2] M. M. T. Chakravarty, G. Keller, S. P. Jones, and S. Marlow. Associated types with class. *ACM SIGPLAN Notices*, 40(1):1–13, Jan. 2005.
- [3] C. Chen and H. Xi. Combining programming with theorem proving. In O. Danvy and B. C. Pierce, editors, *Proceedings of the 10th ACM SIGPLAN International Conference on Functional Programming, ICFP 2005, Tallinn, Estonia, September 26-28, 2005*, pages 66–77. ACM, 2005.
- [4] A. Chlipala. Ur: statically-typed metaprogramming with type-level record computation. In B. G. Zorn and A. Aiken, editors, *Proceedings of the 2010 ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2010, Toronto, Ontario, Canada, June 5-10, 2010*, pages 122–133. ACM, 2010.
- [5] J. Condit, M. Harren, Z. R. Anderson, D. Gay, and G. C. Necula. Dependent types for low-level programming. In R. D. Nicola, editor, *Proceedings of the 16th European Symposium on Programming, ESOP 2007*, volume 4421 of *Lecture Notes in Computer Science*, pages 520–535. Springer, 2007.
- [6] S. Erdweg, T. Rendel, C. Kästner, and K. Ostermann. Sugarj: Library-based syntactic language extensibility. *ACM SIGPLAN Notices*, 46(10):391–406, 2011.
- [7] M. Fowler and R. Parsons. *Domain-Specific Languages*. Addison-Wesley Professional, 2010.
- [8] R. Harper. *Practical foundations for programming languages*. Cambridge University Press, 2012.
- [9] A. Kennedy. Types for units-of-measure: Theory and practice. In Z. Horváth, R. Plasmeijer, and V. Zsóok, editors, *CEFP*, volume 6299 of *Lecture Notes in Computer Science*, pages 268–305. Springer, 2009.
- [10] J. McCarthy. History of lisp. In *History of programming languages I*, pages 173–185. ACM, 1978.
- [11] T. Sheard. Using MetaML: A staged programming language. *Lecture Notes in Computer Science*, 1608:207–??, 1999.
- [12] T. Sheard and N. Linger. Programming in omega. In Z. Horváth, R. Plasmeijer, A. Soós, and V. Zsóok, editors, *CEFP*, volume 5161 of *Lecture Notes in Computer Science*, pages 158–227. Springer, 2007.
- [13] T. L. Veldhuizen and D. Gannon. Active libraries: Rethinking the roles of compilers and libraries. In *Proc. 1998 SIAM Workshop on Object Oriented Methods for Inter-operable Scientific and Engineering Computing*, 1998.
- [14] M. P. Ward. Language-oriented programming. *Software - Concepts and Tools*, 15(4):147–161, 1994.

$\Delta \vdash_{\Sigma} \tau : \kappa$	
VAR-KIND $\Delta, \mathbf{t} : \kappa \vdash_{\Sigma} \mathbf{t} : \kappa$	K-ARROW-INTRO $\Delta, \mathbf{t} : \kappa_1 \vdash_{\Sigma} \tau : \kappa_2$ $\Delta \vdash_{\Sigma} \lambda \mathbf{t} : \kappa_1. \tau : \kappa_1 \rightarrow \kappa_2$
K-ARROW-ELIM $\Delta \vdash_{\Sigma} \tau_1 : \kappa_1 \rightarrow \kappa_2 \quad \Delta \vdash_{\Sigma} \tau_2 : \kappa_1$ $\Delta \vdash_{\Sigma} \tau_1 \tau_2 : \kappa_2$	
(standard statics for lists, integers, strings and products omitted)	
TL-EQUALITY $\kappa \text{ eq} \quad \Delta \vdash_{\Sigma} \tau_1 : \kappa \quad \Delta \vdash_{\Sigma} \tau_2 : \kappa$ $\Delta \vdash_{\Sigma} \tau_3 : \kappa' \quad \Delta \vdash_{\Sigma} \tau_4 : \kappa'$ $\Delta \vdash_{\Sigma} \text{if } \tau_1 \equiv_{\kappa} \tau_2 \text{ then } \tau_3 \text{ else } \tau_4 : \kappa'$	TYPE-INTRO $\text{FAM}[\kappa_{\text{idx}}, \Theta] \in \Sigma$ $\Delta \vdash_{\Sigma} \tau_{\text{idx}} : \kappa_{\text{idx}}$ $\Delta \vdash_{\Sigma} \text{FAM}(\tau_{\text{idx}}) : \star$
TYPE-ELIM $\text{FAM}[\kappa_{\text{idx}}, \Theta] \in \Sigma \quad \Delta \vdash_{\Sigma} \tau : \star$ $\Delta, \mathbf{x} : \kappa_{\text{idx}} \vdash_{\Sigma} \tau_0 : \kappa \quad \Delta \vdash_{\Sigma} \tau_1 : \kappa$ $\Delta \vdash_{\Sigma} \text{case } \tau \text{ of } \text{FAM}(\mathbf{x}) \Rightarrow \tau_0 \text{ ow } \tau_1 : \kappa$	DEN-INTRO-VALID $\Delta \vdash_{\Sigma} \tau_1 : \text{ITm}$ $\Delta \vdash_{\Sigma} \tau_2 : \star$ $\Delta \vdash_{\Sigma} \llbracket \tau_1 \text{ as } \tau_2 \rrbracket : \text{Elab}$
DEN-INTRO-ERR $\Delta \vdash_{\Sigma} \text{error} : \text{Elab}$	DEN-ELIM $\Delta \vdash_{\Sigma} \tau : \text{Elab}$ $\Delta, \mathbf{x} : \text{ITm}, \mathbf{t} : \star \vdash_{\Sigma} \tau_1 : \kappa \quad \Delta \vdash_{\Sigma} \tau_2 : \kappa$ $\Delta \vdash_{\Sigma} \text{case } \tau \text{ of } \llbracket \mathbf{x} \text{ as } \mathbf{t} \rrbracket \Rightarrow \tau_1 \text{ ow } \tau_2 : \kappa$
ITERM-INTRO $\Delta \emptyset \vdash_{\Sigma} \gamma \text{ iterm}$ $\Delta \vdash_{\Sigma} \nabla(\gamma) : \text{ITm}$	ITYPE-INTRO $\Delta \vdash_{\Sigma} \sigma \text{ itype}$ $\Delta \vdash_{\Sigma} \blacktriangledown(\sigma) : \text{ITy}$
$\kappa \text{ eq}$	
T-EQ $\star \text{ eq}$	Z-EQ $\mathbb{Z} \text{ eq}$
S-EQ $\text{Str} \text{ eq}$	U-EQ 1 eq
P-EQ $\kappa_1 \text{ eq} \quad \kappa_2 \text{ eq}$ $\kappa_1 \times \kappa_2 \text{ eq}$	L-EQ $\kappa \text{ eq}$ $\text{list}[\kappa] \text{ eq}$
$\Delta \Omega \vdash_{\Sigma} \gamma \text{ iterm}$	
I-VAR-KINDING $\Delta \Omega, x \vdash_{\Sigma} x \text{ iterm}$	I-LAM-KINDING $\Delta \vdash_{\Sigma} \sigma \text{ itype}$ $\Delta \Omega, x \vdash_{\Sigma} \gamma \text{ iterm}$ $\Delta \Omega \vdash_{\Sigma} \lambda x : \sigma. \gamma \text{ iterm}$
I-FIX-KINDING $\Delta \vdash_{\Sigma} \sigma \text{ itype}$ $\Delta \Omega, x \vdash_{\Sigma} \gamma \text{ iterm}$ $\Delta \Omega \vdash_{\Sigma} \text{fix } f : \sigma \text{ is } \gamma \text{ iterm}$	
(omitted forms have trivially recursive rules)	
ITERM-DEREIFY-KINDING $\Delta \vdash_{\Sigma} \tau : \text{ITm}$ $\Delta \Omega \vdash_{\Sigma} \Delta(\tau) \text{ iterm}$	ABS-TRANS-KINDING $\Delta \vdash_{\Sigma} \tau_{\text{iterm}} : \text{ITm} \quad \Delta \vdash_{\Sigma} \tau_{\text{itype}} : \star$ $\Delta \Omega \vdash_{\Sigma} \text{abs}(\llbracket \tau_{\text{iterm}} \text{ as } \tau_{\text{itype}} \rrbracket) \text{ iterm}$
$\Delta \vdash_{\Sigma} \sigma \text{ itype}$	
I-INT-KINDING $\Delta \vdash_{\Sigma} \mathbb{Z} \text{ itype}$	I-PROD-KINDING $\Delta \vdash_{\Sigma} \sigma_1 \text{ itype}$ $\Delta \vdash_{\Sigma} \sigma_2 \text{ itype}$ $\Delta \vdash_{\Sigma} \sigma_1 \times \sigma_2 \text{ itype}$
I-ARROW-KINDING $\Delta \vdash_{\Sigma} \sigma_1 \text{ itype}$ $\Delta \vdash_{\Sigma} \sigma_2 \text{ itype}$ $\Delta \vdash_{\Sigma} \sigma_1 \rightarrow \sigma_2 \text{ itype}$	ITYPE-DEREIFY-KINDING $\Delta \vdash_{\Sigma} \tau : \text{ITy}$ $\Delta \vdash_{\Sigma} \blacktriangle(\tau) \text{ itype}$
ABS-REP-KINDING $\Delta \vdash_{\Sigma} \tau : \star$ $\Delta \vdash_{\Sigma} \text{rep}(\tau) \text{ itype}$	

Figure 6. Kinding for type-level terms

$\tau \Downarrow \tau'$	
TL-LAM-EVAL $\lambda \mathbf{t} : \kappa. \tau \Downarrow \lambda \mathbf{t} : \kappa. \tau$	TL-AP-EVAL $\tau_1 \Downarrow \lambda \mathbf{t} : \kappa. \tau \quad \tau_2 \Downarrow \tau'_2 \quad [\tau'_2 / \mathbf{t}] \tau \Downarrow \tau'$ $\tau_1 \tau_2 \Downarrow \tau'$
(standard evaluation rules for integers, strings, products and lists omitted)	
TL-EQ-EVAL-EQUAL $\tau_1 \Downarrow \tau'_1 \quad \tau_2 \Downarrow \tau'_1 \quad \tau_3 \Downarrow \tau'_3$ $\text{if } \tau_1 \equiv_{\kappa} \tau_2 \text{ then } \tau_3 \text{ else } \tau_4 \Downarrow \tau'_3$	
TL-EQ-EVAL-INEQUAL $\tau_1 \Downarrow \tau'_1 \quad \tau_2 \Downarrow \tau'_2 \quad \tau'_1 \neq \tau'_2 \quad \tau_4 \Downarrow \tau'_4$ $\text{if } \tau_1 \equiv_{\kappa} \tau_2 \text{ then } \tau_3 \text{ else } \tau_4 \Downarrow \tau'_4$	
TYPE-EVAL $\tau_{\text{idx}} \Downarrow \tau'_{\text{idx}}$ $\text{FAM}(\tau_{\text{idx}}) \Downarrow \text{FAM}(\tau'_{\text{idx}})$	FAMCASE-EVAL-MATCH $\tau \Downarrow \text{FAM}(\tau_{\text{idx}}) \quad [\tau_{\text{idx}} / \mathbf{x}] \tau_1 \Downarrow \tau'_1$ $\text{case } \tau \text{ of } \text{FAM}(\mathbf{x}) \Rightarrow \tau_1 \text{ ow } \tau_2 \Downarrow \tau'_1$
FAMCASE-EVAL-FAIL $\tau \Downarrow \text{FAM}'(\tau_{\text{idx}}) \quad \text{FAM} \neq \text{FAM}' \quad \tau_2 \Downarrow \tau'_2$ $\text{case } \tau \text{ of } \text{FAM}(\mathbf{x}) \Rightarrow \tau_1 \text{ ow } \tau_2 \Downarrow \tau'_2$	
DEN-VALID-EVAL $\tau_1 \Downarrow \tau'_1 \quad \tau_2 \Downarrow \tau'_2$ $\llbracket \tau_1 \text{ as } \tau_2 \rrbracket \Downarrow \llbracket \tau'_1 \text{ as } \tau'_2 \rrbracket$	DEN-ERR-EVAL $\text{error} \Downarrow \text{error}$
DENCASE-EVAL-VALID $\tau \Downarrow \llbracket \tau_{\text{iterm}} \text{ as } \tau_{\text{itype}} \rrbracket$ $[\nabla(\text{abs}(\llbracket \tau_{\text{iterm}} \text{ as } \tau_{\text{itype}} \rrbracket)) / \mathbf{x}, \tau_{\text{itype}} / \mathbf{t}] \tau_1 \Downarrow \tau'_1$ $\text{case } \tau \text{ of } \llbracket \mathbf{x} \text{ as } \mathbf{t} \rrbracket \Rightarrow \tau_1 \text{ ow } \tau_2 \Downarrow \tau'_1$	
DENCASE-EVAL-ERR $\tau \Downarrow \text{error} \quad \tau_2 \Downarrow \tau'_2$ $\text{case } \tau_{\text{den}} \text{ of } \llbracket \mathbf{y} \text{ as } \mathbf{x} \rrbracket \Rightarrow \tau_1 \text{ ow } \tau_2 \Downarrow \tau'_2$	ITERM-REIFY $\gamma \Downarrow \gamma'$ $\nabla(\gamma) \Downarrow \nabla(\gamma')$
ITYPE-REIFY $\sigma \Downarrow \sigma'$ $\blacktriangledown(\sigma) \Downarrow \blacktriangledown(\sigma')$	
$\gamma \Downarrow \gamma'$	
I-VAR-EVAL $x \Downarrow x$	I-LAM-EVAL $\sigma \Downarrow \sigma' \quad \gamma \Downarrow \gamma'$ $\lambda x : \sigma. \gamma \Downarrow \lambda x : \sigma'. \gamma'$
(omitted forms have trivially recursive rules)	
ITERM UNQUOTE EVAL $\tau \Downarrow \tau'$ $\Delta(\tau) \Downarrow \Delta(\tau')$	
VAL FROM DEN EVAL $\tau_{\text{iterm}} \Downarrow \tau'_{\text{iterm}} \quad \tau_{\text{itype}} \Downarrow \tau'_{\text{itype}}$ $\text{abs}(\llbracket \tau_{\text{iterm}} \text{ as } \tau_{\text{itype}} \rrbracket) \Downarrow \text{abs}(\llbracket \tau'_{\text{iterm}} \text{ as } \tau'_{\text{itype}} \rrbracket)$	
$\sigma \Downarrow \sigma'$	
I-INT-EVAL $\mathbb{Z} \Downarrow \mathbb{Z}$	I-PROD-EVAL $\sigma_1 \Downarrow \sigma'_1 \quad \sigma_2 \Downarrow \sigma'_2$ $\sigma_1 \times \sigma_2 \Downarrow \sigma'_1 \times \sigma'_2$
I-ARROW-EVAL $\sigma_1 \Downarrow \sigma'_1 \quad \sigma_2 \Downarrow \sigma'_2$ $\sigma_1 \rightarrow \sigma_2 \Downarrow \sigma'_1 \rightarrow \sigma'_2$	
ITYPE-EVAL $\tau \Downarrow \tau'$ $\blacktriangle(\tau) \Downarrow \blacktriangle(\tau')$	ABS-REP-EVAL $\tau \Downarrow \tau'$ $\text{rep}(\tau) \Downarrow \text{rep}(\tau')$

Figure 7. Evaluation semantics for type-level terms

$$\boxed{\vdash_{\Phi}^{\Xi} \tau \rightsquigarrow \sigma}$$

$$\begin{array}{c}
\text{GET-REP} \\
\frac{\text{FAM}[\theta, \mathbf{i}, \tau] \in \Phi \quad [\tau_{\text{idx}}/\mathbf{i}]\tau \Downarrow \nabla(\sigma) \quad \vdash_{\Phi}^{\Xi} \sigma \rightsquigarrow \sigma_{\text{conc}}}{\vdash_{\Phi}^{\Xi} \text{FAM}\langle \tau_{\text{idx}} \rangle \rightsquigarrow \sigma_{\text{conc}}}
\end{array}$$

$$\boxed{\vdash_{\Phi}^{\Xi} \sigma \rightsquigarrow \sigma'} \quad \Phi ::= \Phi_0 \mid \Phi, \text{FAM}[\theta, \mathbf{i}, \tau_{\text{rep}}] \quad \Xi ::=$$

$$\Xi_0 \mid \Xi, \text{FAM} \quad \Xi_0 := \text{ARROW}$$

$$\begin{array}{c}
\text{ABS-INT} \\
\vdash_{\Phi}^{\Xi} \mathbb{Z} \rightsquigarrow \mathbb{Z}
\end{array}
\quad
\begin{array}{c}
\text{ABS-ARROW} \\
\frac{\vdash_{\Phi}^{\Xi} \sigma_1 \rightsquigarrow \sigma'_1 \quad \vdash_{\Phi}^{\Xi} \sigma_2 \rightsquigarrow \sigma'_2}{\vdash_{\Phi}^{\Xi} \sigma_1 \rightarrow \sigma_2 \rightsquigarrow \sigma'_1 \rightarrow \sigma'_2}
\end{array}$$

$$\begin{array}{c}
\text{ABS-PROD} \\
\frac{\vdash_{\Phi}^{\Xi} \sigma_1 \rightsquigarrow \sigma'_1 \quad \vdash_{\Phi}^{\Xi} \sigma_2 \rightsquigarrow \sigma'_2}{\vdash_{\Phi}^{\Xi} \sigma_1 \times \sigma_2 \rightsquigarrow \sigma'_1 \times \sigma'_2}
\end{array}
\quad
\begin{array}{c}
\text{ITYPE-INVERSE} \\
\frac{\vdash_{\Phi}^{\Xi} \sigma \rightsquigarrow \sigma_{\text{abs}}}{\vdash_{\Phi}^{\Xi} \blacktriangle(\nabla(\sigma)) \rightsquigarrow \sigma_{\text{abs}}}
\end{array}$$

$$\begin{array}{c}
\text{SHOW-REP} \\
\frac{\text{FAM} \in \Xi \quad \vdash_{\Phi}^{\Xi} \text{FAM}\langle \tau_{\text{idx}} \rangle \rightsquigarrow \sigma_{\text{conc}}}{\vdash_{\Phi}^{\Xi} \text{rep}(\text{FAM}\langle \tau_{\text{idx}} \rangle) \rightsquigarrow \sigma_{\text{conc}}}
\end{array}$$

$$\begin{array}{c}
\text{HIDE-REP} \\
\frac{\text{FAM} \notin \Xi}{\vdash_{\Phi}^{\Xi} \text{rep}(\text{FAM}\langle \tau_{\text{idx}} \rangle) \rightsquigarrow \text{rep}(\text{FAM}\langle \tau_{\text{idx}} \rangle)}
\end{array}$$

$$\boxed{\vdash_{\Phi}^{\Xi} \Gamma \rightsquigarrow \Psi} \quad \Psi ::= \emptyset \mid \Psi, x : \sigma$$

$$\begin{array}{c}
\text{ABS-EMPTY} \\
\vdash_{\Phi}^{\Xi} \emptyset \rightsquigarrow \emptyset
\end{array}
\quad
\begin{array}{c}
\text{ABS-CTX} \\
\frac{\vdash_{\Phi}^{\Xi} \Gamma \rightsquigarrow \Psi \quad \vdash_{\Phi}^{\Xi} \text{rep}(\tau) \rightsquigarrow \sigma}{\vdash_{\Phi}^{\Xi} \Gamma, x : \tau \rightsquigarrow \Psi, x : \sigma}
\end{array}$$

$$\boxed{\Psi \vdash_{\Phi}^{\Xi} \gamma \sim \sigma}$$

$$\begin{array}{c}
\text{ABS-I-VAR} \\
\frac{\Psi, x : \sigma \vdash_{\Phi}^{\Xi} x \sim \sigma}{\Psi \vdash_{\Phi}^{\Xi} \lambda x : \sigma_1. \gamma \sim \sigma_1 \rightarrow \sigma_2}
\end{array}
\quad
\begin{array}{c}
\text{ABS-I-LAM} \\
\frac{\vdash_{\Phi}^{\Xi} \sigma_1 \rightsquigarrow \sigma'_1 \quad \Psi, x : \sigma'_1 \vdash_{\Phi}^{\Xi} \gamma \sim \sigma_2}{\Psi \vdash_{\Phi}^{\Xi} \lambda x : \sigma_1. \gamma \sim \sigma'_1 \rightarrow \sigma_2}
\end{array}$$

$$\begin{array}{c}
\text{ABS-I-AP} \\
\frac{\Psi \vdash_{\Phi}^{\Xi} \gamma_1 \sim \sigma_1 \rightarrow \sigma_2 \quad \Psi \vdash_{\Phi}^{\Xi} \gamma_2 \sim \sigma_1}{\Psi \vdash_{\Phi}^{\Xi} \gamma_1 \gamma_2 \sim \sigma_2}
\end{array}$$

$$\begin{array}{c}
\text{ABS-I-FIX} \\
\frac{\vdash_{\Phi}^{\Xi} \sigma \rightsquigarrow \sigma' \quad \Psi, x : \sigma' \vdash_{\Phi}^{\Xi} \gamma \sim \sigma'}{\Psi \vdash_{\Phi}^{\Xi} \text{fix } x : \sigma \text{ is } \gamma_{\text{abs}} \sim \sigma'}
\end{array}$$

(standard statics for integers and products omitted)

$$\begin{array}{c}
\text{ABS-ITERM-INVERSE} \\
\frac{\Psi \vdash_{\Phi}^{\Xi} \gamma \sim \sigma}{\Psi \vdash_{\Phi}^{\Xi} \Delta(\nabla(\gamma)) \sim \sigma}
\end{array}$$

$$\begin{array}{c}
\text{SHOW-TRANS} \\
\frac{\text{FAM} \in \Xi \quad \vdash_{\Phi}^{\Xi} \text{FAM}\langle \tau_{\text{idx}} \rangle \rightsquigarrow \sigma \quad \Psi \vdash_{\Phi}^{\Xi} \gamma \sim \sigma}{\Psi \vdash_{\Phi}^{\Xi} \text{abs}(\llbracket \nabla(\gamma) \text{ as FAM}\langle \tau_{\text{idx}} \rangle \rrbracket) \sim \sigma}
\end{array}$$

$$\begin{array}{c}
\text{HIDE-TRANS} \\
\frac{\text{FAM} \notin \Xi \quad \vdash_{\Phi}^{\Xi} \text{FAM}\langle \tau_{\text{idx}} \rangle \rightsquigarrow \sigma \quad \Psi \vdash_{\Phi}^{\Xi} \gamma \sim \sigma}{\Psi \vdash_{\Phi}^{\Xi} \text{abs}(\llbracket \nabla(\gamma) \text{ as FAM}\langle \tau_{\text{idx}} \rangle \rrbracket) \sim \text{rep}(\text{FAM}\langle \tau_{\text{idx}} \rangle)}
\end{array}$$

Figure 8. Abstracted internal typing