

Modularly Composing Typed Language Fragments

Abstract

Researchers often describe typed programming languages as fragments or simple calculi, leaving to language designers the task of composing these to form complete languages. This is not a systematic process: metatheoretic results must be established anew for each composition, guided only notionally by metatheorems derived for simpler systems. As the design space grows, mechanisms that provide stronger modular reasoning principles than this are needed.

In this paper, we begin from first principles with a core calculus, $@\lambda$, specified like many full-scale languages: as a bidirectionally typed translation semantics. Only the \rightarrow type constructor (tycon) is built in; all other external tycons (we show a variant of record types and constrained string types) are defined by extending a *tycon context*. Each tycon defines the semantics of associated term-level operators (e.g. record projection) using functions written in a static language where types and translations are values. The semantics provide strong metatheoretic guarantees, notably *type safety* and *conservativity*: that all *tycon invariants* will be conserved under extension. Mechanized proofs are not needed: problems are caught during typechecking by lifting typed compilation techniques into the semantics and enforcing abstraction barriers around tycons using type abstraction, the same principle that underlies ML-style modules.

1. Introduction

Typed programming languages are most often described as being composed of *fragments*, each contributing to the language’s concrete syntax, abstract syntax, static semantics and dynamic semantics. In his textbook, Harper organizes fragments around type constructors, describing each in a different chapter [12]. Languages are then identified by a set of type constructors, e.g. $\mathcal{L}\{\rightarrow \forall \mu 1 \times +\}$ is the language that includes partial function types, polymorphic types, recursive types, nullary and binary product types and binary sum types (its syntax is shown in Figure 1, discussed below).

Other researchers introduce fragments by defining a simple calculus with a “catch-all” constant and base type to stand notionally for all other terms and types that may also be included in some future complete language.

In contrast, the usual metatheoretic reasoning techniques for programming languages (e.g., rule induction) operate on complete language specifications. Each combination of fragments must formally be treated as its own monolithic language for which metatheorems must be established anew, guided only informally by those derived for the smaller systems from which the language is notionally composed.

This is not an everyday problem for programmers only because fragments like those mentioned above are “general purpose”: they make it possible to *isomorphically embed* many other fragments as “libraries”. For example, list types need not be built in because they are isomorphic to the type $\forall(\alpha.\mu(t.1 + (\alpha \times t)))$ (e.g. datatypes in ML, which combine these into a single declaration construct).

Nevertheless, situations do arise where it is not possible to use these fragments to establish an isomorphic embedding that preserves a desirable fragment’s static and dynamic semantics, including performance bounds specified by a cost semantics. Embeddings are also sometimes too *complex*, as measured by the cost of the extralinguistic computations that are needed to map in and out of the embedding and, if these must be performed mentally by programmers, considering various human factors. Each time a language is extended with a new fragment for one of these reasons, a new *dialect* is born. Within the ML lineage, for example, dialects abound:

1. **General Purpose Fragments:** A number of variations on product types, for example, have been introduced in dialects: n -ary tuples, labeled tuples, records (identified up to reordering), records with width and depth subtyping [6], records with update operators¹ [15], records with mutable fields [15], and records with “methods” (i.e. pure objects [22]). Sum-like types are also exposed in various ways: finite datatypes, open datatypes [16], hierarchically open datatypes [19], polymorphic variants [15] and ML-style exception types. Combinations of these manifest themselves as class-based object systems [15].
2. **Specialized Fragments:** Fragments that track specialized static invariants to provide stronger correctness

¹ The Haskell wiki notes that “No, extensible records are not implemented in GHC. The problem is that the record design space is large, and seems to lack local optima. [...] As a result, nothing much happens.” [1]

find a
really
good
example,
or add
this to
final
PSP
rev.??

or security guarantees, manage unwieldy lower-level abstractions and run-time systems or control cost are also frequently introduced in dialects, e.g. for data parallelism [7], distributed programming [20], reactive programming [17], authenticated data structures [18], databases [21], units of measure [14] and regular string sanitation [11].

3. **Foreign Fragments:** A safe and natural foreign function interface (FFI) can be valuable (particularly given this proliferation of dialects). This requires enforcing the type system of the foreign language in the calling language. For example, MLj builds in a safe FFI to Java [3].

This *dialect-oriented* state of affairs is unsatisfying. While programmers can choose from dialects supporting, e.g., a principled approach to distributed programming, or one that builds in support for statically reasoning about units of measure, there may not be an available dialect supporting both. Using different dialects separately for different components of a program is untenable: components written in different dialects cannot always interface safely (i.e. a safe FFI, item 3 above, is needed).

These problems do not arise when a fragment is expressed as an isomorphic embedding (i.e. as a library) because modern *module systems* enforce abstraction barriers that ensure that the isomorphism need only be established in the “closed world” of the module. There are no “link-time” proof obligations for clients in the “open world”. For example, a module defining the semantics of sets in ML can hold the representation of sets abstract, ensuring that any invariants maintained by the functions in the module (e.g. uniqueness, if using a list representation) will hold no matter which other modules are separately in use by a client.

When library-based embeddings are not possible, as in the examples above, mechanisms are needed that make it possible to define and reason in a similarly modular manner about direct extensions to the semantics of a language. Such a mechanism could ultimately be integrated directly into the language, blurring the distinction between fragments and libraries and decreasing the need for new dialects.

Contributions In this paper, we take foundational steps towards this goal by constructing a minimal but powerful core calculus, $@\lambda$ (the “actively typed” lambda calculus). Its semantics are structured like those of many modern languages, consisting of an *external language* (EL) governed by a typed translation semantics targeting a much simpler *internal language* (IL). Rather than building in a monolithic set of external type constructors, however, the semantics are indexed by a *tycon context*. Each tycon defines the semantics of its operators via functions written in a *static language* (SL) where types and translations are values.

We will begin by giving an overview of the organization and main judgements of $@\lambda$ in Sec. 2, then discuss how types are constructed and introduce our two main examples, one defining labeled product types with a functional update operator, and the other regular string types, based on a

recent core calculus style specification [11], in Sec. 3. We describe how tycons define their associated term-level operator constructors (opcons) in Sec. 4. We next give the key metatheoretic properties of the calculus in Sec. 5, including *type safety* and a key modularity result, which we call *conservativity*: any invariants that can be established about all values of a type under *some* tycon context (i.e. in some “closed world”) are conserved in any further extended tycon context (i.e. in the “open world”). Interestingly, type system providers need not necessarily provide mechanized proofs to maintain these guarantees. Instead, the approach we take relies on type abstraction in the internal language. As a result, we are able to borrow the same results that underly modular reasoning in simply-typed languages like ML to reason modularly about typed language fragments. We describe related work in Sec. 6 and conclude in Sec. 7.

2. Overview of $@\lambda$

External Language Programmers interface with $@\lambda$ by writing *external terms*, e . The syntax of external terms is shown in Figure 2. The static and dynamic semantics are given simultaneously as a *bidirectionally typed translation semantics*, i.e. the key judgements take the form:

$$\Upsilon \vdash_{\Phi} e \Rightarrow \sigma^+ \rightsquigarrow \iota^+ \quad \text{and} \quad \Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota^+$$

These are pronounced “ e (synthesizes / analyzes against) type σ and has translation ι under typing context Υ and tycon context Φ ”. Note that our specifications in this paper are intended to be algorithmic: we indicate “outputs” when introducing judgement forms by *mode annotations*, $^+$; these are not part of the judgement’s syntax. In particular, note that the type is an “output” in the synthetic judgement, but an “input” in the analytic judgement.

This basic separation of the EL and IL is commonly used for full-scale language specifications, e.g. the Harper-Stone semantics for Standard ML [13]. The internal language is purposely kept small, e.g. defining only simple products, to simplify metatheoretic reasoning and compilation. The EL then specifies various useful higher-level constructs, e.g. record types, by translation to the IL. In $@\lambda$, the EL builds in only function types. All other external constructs are defined in the *tycon context*, described in the Sec. 3.

We choose bidirectional typechecking, also sometimes called *local type inference* [23], for two main reasons. The first is once again to justify the practicality of our approach: local type inference is increasingly being used in modern languages (e.g. Scala) because it eliminates the need for type annotations in many situations while remaining decidable in more situations than whole-function type inference and providing what are widely perceived to be higher quality error messages. Secondly, it will give us a clean way to reuse the generalized introductory form, $\text{intro}[\sigma](\bar{e})$, and its associated desugarings, at many types, in a manner that relates to recent mechanisms supporting type-specific syntax

citation

citation
from
Dunfield

internal types

$$\tau ::= \tau \rightarrow \tau \mid \alpha \mid \forall(\alpha.\tau) \mid t \mid \mu(t.\tau) \mid 1 \mid \tau \times \tau \mid \tau + \tau$$

internal terms

$$\begin{aligned} \iota ::= & x \mid \lambda[\tau](x.\iota) \mid \iota(\iota) \mid \text{fix}[\tau](x.\iota) \mid \Lambda(\alpha.\iota) \mid \iota[\tau] \\ & \mid \text{fold}[t.\tau](\iota) \mid \text{unfold}(\iota) \mid () \mid (\iota, \iota) \mid \text{fst}(\iota) \mid \text{snd}(\iota) \\ & \mid \text{inl}[\tau](\iota) \mid \text{inr}[\tau](\iota) \mid \text{case}(\iota; x.\iota; x.\iota) \end{aligned}$$

internal typing contexts $\Gamma ::= \emptyset \mid \Gamma, x : \tau$

internal type formation contexts $\Delta ::= \emptyset \mid \Delta, \alpha \mid \Delta, t$

Figure 1. Syntax of $\mathcal{L}\{\rightarrow \forall \mu 1 \times +\}$, our internal language (IL). Metavariable x ranges over term variables and α and t both range over type variables.

external terms

$$\begin{aligned} e ::= & x \mid \lambda(x.e) \mid e(e) \mid \text{fix}(x.e) \mid e : \sigma \\ & \mid \text{intro}[\sigma](\bar{e}) \mid \text{targ}[\text{op}; \sigma](e; \bar{e}) \end{aligned}$$

argument lists $\bar{e} ::= \cdot \mid \bar{e}, e$

external typing contexts $\Upsilon ::= \emptyset \mid \Upsilon, x \Rightarrow \sigma$

Figure 2. Syntax of the external language (EL).

extensions [22]. For example, when we define regular string types, we will be able to reuse standard string literal syntax.

Unlike the Harper-Stone semantics, where external and internal terms were governed by a common type system, in $@\lambda$ each external type, σ , maps onto an internal type, τ , called the *type translation* of σ . This mapping is specified by the type translation judgement, $\vdash_{\Phi} \sigma \rightsquigarrow \tau$, which will be described in Sec. 3.4. For this reason, this specification style may also be compared to specifications for the first stage of a type-directed compiler, e.g. the TIL compiler for Standard ML [27], here lifted “one level up” into the semantics of the language itself. As we will see, type safety follows from a property analogous to a correctness condition that arises in typed compilers. Modular reasoning will be based on holding the type translation of σ abstract “outside” the tycon.

Internal Language $@\lambda$ relies on a typed internal language supporting type abstraction (i.e. universal quantification over types). We use $\mathcal{L}\{\rightarrow \forall \mu 1 \times +\}$, the syntax for which is shown in Figure 1, as representative of any intermediate language for a typed functional language.

We assume the statics of the IL are specified in the standard way by judgements for type formation $\Delta \vdash \tau$, typing context formation $\Delta \vdash \Gamma$ and type assignment $\Delta \Gamma \vdash \iota : \tau^+$. In examples, we will omit leading \emptyset , used as the base case for finite mappings, and \cdot , used as the base case for finite sequences (e.g. writing $\Gamma_{\text{test}} := x : \tau$).

The internal dynamics are specified also in a standard way as a structural operational semantics with a stepping judgement $\iota \mapsto \iota^+$ and a value judgement $\iota \text{ val}$. The multi-step judgement $\iota \mapsto^* \iota^+$ is the reflexive, transitive closure of the stepping judgement and the evaluation judgement $\iota \Downarrow \iota'$ is defined iff $\iota \mapsto^* \iota'$ and $\iota' \text{ val}$. Both the static and dynamic semantics of the IL can be found in any standard textbook covering typed lambda calculi (we directly follow [12]), so we assume familiarity and key metatheoretic properties.

kinds

$$\begin{aligned} \kappa ::= & \kappa \rightarrow \kappa \mid \alpha \mid \forall(\alpha.\kappa) \mid k \mid \mu_{\text{ind}}(k.\kappa) \mid 1 \mid \kappa \times \kappa \mid \kappa + \kappa \\ & \mid \text{Ty} \mid \text{ITy} \mid \text{ITm} \end{aligned}$$

static terms

$$\begin{aligned} \sigma ::= & x \mid \lambda x :: \kappa. \sigma \mid \sigma(\sigma) \mid \Lambda(\alpha.\sigma) \mid \sigma[\kappa] \\ & \mid \text{fold}[k.\kappa](\sigma) \mid \text{rec}[\kappa](\sigma; x.\sigma) \\ & \mid () \mid (\sigma, \sigma) \mid \text{fst}(\sigma) \mid \text{snd}(\sigma) \\ & \mid \text{inl}[\kappa](\sigma) \mid \text{inr}[\kappa](\sigma) \mid \text{case}(\sigma; x.\sigma; x.\sigma) \\ & \mid c\langle\sigma\rangle \mid \text{tycase}[c](\sigma; x.\sigma; \sigma) \\ & \mid \blacktriangleright(\hat{\tau}) \mid \triangleright(\hat{\iota}) \mid \text{ana}[n](\sigma) \mid \text{syn}[n] \mid \text{raise}[\kappa] \end{aligned}$$

translational internal types and terms

$$\hat{\tau} ::= \blacktriangleleft(\sigma) \mid \text{trans}(\sigma) \mid \hat{\tau} \rightarrow \hat{\tau} \mid \dots$$

$$\hat{\iota} ::= \triangleleft(\sigma) \mid \text{anatrans}[n](\sigma) \mid \text{syntrans}[n] \mid x \mid \lambda[\hat{\tau}](x.\hat{\iota}) \mid \dots$$

kinding contexts $\Gamma ::= \emptyset \mid \Gamma, x :: \kappa$

kind formation contexts $\Delta ::= \emptyset \mid \Delta, \alpha \mid \Delta, k$

Figure 3. Syntax of the static language (SL). Metavariable x ranges over static term variables, α and k over kind variables and m and n over natural numbers.

Static Language The workhorse of $@\lambda$ is the *static language*, which itself forms a typed lambda calculus where *kinds*, κ , classify *static terms*, σ . The syntax of the SL is given in Figure 3. The portion of the SL covered by the first row of kinds and first four rows of static terms in the syntax forms an entirely standard functional programming language consisting of total functions, universal quantification over kinds, inductive kinds (constrained by the standard positivity condition to prevent non-termination), and products and sums. The reader can consider these as forming a total subset of ML or a simply-typed subset of Coq. The semantics also directly follow [12], so we omit the details here. Only three new kinds will be needed: Ty (Sec. 3), ITy (Sec. 3.4) and ITm (Sec. 4.1).

The kinding judgement takes the form $\Delta \Gamma \vdash_{\Phi}^n \sigma :: \kappa^+$, where Δ and Γ are analogous to Δ and Γ and analogous kind and kinding context formation judgements $\Delta \vdash \kappa$ and $\Delta \vdash \Gamma$ are defined. The natural number n is used as a technical device in our semantics to ensure that the forms shown as being indexed by n in the syntax only arise in a controlled manner internally to prevent “out of bounds” issues, as we will discuss; they would have no corresponding concrete syntax so n can be assumed 0 in user-defined terms.

The dynamic semantics of static terms is defined as a structural operational semantics by the stepping judgement $\sigma \mapsto_{\mathcal{A}} \sigma^+$, the value judgement $\sigma \text{ val}_{\mathcal{A}}$ and the error judgement $\sigma \text{ err}_{\mathcal{A}}$. Here, \mathcal{A} ranges over *argument environments*, which we will return to when considering opcons in Sec. 4. The multi-step judgement $\sigma \mapsto_{\mathcal{A}}^* \sigma^+$ is the reflexive, transitive closure of the stepping judgement. The normalization judgement $\sigma \Downarrow_{\mathcal{A}} \sigma'$ is defined iff $\sigma \mapsto_{\mathcal{A}}^* \sigma'$ and $\sigma' \text{ val}_{\mathcal{A}}$.

3. Types

External types, or simply *types*, are static values of kind Ty. The introductory form for kind Ty is $c\langle\sigma\rangle$, where c is a *tycon* and σ is the *type index*. The syntax for tycons in

$$\begin{array}{c}
\text{(k-parr)} \\
\frac{\Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{Ty} \times \text{Ty}}{\Delta \Gamma \vdash_{\Phi}^n \rightarrow \langle \sigma \rangle :: \text{Ty}}
\end{array}
\quad
\begin{array}{c}
\text{(k-ty)} \\
\frac{\text{tycon TC } \{\theta\} \sim \text{tcsig}[\kappa_{\text{tyidx}}][\chi] \in \Phi \quad \Delta \Gamma \vdash_{\Phi}^n \sigma :: \kappa_{\text{tyidx}}}{\Delta \Gamma \vdash_{\Phi}^n \text{TC} \langle \sigma \rangle :: \text{Ty}}
\end{array}
\quad
\begin{array}{c}
\text{(k-otherty)} \\
\frac{\Delta \Gamma \vdash_{\Phi}^n \sigma_{\text{tyidx}} :: \text{Nat} \times \text{ITy}}{\Delta \Gamma \vdash_{\Phi}^n \text{other}[m] \langle \sigma_{\text{tyidx}} \rangle :: \text{Ty}}
\end{array}$$

Figure 5. Kinding rules for types, which take the form $c\langle\sigma_{\text{tyidx}}\rangle$ where c is a tycon and σ_{tyidx} is the type index.

Figure 4. Syntax of tycons. Metavariables TC and op range over user-defined tycon and opcon names, respectively, and m ranges over natural numbers.

The rule (k-parr) specifies that the type index of partial function types must be a pair of types. We thus say that \rightarrow has *index kind* $\text{Ty} \times \text{Ty}$. To recover a conventional syntax, we can introduce a desugaring from $\sigma_1 \rightarrow \sigma_2$ to $\rightarrow((\sigma_1, \sigma_2))$.

All user-defined tycons must be defined in the *tycon context*, Φ , which is simply a list of tycon definitions. Each tycon defines a name, TC, a *tycon structure*, θ , and a *tycon signature*, ψ . We will return to the tycon structure below. Tycon signatures have the form $\text{tcsig}^{\kappa_{\text{tyidx}}} \{\chi\}$, where κ_{tyidx} is the tycon’s index kind and χ is the *opcon signature*, which we discuss in Sec. 4. The first premise of (k-ty) extracts the index kind and the second checks the type index against it.

The rule (k-otherty) governs types constructed by an “other” tycon. These will serve only as technical devices to stand in for types other than those in a given tycon context in Sec. 5. The index of such a type must pair a natural number with a type translation of kind ITy , discussed in Sec. 3.4.

Examples As our first example, consider the user-defined tycon RSTR. The index kind of RSTR is Rx, which classifies static regular expressions and is defined as an inductive sum kind in the usual way. Types constructed by RSTR will classify *regular strings*, which are statically known to be in the regular language specified by the type index [11]. For example, $\sigma_{\text{title}} := \text{RSTR} \langle / . + / \rangle$ classifies non-empty strings and $\sigma_{\text{conf}} := \text{RSTR} \langle / [A-Z]^+ \backslash d \backslash d \backslash d / \rangle$ classifies conference names. The type indices are here written using standard concrete syntax for concision. Previous work has shown how to define type-specific (here, kind-specific) syntax like this composably in libraries [22]. We define

$$\begin{array}{c}
(\text{tcc-ext}) \quad \vdash \Phi \quad \text{TC} \notin \text{dom}(\Phi) \quad \boxed{\vdash \Phi} \\
\emptyset \vdash \kappa_{\text{tyidx}} \text{ eq} \quad \emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{schema}} :: \kappa_{\text{tyidx}} \rightarrow \text{ITy} \\
\frac{\vdash \Phi, \text{tycon TC} \{ \text{trans} = \sigma_{\text{schema}} \text{ in } \omega \} \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{ \chi \} \quad \omega \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{ \chi \}}{\vdash \Phi, \text{tycon TC} \{ \text{trans} = \sigma_{\text{schema}} \text{ in } \omega \} \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{ \chi \}}
\end{array}$$

a tycon context containing only the definition of RSTR,

$$\Phi_{\text{rstr}} := \text{tycon RSTR } \{\theta_{\text{rstr}}\} \sim \text{tcsig}[\text{Rx}] \{\chi_{\text{rstr}}\}.$$

Our second example is the tycon LPROD, which will define a variant of labeled product type (labeled products are like record types, but maintain a row ordering; record types are also definable in a manner discussed in the supplement, but maintaining an ordering simplifies our discussion). We choose the index kind of LPROD to be $\text{List}[\text{Lbl} \times \text{Ty}]$, where list kinds are defined as inductive sums in the usual way, and Lbl classifies static representations of row labels. The tycon context containing only LPROD’s definition is $\Phi_{\text{lprod}} := \text{tycon LPROD} \{ \theta_{\text{lprod}} \} \sim \text{tcsig}[\text{List}[\text{Lbl} \times \text{Ty}]] \{ \chi_{\text{lprod}} \}$.

In the tycon context containing both tycon definitions, $\Phi_{\text{rstr}}\Phi_{\text{lprod}}$, we can define a labeled product type classifying conference papers, $\sigma_{\text{paper}} := \text{LPROD}\langle\{\text{title} : \sigma_{\text{title}}, \text{conf} : \sigma_{\text{conf}}\}\rangle$. Note that σ_{paper} **type** $_{\Phi_{\text{rstr}}\Phi_{\text{lprod}}}$ and we again use kind-specific syntax, in this case for Lbl and $\text{List}[\text{Lbl} \times \text{Ty}]$.

3.1 Type Case Analysis

$$\frac{\begin{array}{c} \text{(k-tycase)} \\ \Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{Ty} \quad \text{tycon TC } \{\theta\} \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{\chi\} \in \Phi \\ \Delta \Gamma, \mathbf{x} :: \kappa_{\text{tyidx}} \vdash_{\Phi}^n \sigma_1 :: \kappa \quad \Delta \Gamma \vdash_{\Phi}^n \sigma_2 :: \kappa \end{array}}{\Delta \Gamma \vdash_{\Phi}^n \text{tycase}[\text{TC}](\sigma; \mathbf{x}.\sigma_1; \sigma_2) :: \kappa}$$

Put another way, types can be thought of as arising from a distinguished “open datatype” defined by the tycon context.

3.2 Tycon Context Well-Definedness

3.3 Type Equivalence

The first check simplifies the handling of type equivalence: type index kinds must be *equality kinds*, i.e. those for which semantic equivalence implies syntactic equality at normal form. We define these by the judgement $\Delta \vdash \kappa \text{ eq}$ (see supplement). Equality kinds are similar to equality types as citation found in Standard ML. The main implication of this choice is that type indices cannot contain static functions.

3.4 Type Translations

Each tycon computes translations for the types it constructs as a function of each type's index by specifying a *translation schema* as the first component of the tycon structure, θ . For a tycon with index kind κ_{tyidx} , the translation schema must have kind $\kappa_{\text{tyidx}} \rightarrow \text{ITy}$, checked by (tcc-ext).

The kind ITy has a single introductory form, $\blacktriangleright(\hat{\tau})$, where $\hat{\tau}$ is a *translational internal type*. Each form in the syntax for internal types, τ , corresponds to a form in the syntax of translational internal types, $\hat{\tau}$. For example, our translation schema for RSTR simply chooses to ignore the type index and represent all regular strings internally as strings, of internal type abbreviated str . We abbreviate the corresponding translational internal type $\hat{\text{str}}$ and define the translation schema as $\theta_{\text{rstr}} := \text{trans} = \lambda \text{tyidx}::\text{Rx}.\blacktriangleright(\hat{\text{str}})$ in ω_{rstr} .

The kinding rules and operational semantics for these shared forms simply proceed recursively, e.g.

$$\frac{(\text{k-ity-prod}) \quad \Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\hat{\tau}_1) :: \text{ITy} \quad \Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\hat{\tau}_2) :: \text{ITy}}{\Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) :: \text{ITy}}$$

The syntax for translational internal types additionally includes an “unquote” form, $\blacktriangleleft(\sigma)$, so that they can be constructed compositionally, as well as a form, $\text{trans}(\sigma)$, that allows one type's translation to refer to the translation of another type σ .

$$\begin{array}{c} (\text{k-ity-unquote}) \\ \frac{\Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{ITy}}{\Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\blacktriangleleft(\sigma)) :: \text{ITy}} \end{array} \quad \begin{array}{c} (\text{k-ity-trans}) \\ \frac{\Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{ITy}}{\Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\text{trans}(\sigma)) :: \text{ITy}} \end{array}$$

The unquote form is eliminated during normalization, while references to the translation of a type are retained in values of kind ITy . The key rules in the dynamics are:

$$\begin{array}{c} (\text{s-ity-unquote-elim}) \\ \frac{\blacktriangleright(\hat{\tau}) \text{ val}_{\mathcal{A}}}{\blacktriangleright(\blacktriangleleft(\blacktriangleright(\hat{\tau}))) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau})} \end{array} \quad \begin{array}{c} (\text{s-ity-trans-val}) \\ \frac{\sigma \text{ val}_{\mathcal{A}}}{\blacktriangleright(\text{trans}(\sigma)) \text{ val}_{\mathcal{A}}} \end{array}$$

We choose a translation schema for LPROD that generates nested binary product types by folding over the type index and referring to the translations of the types therein (though this is not the only workable choice, e.g. we could also have used a list). We assume $\text{listfold} :: \forall(\alpha_1.\forall(\alpha_2.\text{List}[\alpha_1] \rightarrow \alpha_2 \rightarrow (\alpha_1 \rightarrow \alpha_2 \rightarrow \alpha_2) \rightarrow \alpha_2))$ in defining $\theta_{\text{lprod}} := \text{trans} = \sigma_{\text{lprod/trans}} \cdot \omega_{\text{lprod}}$ where $\sigma_{\text{lprod/trans}} :=$

$$\lambda \text{tyidx}::\text{List}[\text{Lbl} \times \text{Ty}].\text{listfold}[\text{Lbl} \times \text{Ty}] [\text{ITy}] \text{tyidx} \blacktriangleright(1) \\ (\lambda h:\text{Lbl} \times \text{Ty}.\lambda r:\text{ITy}.\blacktriangleright(\text{trans}(\text{snd}(h)) \times \blacktriangleleft(r)))$$

Evaluating this translation schema with the index of σ_{paper} , for example, produces the value $\sigma_{\text{paper/trans}} := \blacktriangleright(\hat{\tau}_{\text{paper/trans}})$ where $\hat{\tau}_{\text{paper/trans}} := \text{trans}(\sigma_{\text{title}}) \times (\text{trans}(\sigma_{\text{conf}}) \times 1)$. Note that we do not include logic to optimize away the trailing unit type for simplicity (and to again emphasize that many translations are possible for any given type).

3.4.1 Selective Type Translation Abstraction

References to translations of other types are maintained in values of kind ITy like this to allow us to selectively hold them abstract, which will be the key to our main results. This can be thought of as analogous to the process in ML by which the true identity of an abstract type in a module is held abstract outside the module until after typechecking. The judgement $\hat{\tau} \parallel \mathcal{D} \mapsto_{\Phi}^c \tau^+ \parallel \mathcal{D}^+$ relates a normalized translational internal type $\hat{\tau}$ to an internal type τ , called a corresponding *abstracted type translation* because references to translations of types constructed by a tycon other than those constructed by a “delegated tycon”, c , are replaced by universally quantified type variables, α . For example, $\hat{\tau}_{\text{paper/trans}} \parallel \emptyset \mapsto_{\Phi_{\text{rstr}} \Phi_{\text{lprod}}}^{\text{LPROD}} \tau_{\text{paper/abs}} \parallel \mathcal{D}_{\text{paper/abs}}$ where $\tau_{\text{paper/abs}} := \alpha_1 \times (\alpha_2 \times 1)$ and $\mathcal{D}_{\text{paper/abs}} := \sigma_{\text{title}} \leftrightarrow \text{str}/\alpha_1, \sigma_{\text{conf}} \leftrightarrow \text{str}/\alpha_2$. The type translation store $\mathcal{D} ::= \emptyset \mid \mathcal{D}, \sigma \leftrightarrow \tau/\alpha$ maintains the mapping from types, σ , to their (non-abstract) translations, τ , and the type variables, α , which appear in their place.

Each type translation store induces a *type substitution*, δ , and a corresponding internal type formation context, Δ , according to the judgement $\mathcal{D} \rightsquigarrow \delta : \Delta$. Type substitutions simply define an n -ary substitution for type variables, $\delta ::= \emptyset \mid \delta, \tau/\alpha$. For example, $\mathcal{D}_{\text{paper/abs}} \rightsquigarrow \delta_{\text{paper/abs}} : \Delta_{\text{paper/abs}}$ where $\delta_{\text{paper/abs}} := \text{str}/\alpha_1, \text{str}/\alpha_2$ and $\Delta_{\text{paper/abs}} := \alpha_1, \alpha_2$. We can apply type substitutions to internal types, terms and typing contexts, written $[\delta]\tau$, $[\delta]\iota$ and $[\delta]\Gamma$, respectively. For example, $[\delta_{\text{paper/abs}}]\tau_{\text{paper/abs}}$ is $\tau_{\text{paper}} := \text{str} \times (\text{str} \times 1)$, i.e. the final type translation of σ_{paper} . Indeed, we can now give the rule defining the type translation judgement, $\vdash_{\Phi} \sigma \rightsquigarrow \tau$, mentioned in Sec. 2. We simply determine any selectively abstract translation, then apply the substitution:

$$\frac{(\text{conc-ty-trans}) \quad \sigma \text{ type}_{\Phi} \quad \text{trans}(\sigma) \parallel \emptyset \mapsto_{\Phi}^{\text{TC}} \tau \parallel \mathcal{D} \quad \mathcal{D} \rightsquigarrow \delta : \Delta}{\vdash_{\Phi} \sigma \rightsquigarrow [\delta]\tau}$$

With this intuition, we can now give the rules for the selective type abstraction judgement. We recurse generically over sub-terms of $\hat{\tau}$ until sub-terms of form $\text{trans}(\sigma)$ are encountered (see supplement).

The translation of partial function types is direct and is not held abstract, so that lambdas can be used as the sole binding construct in the external language:

$$\frac{(\text{abs-parr}) \quad \text{trans}(\sigma_1) \parallel \mathcal{D} \mapsto_{\Phi}^c \tau_1 \parallel \mathcal{D}' \quad \text{trans}(\sigma_2) \parallel \mathcal{D}' \mapsto_{\Phi}^c \tau_2 \parallel \mathcal{D}''}{\text{trans}(\rightarrow((\sigma_1, \sigma_2))) \parallel \mathcal{D} \mapsto_{\Phi}^c \tau_1 \rightarrow \tau_2 \parallel \mathcal{D}''}$$

The translation of a user-defined types constructed by the delegated tycon is determined by calling the translation schema and checking that the type translation it generates refers only to type variables generated from \mathcal{D}' :

$$\begin{array}{c}
 \text{(abs-tc-local)} \\
 \text{tycon TC } \{\text{trans} = \sigma_{\text{schema}} \text{ in } \omega\} \sim \psi \in \Phi \\
 \sigma_{\text{schema}}(\sigma_{\text{tyidx}}) \Downarrow \blacktriangleright(\hat{\tau}) \quad \hat{\tau} \parallel \mathcal{D} \xrightarrow{\text{TC}} \tau \parallel \mathcal{D}' \\
 \mathcal{D}' \sim \delta : \Delta \quad \Delta \vdash \tau \\
 \hline
 \text{trans}(\text{TC}(\sigma_{\text{tyidx}})) \parallel \mathcal{D} \xrightarrow{\text{TC}} \tau \parallel \mathcal{D}'
 \end{array}$$

The translation of a user-defined type constructed by any tycon other than the delegated tycon is added to the store (the supplement has the simple rule for retrieving it once there):

$$\begin{array}{c}
 \text{(abs-tc-foreign-new)} \\
 c \neq \text{TC} \quad \text{TC}(\sigma_{\text{tyidx}}) \notin \text{dom}(\mathcal{D}) \\
 \text{tycon TC } \{\text{trans} = \sigma_{\text{schema}} \text{ in } \omega\} \sim \psi \in \Phi \\
 \sigma_{\text{schema}}(\sigma_{\text{tyidx}}) \Downarrow \blacktriangleright(\hat{\tau}) \quad \hat{\tau} \parallel \mathcal{D} \xrightarrow{\text{TC}} \tau \parallel \mathcal{D}' \\
 \mathcal{D}' \sim \delta : \Delta \quad \Delta \vdash \tau \quad (\alpha \text{ fresh}) \\
 \hline
 \text{trans}(\text{TC}(\sigma_{\text{tyidx}})) \parallel \mathcal{D} \xrightarrow{\text{TC}} \alpha \parallel \mathcal{D}', \text{TC}(\sigma_{\text{tyidx}}) \leftrightarrow \tau/\alpha
 \end{array}$$

The translation of an “other” type is given directly in its index (the rule where it is held abstract is in the supplement):

$$\begin{array}{c}
 \text{(abs-tc-other-new)} \\
 \hat{\tau} \parallel \mathcal{D} \xrightarrow{\text{other}[m]} \tau \parallel \mathcal{D}' \\
 \hline
 \text{trans}(\text{other}[m](\langle \sigma_{\text{nat}}, \blacktriangleright(\hat{\tau}) \rangle)) \parallel \mathcal{D} \xrightarrow{\text{other}[m]} \tau \parallel \mathcal{D}'
 \end{array}$$

3.5 Typing Contexts

We must also define a translation judgement for external typing contexts $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ that checks that Υ contains valid types, then generates the corresponding internal typing context Γ from their translations (see supplement).

4. Typing and Translation of External Terms

Having established how types are constructed, and how they determine abstract and from there concrete translations to internal types, we can finally give the typing and translation rules for external terms, shown in Figure 7.

Because we are defining a bidirectional type system, a subsumption rule is needed to allow synthetic terms to be analyzed against an equivalent type. Per above, equivalent types must be syntactically identical at normal form, and we consider analysis only if $\sigma \text{ type}_{\Phi}$, so the rule (subsume) is straightforward. To use an analytic term in a synthetic position, the programmer must provide a type ascription, written $e : \sigma$. The ascription is kind checked and normalized to a type before being used for analysis by rule (ascribe).

Variables and functions behave in the standard manner given our definitions of types and type translations (used to generate the ascription on the lambda in the IL, which is intrinsically rather than bidirectionally typed). We use Plotkin’s fixpoint operator for general recursion (cf. [12]), and define both only in analytic positions for simplicity.

4.1 Generalized Introductory Operations

The translation of the generalized introductory form, $\text{intro}[\sigma_{\text{tmidx}}](\bar{e})$, is determined by the tycon of the type it is being analyzed against as a function of the type’s index, the *term index*, σ_{tmidx} , and the *argument list*, \bar{e} .

Before discussing rules (ana-intro) and (ana-intro-other), we note that we can recover a variety of standard concrete introductory forms by a purely syntactic desugaring to this abstract form (and thus allow their use at more than one type). For example, for regular strings we will be able to use the string literal form, “s”, which desugars to $\text{intro}[\text{"s"}_{\text{SL}}](\cdot)$. The term index is the corresponding static value of kind Str, indicated by a subscript for clarity. Similarly, for labeled products, records, objects and so on, we can define a generalized labeled collection form, $\{\text{lbl}_1 = e_1, \dots, \text{lbl}_n = e_n\}$, that desugars to $\text{intro}[\text{[lbl}_1, \dots, \text{lbl}_n\text{]}](e_1; \dots; e_n)$, i.e. a list constructed from the row labels is the term index and the corresponding row values are the arguments. In both cases, the term index captures static portions of the concrete form and the arguments capture all external sub-terms. Additional desugarings are shown in the supplement and a technique based on [22] could be introduced to allow tycon providers to define more such desugarings compositably.

Let us derive $\Upsilon_{\text{ex}} \vdash_{\Phi_{\text{rstr}} \Phi_{\text{lprod}}} e_{\text{ex}} \Leftarrow \sigma_{\text{paper}} \rightsquigarrow \iota_{\text{ex}}$ where we define an example typing context $\Upsilon_{\text{ex}} := \text{title} \Rightarrow \sigma_{\text{title}}$ and $e_{\text{ex}} := \{\text{title} = \text{title}, \text{conf} = \text{"EXMPL 2015"}\}$. The translation will be $\iota_{\text{ex}} := (\text{title}, (\text{"EXMPL 2015"}_{\text{IL}}, ()))$, where $\text{"EXMPL 2015"}_{\text{IL}}$ is an internal string (of internal type str).

The first premise of (ana-intro) extracts the tycon definition for the tycon of the type the intro form is being analyzed against. In this example, this is LPROD. We will use this as the *delegated tycon* in the final premises of the rule.

The second premise extracts the *intro term index kind*, κ_{tmidx} , from the *opcon signature*, χ , of the tycon signature and the third premise checks the provided term index against this kind. This is simply the kind of term index expected by the tycon, e.g., LPROD specifies $\text{List}[\text{Lbl}]$, so that it can use the labeled collection form, while RSTR specifies an intro index kind of Str, so that it can use the string literal form.

The fourth premise extracts the *intro opcon definition* from the *opcon structure*, ω , of the tycon structure, calling it σ_{def} . This is a static function that is applied, in the seventh premise, to determine whether the term is well-typed, raising an error if not or determining the translation of the term if so. The function has access to the type index, the term index and an interface to the list of arguments, and its kind is checked by the judgement $\vdash_{\Phi} \omega \sim \psi$, which appeared as the final premise of the rule (tcc-ext) and is defined in Figure 8.

For example, the opcon structure of RSTR is $\omega_{\text{rstr}} := \text{ana intro} = \sigma_{\text{rstr/intro}}; \omega_{\text{rstr/targops}}$ where $\sigma_{\text{rstr/intro}} :=$

```

 $\lambda \text{tyidx} :: \text{Rx} . \lambda \text{tmidx} :: \text{Str} . \lambda \text{args} :: \text{List}[\text{Arg}] .$ 
  let  $\text{aok} :: 1 = \text{arity0 args}$  in
  let  $\text{rok} :: 1 = \text{rmatch tyidx tmidx in str\_of\_Str tmidx}$ 

```

$$\boxed{\Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota} \quad \boxed{\Upsilon \vdash_{\Phi} e \Rightarrow \sigma \rightsquigarrow \iota}$$

$$\begin{array}{c}
\text{(subsume)} \\
\frac{\Upsilon \vdash_{\Phi} e \Rightarrow \sigma \rightsquigarrow \iota}{\Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota}
\end{array}
\quad
\begin{array}{c}
\text{(ascribe)} \\
\frac{\emptyset \emptyset \vdash_{\Phi}^0 \sigma :: \text{Ty} \quad \sigma \Downarrow \sigma'}{\Upsilon \vdash_{\Phi} e \Leftarrow \sigma' \rightsquigarrow \iota}
\end{array}
\quad
\begin{array}{c}
\text{(syn-var)} \\
\frac{x \Rightarrow \sigma \in \Upsilon}{\Upsilon \vdash_{\Phi} x \Rightarrow \sigma \rightsquigarrow x}
\end{array}
\quad
\begin{array}{c}
\text{(ana-fix)} \\
\frac{\Upsilon, x \Rightarrow \sigma \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota \quad \vdash_{\Phi} \sigma \rightsquigarrow \tau}{\Upsilon \vdash_{\Phi} \text{fix}(x.e) \Leftarrow \sigma \rightsquigarrow \text{fix}[\tau](x.\iota)}
\end{array}$$

$$\begin{array}{c}
\text{(ana-lam)} \\
\frac{\Upsilon, x \Rightarrow \sigma_1 \vdash_{\Phi} e \Leftarrow \sigma_2 \rightsquigarrow \iota \quad \vdash_{\Phi} \sigma_1 \rightsquigarrow \tau_1}{\Upsilon \vdash_{\Phi} \lambda(x.e) \Leftarrow \rightarrow\langle(\sigma_1, \sigma_2)\rangle \rightsquigarrow \lambda[\tau_1](x.\iota)}
\end{array}
\quad
\begin{array}{c}
\text{(syn-ap)} \\
\frac{\Upsilon \vdash_{\Phi} e_1 \Rightarrow \rightarrow\langle(\sigma_1, \sigma_2)\rangle \rightsquigarrow \iota_1 \quad \Upsilon \vdash_{\Phi} e_2 \Leftarrow \sigma_2 \rightsquigarrow \iota_2}{\Upsilon \vdash_{\Phi} e_1(e_2) \Rightarrow \sigma_2 \rightsquigarrow \iota_1(\iota_2)}
\end{array}$$

$$\begin{array}{c}
\text{(ana-intro)} \\
\text{tycon TC } \{\text{trans} = _ \text{ in } \omega\} \sim \text{tcsig}[_] \{\chi\} \in \Phi \\
\text{intro}[\kappa_{\text{tmidx}}] \in \chi \quad \emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{tmidx}} :: \kappa_{\text{tmidx}} \\
\text{ana intro} = \sigma_{\text{def}} \in \omega \quad |\bar{e}| = n \quad \text{args}(n) = \sigma_{\text{args}} \\
\sigma_{\text{def}} \sigma_{\text{tyidx}} \sigma_{\text{tmidx}} \sigma_{\text{args}} \Downarrow \bar{e}; \Upsilon; \Phi \triangleright (\hat{\iota}) \\
\hat{\iota} \parallel \emptyset \emptyset \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} \tau_{\text{abs}} \parallel \mathcal{D} \mathcal{G} \quad \mathcal{G} \rightsquigarrow \gamma : \Gamma_{\text{abs}} \\
\text{trans}(\text{TC}(\sigma_{\text{tyidx}})) \parallel \mathcal{D} \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} \tau_{\text{abs}} \parallel \mathcal{D}' \quad \mathcal{D}' \rightsquigarrow \delta : \Delta_{\text{abs}} \\
\Delta_{\text{abs}} \Gamma_{\text{abs}} \vdash \iota_{\text{abs}} : \tau_{\text{abs}} \\
\hline
\Upsilon \vdash_{\Phi} \text{intro}[\sigma_{\text{tmidx}}](\bar{e}) \Leftarrow \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow [\delta][\gamma]\iota
\end{array}
\quad
\begin{array}{c}
\text{(syn-targ)} \\
\Upsilon \vdash_{\Phi} e_{\text{targ}} \Rightarrow \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow \iota_{\text{targ}} \\
\text{tycon TC } \{\text{trans} = _ \text{ in } \omega\} \sim \text{tcsig}[_] \{\chi\} \in \Phi \\
\text{op}[\kappa_{\text{tmidx}}] \in \chi \quad \emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{tmidx}} :: \kappa_{\text{tmidx}} \\
\text{syn op} = \sigma_{\text{def}} \in \omega \quad |e_{\text{targ}}; \bar{e}| = n \quad \text{args}(n) = \sigma_{\text{args}} \\
\sigma_{\text{def}} \sigma_{\text{tyidx}} \sigma_{\text{tmidx}} \sigma_{\text{args}} \Downarrow (e_{\text{targ}}; \bar{e}); \Upsilon; \Phi (\sigma, \triangleright(\hat{\iota})) \\
\hat{\iota} \parallel \emptyset \emptyset \Downarrow_{(e_{\text{targ}}; \bar{e}); \Upsilon; \Phi}^{\text{TC}} \tau_{\text{abs}} \parallel \mathcal{D} \mathcal{G} \quad \mathcal{G} \rightsquigarrow \gamma : \Gamma_{\text{abs}} \\
\text{trans}(\sigma) \parallel \mathcal{D} \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} \tau_{\text{abs}} \parallel \mathcal{D}' \quad \mathcal{D}' \rightsquigarrow \delta : \Delta_{\text{abs}} \\
\Delta_{\text{abs}} \Gamma_{\text{abs}} \vdash \iota_{\text{abs}} : \tau_{\text{abs}} \\
\hline
\Upsilon \vdash_{\Phi} \text{targ}[\text{op}; \sigma_{\text{tmidx}}](e_{\text{targ}}; \bar{e}) \Rightarrow \sigma \rightsquigarrow [\delta][\gamma]\iota_{\text{abs}}
\end{array}$$

$$\begin{array}{c}
\text{(ana-intro-other)} \\
\emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{def}} :: \text{List}[\text{Arg}] \rightarrow \text{ITm} \\
|\bar{e}| = n \quad \text{args}(n) = \sigma_{\text{args}} \quad \sigma_{\text{def}} \sigma_{\text{args}} \Downarrow \bar{e}; \Upsilon; \Phi \triangleright (\hat{\iota}) \\
\hat{\iota} \parallel \emptyset \emptyset \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{other}[m]} \tau_{\text{abs}} \parallel \mathcal{D} \mathcal{G} \\
\text{trans}(\text{other}[m](\sigma_{\text{tyidx}})) \parallel \mathcal{D} \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{other}[m]} \tau_{\text{abs}} \parallel \mathcal{D}' \\
\mathcal{D}' \rightsquigarrow \delta : \Delta_{\text{abs}} \quad \mathcal{G} \rightsquigarrow \gamma : \Gamma_{\text{abs}} \quad \Delta_{\text{abs}} \Gamma_{\text{abs}} \vdash \iota_{\text{abs}} : \tau_{\text{abs}} \\
\hline
\Upsilon \vdash_{\Phi} \text{intro}[\sigma_{\text{def}}](\bar{e}) \Leftarrow \text{other}[m](\sigma_{\text{tyidx}}) \rightsquigarrow [\delta][\gamma]\iota_{\text{abs}}
\end{array}
\quad
\begin{array}{c}
\text{(syn-targ-other)} \\
\Upsilon \vdash_{\Phi} e_{\text{targ}} \Rightarrow \text{other}[m](\sigma_{\text{tyidx}}) \rightsquigarrow \iota_{\text{targ}} \\
\emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{def}} :: \text{List}[\text{Arg}] \rightarrow (\text{Ty} \times \text{ITm}) \\
|e_{\text{targ}}; \bar{e}| = n \quad \text{args}(n) = \sigma_{\text{args}} \quad \sigma_{\text{def}} \sigma_{\text{args}} \Downarrow (e_{\text{targ}}; \bar{e}); \Upsilon; \Phi (\sigma, \triangleright(\hat{\iota})) \\
\hat{\iota} \parallel \emptyset \emptyset \Downarrow_{(e_{\text{targ}}; \bar{e}); \Upsilon; \Phi}^{\text{other}[m]} \tau_{\text{abs}} \parallel \mathcal{D} \mathcal{G} \\
\text{trans}(\sigma) \parallel \mathcal{D} \Downarrow_{\bar{e}; \Upsilon; \Phi}^{\text{other}[m]} \tau_{\text{abs}} \parallel \mathcal{D}' \\
\mathcal{D}' \rightsquigarrow \delta : \Delta_{\text{abs}} \quad \mathcal{G} \rightsquigarrow \gamma : \Gamma_{\text{abs}} \quad \Delta_{\text{abs}} \Gamma_{\text{abs}} \vdash \iota_{\text{abs}} : \tau_{\text{abs}} \\
\hline
\Upsilon \vdash_{\Phi} \text{targ}[\text{op}; \sigma_{\text{def}}](e_{\text{targ}}; \bar{e}) \Rightarrow \sigma \rightsquigarrow [\delta][\gamma]\iota_{\text{abs}}
\end{array}$$

Figure 7. Typing

$$\boxed{\vdash_{\Phi} \omega \sim \psi}$$

$$\begin{array}{c}
\text{(ocstruct-intro)} \\
\text{intro}[\kappa_{\text{tmidx}}] \in \chi \quad \emptyset \vdash \kappa_{\text{tmidx}} \\
\emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{def}} :: \kappa_{\text{tyidx}} \rightarrow \kappa_{\text{tmidx}} \rightarrow \text{List}[\text{Arg}] \rightarrow \text{ITm} \\
\hline
\vdash_{\Phi} \text{ana intro} = \sigma_{\text{def}} \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{\chi\}
\end{array}
\quad
\begin{array}{c}
\text{(ocstruct-targ)} \\
\vdash_{\Phi} \omega \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{\chi\} \quad \text{op} \notin \text{dom}(\chi) \quad \emptyset \vdash \kappa_{\text{tmidx}} \\
\emptyset \emptyset \vdash_{\Phi}^0 \sigma_{\text{def}} :: \kappa_{\text{tyidx}} \rightarrow \kappa_{\text{tmidx}} \rightarrow \text{List}[\text{Arg}] \rightarrow (\text{Ty} \times \text{ITm}) \\
\hline
\vdash_{\Phi} \omega; \text{syn op} = \sigma_{\text{def}} \sim \text{tcsig}[\kappa_{\text{tyidx}}] \{\chi, \text{op}[\kappa_{\text{tmidx}}]\}
\end{array}$$

Figure 8. Checking opcon structures against tycon signatures.

Because regular strings are implemented as strings, this intro opcon definition is straightforward. It begins by making sure that no arguments were passed in (we will return to arguments and the fifth and sixth premises of (ana-intro) with the next example), using the helper function $\text{arity0} :: \text{List}[\text{Arg}] \rightarrow 1$ defined such that any non-empty list will raise an error, via the static term $\text{raise}[1]$. In practice, the tycon provider would specify an error message here.

Next, it checks the string provided as the term index against the regular expression given as the type index using $\text{rmatch} :: \text{Rx} \rightarrow \text{Str} \rightarrow 1$, which we assume is defined in the usual way and again raises an error on failure. Finally, a translation corresponding to the term index is generated via helper function $\text{str_of_Str} :: \text{Str} \rightarrow \text{ITm}$.

Term Translations The only introductory form for kind ITm is $\triangleright(\hat{\iota})$, where $\hat{\iota}$ is a *translational internal term*. This form is analogous to the introductory form for kind ITy described in Sec. 3.4, $\blacktriangleright(\hat{\tau})$. Each form in the syntax of ι has a corresponding form in the syntax for $\hat{\iota}$ and both the kinding rules and operational semantics simply recurse through these in the same manner as shown in Sec. 3.4. Also analogously, there is an unquote form, $\triangleleft(\sigma)$, that permits translational internal terms to be constructed compositionally. The supplement gives the analogous rules.

The final two forms of translational internal term are $\text{anatrans}[n](\sigma)$ and $\text{syntrns}[n]$. These stand in for the translation of argument n , the first if it arises via analysis against type σ and the second if it arises via type synthesis. They are not intended to be written directly by tycon providers,

arising only internally via argument interface lists. Before giving the rules, let us motivate the mechanism with the intro opcon definition for LPROD. We have $\omega_{\text{lprod}} := \text{ana intro} = \sigma_{\text{lprod/intro}}; \omega_{\text{lprod/targops}}$ where $\sigma_{\text{lprod/intro}} :=$

```

λtyidx:List[Lbl × Ty].λtmidx:List[Lbl].λargs:List[Arg].
  let inhabited : 1 = uniqmap tyidx
  listrec3 [Lbl × Ty] [Lbl] [Arg] [ITm] tyidx tmidx args ▷ (())
  λrowtyidx:Lbl × Ty.λrowtmidx:Lbl.λrowarg:Arg.λr:ITm.
    letpair (rowlbl, rowty) = rowtyidx
    let lok :: 1 = lbleq rowlbl rowtmidx
    let rowtr :: ITm = ana rowarg rowty
    ▷ ((<(rowtr), <(r)))

```

The first line checks that the type provided is inhabited, in this case by checking that there are no duplicate labels via the helper function *uniqmap* :: List[Lbl × Ty] → 1, raising an error if there are. An alternative strategy may have been to use an abstract kind that ensured that such type indices could not have been constructed, but to be compatible with our equality kind restriction, this would require support for abstract equality kinds, analogous to abstract equality types in SML. We chose not to formalize these for simplicity, and to demonstrate this general technique. An analogous technique could be used to implement record types by requiring that the index be sorted.

The rest of this opcon definition folds over the three lists provided as input: the list mapping labels to types provided as the type index, the list of labels provided as the term index, and the list of argument interfaces. We assume a straightforward helper function, *listfold3*, that raises an error if the three lists are not of the same length. The base case is the translational empty product. The recursive case first checks that the label provided in the term index matches the label provided in the type index, using a helper function *lbleq* :: Lbl → Lbl → 1. Then, we request type analysis of the corresponding argument, *rowarg*, against the type in the type index, *rowty*, by writing *ana rowarg rowty*. Here, *ana* is a helper function defined below that produces a translational internal term of the form $\triangleright(\text{anatrans}[n](\sigma))$ where σ is the value of *rowty* if the row value is well-typed and fails if not. The final line constructs a nested tuple based on this translation and the recursive result. Taken together, the translational internal term that will be derived for our example involving e_{ex} above is $\hat{e}_{\text{ex}} := (\text{anatrans}[0](\sigma_{\text{title}}), (\text{anatrans}[1](\sigma_{\text{conf}}), ()))$.

Argument Interface Lists We define Arg, the kind of *argument interfaces*, as a simple product of functions, $\text{Arg} := (\text{Ty} \rightarrow \text{ITm}) \times (1 \rightarrow \text{Ty} \times \text{ITm})$, one for analysis and the other for synthesis, both having the expected kind for these operations. The helper functions *ana* and *syn* simply project the corresponding function out, *ana* :: $\lambda \text{arg} :: \text{Arg}. \text{fst}(\text{arg})$ and *syn* :: $\lambda \text{arg} :: \text{Arg}. \text{snd}(\text{arg})$.

The argument interface list, called σ_{args} in the rules, is simply a static list of kind List[Arg] where the i th entry is $(\lambda \text{ty} :: \text{Ty}. \text{ana}[i](\text{ty}), \lambda \text{.} :: 1. \text{syn}[i])$. It is generated by the

judgement $\text{args}(n) = \sigma_{\text{args}}$, where n is given as the length of the argument list, written $|\bar{e}| = n$ (see supplement).

Recall that the kinding judgement is indexed by n , which is an upper bound on the argument index of terms of the form $\text{ana}[n](\sigma)$ and $\text{syn}[n]$. This is enforced in the kinding rules:

$$\begin{array}{c}
 \text{(k-ana)} \\
 \frac{n' < n \quad \Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{Ty}}{\Delta \Gamma \vdash_{\Phi}^n \text{ana}[n'](\sigma) :: \text{ITm}}
 \end{array}
 \qquad
 \begin{array}{c}
 \text{(k-syn)} \\
 \frac{n' < n}{\Delta \Gamma \vdash_{\Phi}^n \text{syn}[n'] :: \text{Ty} \times \text{ITm}}
 \end{array}$$

The following lemma characterizes thus characterizes the argument list interface generation judgement:

Lemma 1. *If $\text{args}(n) = \sigma_{\text{args}}$ then $\emptyset \emptyset \vdash_{\Phi}^n \sigma_{\text{args}} :: \text{List}[\text{Arg}]$.*

The rule (ocstruct-intro) ruled out writing either of these forms explicitly in an opcon definition by checking against the bound $n = 0$. This is to prevent out-of-bounds errors – tycon providers do not write these forms directly, but only access them via the argument interface list, guaranteed to have the correct length. This is possible because:

Lemma 2. *If $\Delta \Gamma \vdash_{\Phi}^{n'} \sigma :: \kappa$ and $n > n'$ then $\Delta \Gamma \vdash_{\Phi}^n \sigma :: \kappa$.*

The dynamics for these forms couples the dynamics of the static language to the statics of the external language. For $\text{ana}[n](\sigma)$, after stepping the type σ to a value and propagating errors, the argument environment, which stores the arguments themselves and the typing and tycon contexts, $\bar{e}; \Upsilon; \Phi$, is consulted to retrieve the n th argument and analyze it against σ . If this succeeds, the translational internal term $\triangleright(\text{anatrans}[n](\sigma))$ is generated to refer to it.

$$\begin{array}{c}
 \text{(s-ana-success)} \\
 \frac{\sigma \text{ val} \quad \text{nth}[n](\bar{e}) = e \quad \Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota}{\text{ana}[n](\sigma) \mapsto_{\bar{e}; \Upsilon; \Phi} \triangleright(\text{anatrans}[n](\sigma))}
 \end{array}$$

If it fails, an error is raised:

$$\begin{array}{c}
 \text{(s-ana-fail)} \\
 \frac{\sigma \text{ val} \quad \text{nth}[n](\bar{e}) = e \quad [\Upsilon \vdash_{\Phi} e \not\Leftarrow \sigma]}{\text{ana}[n](\sigma) \text{ err}_{\bar{e}; \Upsilon; \Phi}}
 \end{array}$$

We write $[\Upsilon \vdash_{\Phi} e \not\Leftarrow \sigma]$ to indicate that e fails to analyze against σ . We do not define this inductively, so we also allow that this premise be omitted, leaving a non-deterministic semantics nevertheless sufficient for our metatheory.

The dynamics for $\text{syn}[n]$ are analogous, evaluating to a pair $(\sigma, \triangleright(\text{syntrans}[n]))$ where σ is the synthesized type.

The kinding rules also prevent these translational internal term forms generated by these static operations from being well-kinded when $n = 0$. Like the translational internal type form $\text{trans}(\sigma)$, these are retained in values of kind ITm:

$$\begin{array}{c}
 \text{(s-itm-anatrans-v)} \\
 \frac{\sigma \text{ val} \quad \text{nth}[n](\bar{e}) = e}{\triangleright(\text{anatrans}[n](\sigma)) \text{ val}_{\bar{e}; \Upsilon; \Phi}}
 \end{array}
 \qquad
 \begin{array}{c}
 \text{(s-itm-syntrans-v)} \\
 \frac{\text{nth}[n](\bar{e}) = e}{\triangleright(\text{syntrans}[n]) \text{ val}_{\bar{e}; \Upsilon; \Phi}}
 \end{array}$$

Abstracted Term Translations The reason for this is again because we will hold these abstract by replacing them with

variables. The judgement $\hat{\iota} \parallel \mathcal{D} \mathcal{G} \mapsto_{\mathcal{A}}^{\text{TC}} \iota \parallel \mathcal{D}' \mathcal{G}'$, appearing as the eighth premise of (ana-intro), relates a translational internal term $\hat{\iota}$ to an internal term ι called the corresponding *abstracted term translation*, where all references to the translation of an argument (having any type) are replaced with a unique variable and all references to the translation of a type constructed by a user-defined tycon other than the “delegated tycon” TC are replaced with an abstract type variable as described in Sec. 3.4.1. The type translation store, \mathcal{D} , discussed previously, and term translation store, \mathcal{G} , track these mappings. Term translation stores have the following syntax: $\mathcal{G} ::= \emptyset \mid \mathcal{G}, n : \sigma \rightsquigarrow \iota/x : \tau$. Each entry can be read “argument n having type σ and translation ι appears as variable x with type τ ”. For example,

$$\hat{\iota}_{\text{ex}} \parallel \emptyset \mapsto_{(\text{title}, \text{"EXMPL 2015"}), \Upsilon, \Phi_{\text{rstr}}, \Phi_{\text{lprod}}}^{\text{LPROD}} \iota_{\text{ex}/\text{abs}} \parallel \mathcal{D}_{\text{ex}/\text{abs}} \mathcal{G}_{\text{ex}/\text{abs}}$$

where $\iota_{\text{ex}/\text{abs}} := (x_0, (x_1, ()))$, the term translation store is $\mathcal{G}_{\text{ex}/\text{abs}} := 0 : \sigma_{\text{title}} \rightsquigarrow \text{title}/x_0 : \alpha_0, 1 : \sigma_{\text{conf}} \rightsquigarrow \text{"EXMPL 2015"}_{\text{IL}}/x_1 : \alpha_1$ and the type translation store $\mathcal{D}_{\text{ex}/\text{abs}}$ is alpha-equivalent to $\mathcal{D}_{\text{paper}/\text{abs}}$ from Sec. 3.4.1.

The rules for the abstracted term translation judgement follow recursively for shared forms just like those for the abstracted type translation judgement (see supplement). The rules for references to argument translations derived via type analysis operates as follows (syntrans[n] is analogous):

$$\begin{array}{c} \text{(abs-anatrans-new)} \\ \Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota \quad n \notin \text{dom}(\mathcal{G}) \quad \text{nth}[n](\bar{e}) = e \quad \text{trans}(\sigma) \parallel \mathcal{D} \mapsto_{\Phi}^{\text{c}} \tau \parallel \mathcal{D}' \quad (x \text{ fresh}) \\ \hline \text{anatrans}[n](\sigma) \parallel \mathcal{D} \mathcal{G} \mapsto_{\bar{e}; \Upsilon; \Phi}^{\text{c}} x \parallel \mathcal{D}' \mathcal{G}, n : \sigma \rightsquigarrow \iota/x : \tau \end{array}$$

Like \mathcal{D} , each \mathcal{G} induces an internal term substitution, $\gamma := \emptyset \mid \gamma, \iota/x$, and corresponding internal typing context Γ by the judgement $\mathcal{G} \rightsquigarrow \gamma : \Gamma$, appearing as the ninth premise. In this case, $\gamma_{\text{ex}/\text{abs}} := \text{title}/x_0, \text{"EXMPL 2015"}_{\text{IL}}/x_1$ and $\Gamma_{\text{ex}/\text{abs}} := x_0 : \alpha_0, x_1 : \alpha_1$.

The tenth premise of (ana-intro) determines an abstract type translation for the type provided for analysis, in this case $\tau_{\text{ex}/\text{abs}} := \alpha_0 \times (\alpha_1 \times 1)$ (alpha-equivalent to $\tau_{\text{conf}/\text{abs}}$ in Sec. 3.4.1), and the eleventh premise extracts a type substitution, $\delta_{\text{ex}/\text{abs}}$, and type formation context, $\Delta_{\text{ex}/\text{abs}}$, from $\mathcal{D}_{\text{ex}/\text{abs}}$, again equivalent to $\delta_{\text{conf}/\text{abs}}$ and $\Delta_{\text{conf}/\text{abs}}$ from Sec. 3.4.1.

Finally, the twelfth premise checks the abstracted term translation against the abstracted type translation. Here, we are checking $\Delta_{\text{ex}/\text{abs}} \Gamma_{\text{ex}/\text{abs}} \vdash \iota_{\text{ex}/\text{abs}} : \tau_{\text{ex}/\text{abs}}$, i.e.:

$$(\alpha_0, \alpha_1) (x_0 : \alpha_0, x_1 : \alpha_1) \vdash (x_0, (x_1, ())) : \alpha_0 \times (\alpha_1 \times 1)$$

In other words, the translation of the labeled product e_{ex} generated by LPROD is checked with all references to term and type translations of regular strings replaced by variables and type variables, respectively. Nevertheless, because our intro opcon definition treated arguments parametrically, the check succeeds. We will describe an ill-behaved operator, which would lead to a violation of the regular string tycon invariant were it not disallowed by this check, in Sec. 5.

Applying the substitutions $\gamma_{\text{ex}/\text{abs}}$ and $\delta_{\text{ex}/\text{abs}}$ in the conclusion of the rule, we arrive at ι_{ex} .

Other Intro The rule (ana-intro-other) is used to introduce values of a type constructed by an “other” tycon. The term index, rather than the tycon context, directly specifies the static function that maps the arguments to a translation. In all other respects, the rule is analogous. It is used as a technical device in our proof of conservativity in Sec. 5.

4.2 Generalized Targeted Operations

All non-introductory opcons associated with user-defined tycons go through another generalized form, in this case for *targeted operations*, $\text{targ}[\text{op}; \sigma_{\text{tmidx}}](e_{\text{targ}}; \bar{e})$, where **op** ranges over opcon names, σ_{tmidx} is the term index, e_{targ} is the *target argument* and \bar{e} are the remaining arguments.

Desugaring to this form include an explicit form, $e_{\text{targ}}.\text{op}(\sigma_{\text{tmidx}})(\bar{e})$ (and variants of where the term index or the arguments are omitted), projection syntax for use by record-like types, $e_{\text{targ}}\#1\text{bl}$, which desugars to $\text{targ}[\text{prj}; 1\text{bl}](e_{\text{targ}}; \cdot)$, and concatenative syntax $e_{\text{targ}} \cdot e_{\text{arg}}$, which desugars to $\text{targ}[\text{conc}; ()](e_{\text{targ}}; e_{\text{arg}})$. We show other desugarings, including case analysis, in the supplement.

Whereas introductory operations were analytic, targeted operations are synthetic in $@\lambda$. The type and translation is determined by the tycon of the type synthesized by the target argument. The rule (syn-targ) is otherwise similar to (ana-intro) in its structure. The first premise synthesizes a type, $\text{TC}(\sigma_{\text{tyidx}})$, for the target argument. The second premise extracts the tycon definition for TC from the tycon context. The third extracts the *operator index kind* from its opcon signature, and the fourth checks the term index against this kind. For example, we may associate the following opcon signatures with the tycons RSTR and LPROD:

$$\begin{array}{ll} \chi_{\text{rstr}} & := \text{intro}[\text{Str}], \text{conc}[1], \text{case}[\text{List}[\text{StrPattern}]], \\ & \quad \text{coerce}[\text{Rx}], \text{check}[\text{Rx}], \text{replace}[\text{Rx}] \\ \chi_{\text{lprod}} & := \text{intro}[\text{List}[\text{Lbl}]], \text{prj}[\text{Lbl}], \text{conc}[1], \text{drop}[\text{List}[\text{Lbl}]] \end{array}$$

The opcons associated with RSTR are taken directly from Fulton et al.’s specification of regular string types [11], with the exception of **case**, which generalizes case analysis as defined there to arbitrary string patterns, which we discuss in the supplement. The opcons associated with LPROD are also straightforward: **prj** projects out the row with the provided label, **conc** concatenates two labeled products (updating common fields with the value of the right argument), and **drop** creates a new labeled product from the target with some fields dropped. Note that both regular strings and labeled products define concatenation. Targeted operations use no new mechanisms, so we will only show regular string concatenation here; others are in the supplement.

The fifth premise of (syn-targ) extracts the targeted opcon definition of **op** from the opcon structure, ω . Like the intro opcon definition, this is a static function that generates a translational internal term on the basis of the target tycon’s

type index, the term index and an argument interface list. Targeted opcon definitions additionally synthesize a type. The rule (ocstruct-targ) in Figure 8 ensures that no tycon defines an opcon twice and that the opcon definitions are well-kinded. For example, $\omega_{\text{rstr}/\text{targops}}$ defines:

```
syn conc = λtyidx::Rx.λtmidx::1.λargs::List[Arg].
  letpair (arg1, arg2) = arity2 args
  letpair (_, tr1) = syn arg1
  letpair (ty2, tr2) = syn arg2
  typecase[RSTR](ty2, tyidx2.
    (RSTR(rxconcat tyidx tyidx2), ▷(sconcat ◁(tr1) ◁(tr2)))
  ; raise[Ty × ITm])
```

The first line of the body checks that exactly two arguments, including the target argument, were passed in via helper function *arity2*. We then request synthesis of both arguments. We can ignore the type synthesized by the first because by definition it is a regular string type with type index *tyidx*. We case analyze the second against RSTR, extracting its index regular expression if so and failing if not. We finally synthesize the resulting regular string type, using the helper function *rxconcat* :: Rx → Rx → Rx which generates the concatenated regular expression, and the translation, using an internal helper function *sconcat* : str → str → str, the translational internal term for which we assume has been substituted in directly above.

The remaining premises of (syn-targ) are analagous to the corresponding premises in (ana-intro), with the only difference being that we check the abstract term translation against the abstract type translation of the synthesized type. Like (ana-intro-other), the rule (syn-targ-other) is used when the target synthesizes an “other” type. The mapping from the arguments to a type and translation is given directly in the term index (the operator name is ignored).

5. Metatheory

We will now state the key metatheoretic properties of @λ. The full proofs are in the supplement.

Kind Safety Kind safety ensures that normalization of well-kinded static terms cannot go wrong. We can take a standard progress and preservation based approach.

Theorem 1 (Static Progress). *If $\emptyset \vdash \sigma :: \kappa$ and $\vdash \Phi$ and $|\bar{e}| = n$ and $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ then $\sigma \text{ val}_{\bar{e};\Phi;\Upsilon}$ or $\sigma \text{ err}_{\bar{e};\Phi;\Upsilon}$ or $\sigma \mapsto_{\bar{e};\Phi;\Upsilon} \sigma'$.*

Theorem 2 (Static Preservation). *If $\emptyset \vdash \sigma :: \kappa$ and $\vdash \Phi$ and $|\bar{e}| = n$ and $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ and $\sigma \mapsto_{\bar{e};\Phi;\Upsilon} \sigma'$ then $\emptyset \vdash \sigma' :: \kappa$.*

The case in the proof for the form $\text{syn}[n]$ requires that the following theorem be simultaneously defined. The mutual induction is well-founded because the total length of all the argument lists in the terms being considered decreases.

Theorem 3 (Type Synthesis). *If $\vdash \Phi$ and $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ and $\Upsilon \vdash_{\Phi} e \Rightarrow \sigma \rightsquigarrow \iota$ then $\vdash_{\Phi} \sigma \rightsquigarrow \tau$ (and thus $\sigma \text{ type}_{\Phi}$).*

Type Safety Type safety in a typed translation semantics requires that well-typed external terms translate to well-typed internal terms. Type safety for the IL [12] then implies that well-typed external terms cannot go wrong. To prove this, we must prove a stronger theorem, type-preserving translation (the analog of type-preserving compilation in the typed compilation literature [27]):

Theorem 4 (Type Preserving Translation). *If $\vdash \Phi$ and $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ and $\vdash_{\Phi} \sigma \rightsquigarrow \tau$ and $\Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota$ then $\emptyset \vdash \iota : \tau$.*

Proof Sketch. We induct on the typing judgement. The interesting cases are (ana-intro), (ana-intro-other), (syn-trans) and (syn-trans-other); the latter two arise via subsumption. The result follows directly from the final premise of each rule, combined with a lemma that states that the substitutions γ and δ generated by the abstracted type and term translation stores satisfy the corresponding Δ_{abs} and Γ_{abs} , and so applying them in the conclusion gives a well-typed term. \square

Hygienic Translation Note in the just mentioned rules that the domains of Υ and Γ_{abs} are disjoint. This serves to ensure *hygienic translation* – translations cannot refer to variables in the surrounding scope directly, so uniformly renaming a variable cannot change the meaning of a program. Variables in Υ can occur in arguments (e.g. *title* in the previous example), but the translations of the arguments only appear *after* the substitution γ has been applied. We assume that substitution is capture-avoiding in the usual manner.

Unicity The rules are structured so that if a term is well-typed, both its type and translation are unique.

Theorem 5 (Unicity). *If $\vdash \Phi$ and $\vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma$ and $\vdash_{\Phi} \sigma \rightsquigarrow \tau$ and $\vdash_{\Phi} \sigma' \rightsquigarrow \tau'$ and $\Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota$ and $\Upsilon \vdash_{\Phi} e \Leftarrow \sigma' \rightsquigarrow \iota'$ then $\sigma = \sigma'$ and $\iota = \iota'$.*

Stability Extending the tycon context does not change the meaning of any terms that were previously well-typed.

Theorem 6 (Stability). *Letting $\Phi' := \Phi, \text{tycon TC } \{\theta\} \sim \psi$, if $\vdash \Phi'$ and $\vdash_{\Phi'} \Upsilon \rightsquigarrow \Gamma$ and $\vdash_{\Phi} \sigma \rightsquigarrow \tau$ and $\Upsilon \vdash_{\Phi} e \Leftarrow \sigma \rightsquigarrow \iota$ then, $\vdash_{\Phi'} \Upsilon \rightsquigarrow \Gamma$ and $\vdash_{\Phi'} \sigma \rightsquigarrow \tau$ and $\Upsilon \vdash_{\Phi'} e \Leftarrow \sigma \rightsquigarrow \iota$.*

Conservativity Extending the tycon context also conserves all *tycon invariants* maintained in any smaller tycon context. An example of a tycon invariant is the following:

Tycon Invariant 1 (Regular String Soundness). *If $\emptyset \vdash_{\Phi_{\text{rstr}}} e \Leftarrow \text{RSTR}\langle r \rangle \rightsquigarrow \iota$ and $\iota \Downarrow \iota'$ then $\iota' = \text{"s"}$ and “s” is in the regular language $\mathcal{L}(r)$.*

Proof Sketch. The proof is not unusually difficult because we have fixed the tycon context Φ_{rstr} , so we can simply treat the calculus like a type-directed compiler for a calculus with only two tycons, \rightarrow and RSTR. Such a calculus and compiler specification was given by Fulton et al. [11] and extending this to our setting requires showing only that the opcon definitions in RSTR adequately capture these specifications

right
citation
for
TIL?

could
move
this
whole
thing
to
supple-
ment
if
room
needed

using standard program correctness techniques for the SL, which is a simple functional language. The only “wrinkle” is that the rule (syn-targ-other) can synthesize a regular string type, but the abstracted translation will always be checked against $\tau_{\text{abs}} = \alpha$, so no strings that could not arise by another rule could be generated, maintaining the invariant. \square

The reason why the rule (syn-targ-other) is not a problem in proving a tycon invariant turns out to be the same reason extending the tycon context also cannot violate such a tycon invariant. Any newly introduced tycon defining a targeted operator that synthesizes a regular string type, e.g. σ_{paper} , and generating a translation that is not in the corresponding regular language, e.g. “”, could be defined, but when used, the rule (syn-targ) would check this translation against $\tau_{\text{abs}} = \alpha$ and the check would fail. We can state this conservativity property more generally.

Theorem 7 (Conservativity). *If $\vdash \Phi$ and $\text{TC} \in \text{dom}(\Phi)$ and a tycon invariant for TC holds under Φ :*

- *If $\emptyset \vdash_{\Phi} e \Leftarrow \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow \iota$ and $\vdash_{\Phi} \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow \tau$ and $\iota \Downarrow \iota'$ then $P(\sigma_{\text{tyidx}}, \iota')$.*

then for all $\Phi' = \Phi, \text{tycon TC} \{ \theta' \} \sim \psi'$ such that $\vdash \Phi'$, the same tycon invariant holds under Φ' :

- *If $\emptyset \vdash_{\Phi'} e \Leftarrow \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow \iota$ and $\vdash_{\Phi'} \text{TC}(\sigma_{\text{tyidx}}) \rightsquigarrow \tau$ and $\iota \Downarrow \iota'$ then $P(\sigma_{\text{tyidx}}, \iota')$.*

Proof Sketch. The proof maps every well-typed term in the extended tycon context to a well-typed term in the original tycon context with the same translation, so we can apply the original invariant to arrive at the property $P(\sigma_{\text{tyidx}}, \iota')$. This term is constructed by replacing all types constructed by TC' with a type constructed by $\text{other}[m]$, for an m uniquely associated with TC' , and all introductory and targeted sub-terms associated with TC' with an equivalent one that passes through the rules (ana-intro-other) and (syn-targ-other) by simply partially applying the intro and targeted opcon definitions to generate the term indices. Types constructed by $\text{other}[m]$ are indistinguishable from types constructed by a “future” tycon because the type case analysis construct is not defined for this tycon form, so the translations coincide. \square

6. Related Work

Term Rewriting Language-integrated static term rewriting (“macro”) systems, like Template Haskell [26] and Scala’s static macros [5], can be used to decrease complexity when an isomorphic embedding into the underlying type system is possible, but cannot extend the type system directly.

Optimization When a simple embedding that preserves the static semantics exists, but a different embedding would better preserve a desired cost semantics, term rewriting techniques can also be used to perform “optimizations”, thus achieving an isomorphic embedding. Care must be taken, however, to ensure that the optimized value is not manipulated directly if the rewriting does not preserve the

static semantics (i.e. the rewritten term has a different, less constrained type). Type abstraction has been used directly to achieve this property in work on *lightweight modular staging* [24]. This does not allow new type systems to be defined (in LMS, Scala’s type system).

Type Refinements When new static distinctions need to be made for values of an existing type, but new primitive operations are not needed, one solution is to develop a system of *type refinements* based on the fragment of interest [10]. For example, one might refine the type of integers to distinguish negative integers. Most proposals for “pluggable type systems” describe such type refinement systems [4], which require only additional annotations supplied as comments or metadata (e.g. JavaCOP [2]). Refinement systems do not support holding the refined type abstract as we did in our regular string example (RSTR could have picked a non-string representation for, e.g., fixed size regular languages)

Language Frameworks The most general situation, of which type refinements are a special case, is when exposing a fragment requires defining new types as well as new operators, where the static and dynamic semantics governing these new operators make non-trivial use of statically valued information. We saw labeled product types added, which require a new type constructor and new operator constructors.

A variety of *language frameworks* have been developed to make it easier to define new languages, in some cases together with a compiler and other tooling for these languages; see [9]. These tools can decrease the effort needed to define a new language, but they either do not support forming languages from separately defined fragments or provide few modular reasoning principles that make it possible to reason separately about these fragments.

Module Systems Our work has some technical similarities to work on maintaining type abstraction using a form of effect system [8], and on translating modules to System F [25]. Unlike modules, tycons are indexed by arbitrary static values, permit the representation of a type to be computed based on this index, and compute both the type and translation of each operation, rather than requiring that the operations be expressible as a finite collection of functions.

7. Discussion

unlimited
refer-
ences
so
add a
bunch
here

References

- [1] GHC/FAQ. http://www.haskell.org/haskellwiki/GHC:FAQ#Extensible_Records.
- [2] C. Andreae, J. Noble, S. Markstrum, and T. Millstein. A framework for implementing pluggable type systems. In *OOPSLA, OOPSLA '06*, pages 57–74, New York, NY, USA, 2006. ACM.
- [3] N. Benton and A. Kennedy. Interlanguage Working Without Tears: Blending SML with Java. In *ICFP '99*, pages 126–137, New York, NY, USA, 1999. ACM.
- [4] G. Bracha. Pluggable type systems, Oct. 2004. *OOPSLA Workshop on Revival of Dynamic Languages*.
- [5] E. Burmako. Scala Macros: Let Our Powers Combine!: On How Rich Syntax and Static Types Work with Metaprogramming. In *Proceedings of the 4th Workshop on Scala, SCALA '13*, pages 3:1–3:10, New York, NY, USA, 2013. ACM.
- [6] L. Cardelli. A semantics of multiple inheritance. In *Proc. Of the International Symposium on Semantics of Data Types*, pages 51–67, New York, NY, USA, 1984. Springer-Verlag New York, Inc.
- [7] M. Chakravarty, R. Leshchinskiy, S. Jones, G. Keller, and S. Marlow. Data Parallel Haskell: a status report. In *Proceedings of the 2007 Workshop on Declarative Aspects of Multi-core Programming*, pages 10–18. ACM, 2007.
- [8] K. Crary, R. Harper, and D. Dreyer. A type system for higher-order modules. In *ACM SIGPLAN–SIGACT Symposium on Principles of Programming Languages (POPL)*, Portland, Oregon, 2002.
- [9] S. Erdweg, T. van der Storm, M. Völter, M. Boersma, R. Bosman, W. R. Cook, A. Gerritsen, A. Hulshout, S. Kelly, A. Loh, et al. The state of the art in language workbenches. In *Software Language Engineering*. 2013.
- [10] T. Freeman and F. Pfenning. Refinement types for ML. In *Proceedings of the SIGPLAN '91 Symposium on Language Design and Implementation*, pages 268–277, Toronto, Ontario, June 1991. ACM Press.
- [11] N. Fulton, C. Omar, and J. Aldrich. Statically typed string sanitation inside a python. In *First International Workshop on Privacy and Security in Programming (PSP 2014)*. ACM, 2014.
- [12] R. Harper. *Practical Foundations for Programming Languages*. Cambridge University Press, 2012.
- [13] R. Harper and C. Stone. A Type-Theoretic Interpretation of Standard ML. In *IN Proof, Language and Interaction: Essays in Honour of Robin Milner*. MIT Press, 2000.
- [14] A. Kennedy. Types for units-of-measure: Theory and practice. In Z. Horváth, R. Plasmeijer, and V. Zsók, editors, *CEFP*, volume 6299 of *Lecture Notes in Computer Science*, pages 268–305. Springer, 2009.
- [15] X. Leroy, D. Doligez, A. Frisch, J. Garrigue, D. Rémy, and J. Vouillon. *The OCaml system release 4.01 Documentation and user's manual*. Institut National de Recherche en Informatique et en Automatique, September 2013.
- [16] A. Löb and R. Hinze. Open data types and open functions. In *PPDP*, pages 133–144. ACM, 2006.
- [17] L. Mandel and M. Pouzet. ReactiveML: a reactive extension to ML. In *PPDP*, pages 82–93. ACM, 2005.
- [18] A. Miller, M. Hicks, J. Katz, and E. Shi. Authenticated Data Structures, Generically. In *POPL '14*, pages 411–423, New York, NY, USA, 2014. ACM.
- [19] T. D. Millstein, C. Bleckner, and C. Chambers. Modular typechecking for hierarchically extensible datatypes and functions. *ACM Trans. Program. Lang. Syst.*, 26(5):836–889, 2004.
- [20] T. Murphy, VII., K. Crary, and R. Harper. Type-safe Distributed Programming with ML5. In *Proceedings of the 3rd Conference on Trustworthy Global Computing, TGC'07*, pages 108–123, Berlin, Heidelberg, 2008. Springer-Verlag.
- [21] A. Ohori and K. Ueno. Making Standard ML a Practical Database Programming Language. In *ICFP, ICFP '11*, pages 307–319, New York, NY, USA, 2011. ACM.
- [22] C. Omar, D. Kurilova, L. Nistor, B. Chung, A. Potanin, and J. Aldrich. Safely composable type-specific languages. In *ECOOP*, 2014.
- [23] B. C. Pierce and D. N. Turner. Local type inference. *ACM Trans. Program. Lang. Syst.*, 22(1):1–44, Jan. 2000.
- [24] T. Rompf and M. Odersky. Lightweight modular staging: a pragmatic approach to runtime code generation and compiled DSLs. *Communications of the ACM*, 55(6):121–130, June 2012.
- [25] A. Rossberg, C. V. Russo, and D. Dreyer. F-ing modules. In A. Kennedy and N. Benton, editors, *TLDI*, pages 89–102. ACM, 2010.
- [26] T. Sheard and S. Peyton Jones. Template metaprogramming for Haskell. In M. M. T. Chakravarty, editor, *ACM SIGPLAN Haskell Workshop 02*, pages 1–16. ACM Press, Oct. 2002.
- [27] D. Tarditi, G. Morrisett, P. Cheng, C. Stone, R. Harper, and P. Lee. TIL: A type-directed optimizing compiler for ML. In *PLDI '96*, Philadelphia, PA, May 1996.

A. Appendix

$$\begin{array}{c}
\text{(s-ty-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{c\langle\sigma\rangle \mapsto_{\mathcal{A}} c\langle\sigma'\rangle}
\end{array}
\quad
\begin{array}{c}
\text{(s-ty-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{c\langle\sigma\rangle \text{ err}_{\mathcal{A}}}
\end{array}
\quad
\begin{array}{c}
\text{(s-ty-v)} \\
\frac{\sigma \text{ val}_{\mathcal{A}}}{c\langle\sigma\rangle \text{ val}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-otherty-v)} \\
\hline
\text{otherty}[m; \tau] \text{ val}_{\mathcal{A}}
\end{array}$$

$$\begin{array}{c}
\text{(s-tycase-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\text{tycase}[c](\sigma; \mathbf{x}.\sigma_1; \sigma_2) \mapsto_{\mathcal{A}} \text{tycase}[c](\sigma'; \mathbf{x}.\sigma_1; \sigma_2)}
\end{array}$$

$$\begin{array}{c}
\text{(s-tycase-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\text{tycase}[c](\sigma; \mathbf{x}.\sigma_1; \sigma_2) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-tycase-match)} \\
\frac{c\langle\sigma\rangle \text{ val}_{\mathcal{A}}}{\text{tycase}[c](c\langle\sigma\rangle; \mathbf{x}.\sigma_1; \sigma_2) \mapsto_{\mathcal{A}} [\sigma/x]\sigma_1}
\end{array}$$

$$\begin{array}{c}
\text{(s-tycase-fail)} \\
\frac{\sigma \neq c\langle\sigma'\rangle}{\text{tycase}[c](\sigma; \mathbf{x}.\sigma_1; \sigma_2) \mapsto_{\mathcal{A}} \sigma_2}
\end{array}$$

$$\begin{array}{c}
\text{(keq-k)} \\
\frac{k \in \Delta}{\Delta \vdash k \text{ eq}}
\end{array}
\quad
\begin{array}{c}
\text{(keq-ind)} \\
\frac{\Delta, k \vdash \kappa \text{ eq}}{\Delta \vdash \mu_{\text{ind}}(k.\kappa) \text{ eq}}
\end{array}
\quad
\begin{array}{c}
\text{(keq-unit)} \\
\frac{}{\Delta \vdash 1 \text{ eq}}
\end{array}$$

$$\begin{array}{c}
\text{(keq-prod)} \\
\frac{\Delta \vdash \kappa_1 \text{ eq} \quad \Delta \vdash \kappa_2 \text{ eq}}{\Delta \vdash \kappa_1 \times \kappa_2 \text{ eq}}
\end{array}
\quad
\begin{array}{c}
\text{(keq-sum)} \\
\frac{\Delta \vdash \kappa_1 \text{ eq} \quad \Delta \vdash \kappa_2 \text{ eq}}{\Delta \vdash \kappa_1 + \kappa_2 \text{ eq}}
\end{array}$$

$$\begin{array}{c}
\text{(keq-ty)} \\
\hline
\Delta \vdash \text{Ty eq}
\end{array}$$

$$\begin{array}{c}
\text{(k-ity-alpha)} \\
\hline
\Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\alpha) :: \text{ITy}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-lam-step-1)} \\
\frac{\blacktriangleright(\hat{\tau}_1) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau}'_1)}{\blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau}'_1 \times \hat{\tau}_2)}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-lam-step-2)} \\
\frac{\blacktriangleright(\hat{\tau}_1) \text{ val}_{\mathcal{A}} \quad \blacktriangleright(\hat{\tau}_2) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau}'_2)}{\blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau}_1 \times \hat{\tau}'_2)}
\end{array}
\quad
\begin{array}{c}
\text{(s-ity-lam-err-1)} \\
\frac{\blacktriangleright(\hat{\tau}_1) \text{ err}_{\mathcal{A}}}{\blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-lam-err-2)} \\
\frac{\blacktriangleright(\hat{\tau}_2) \text{ err}_{\mathcal{A}}}{\blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) \text{ err}_{\mathcal{A}}}
\end{array}
\quad
\begin{array}{c}
\text{(s-ity-lam-v)} \\
\frac{\blacktriangleright(\hat{\tau}_1) \text{ val}_{\mathcal{A}} \quad \blacktriangleright(\hat{\tau}_2) \text{ val}_{\mathcal{A}}}{\blacktriangleright(\hat{\tau}_1 \times \hat{\tau}_2) \text{ val}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-alpha-v)} \\
\hline
\blacktriangleright(\alpha) \text{ val}_{\mathcal{A}}
\end{array}$$

$$\begin{array}{c}
\text{(k-tycase-parr)} \\
\frac{\Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{Ty} \quad \Delta \Gamma, \mathbf{x} :: \text{Ty} \times \text{Ty} \vdash_{\Phi}^n \sigma_1 :: \kappa \quad \Delta \Gamma \vdash_{\Phi}^n \sigma_2 :: \kappa}{\Delta \Gamma \vdash_{\Phi}^n \text{tycase}[\rightarrow](\sigma; \mathbf{x}.\sigma_1; \sigma_2) :: \kappa}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-unquote-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\blacktriangleright(\blacktriangleleft(\sigma)) \mapsto_{\mathcal{A}} \blacktriangleright(\blacktriangleleft(\sigma'))}
\end{array}
\quad
\begin{array}{c}
\text{(s-ity-unquote-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\blacktriangleright(\blacktriangleleft(\sigma)) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-ity-trans-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\blacktriangleright(\text{trans}(\sigma)) \mapsto_{\mathcal{A}} \blacktriangleright(\text{trans}(\sigma'))}
\end{array}
\quad
\begin{array}{c}
\text{(s-ity-trans-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\blacktriangleright(\text{trans}(\sigma)) \text{ err}_{\mathcal{A}}}
\end{array}$$

This judgement is defined by the following straightforward rules:

$$\begin{array}{c}
\text{(tstore-emp)} \\
\frac{}{\emptyset \rightsquigarrow \emptyset : \emptyset}
\end{array}
\quad
\begin{array}{c}
\text{(tstore-ext)} \\
\frac{\mathcal{D} \rightsquigarrow \delta : \Delta}{(\mathcal{D}, \sigma \leftrightarrow \tau/\alpha) \rightsquigarrow (\delta, \tau/\alpha) : (\Delta, \alpha)}
\end{array}$$

Description	Concrete Form	Desugared Form
sequences	(e_1, \dots, e_n) or $[e_1, \dots, e_n]$	$\text{intro}[(\cdot)](e_1; \dots; e_n)$
labeled sequences	$\{\text{lbl}_1 = e_1, \dots, \text{lbl}_n = e_n\}$	$\text{intro}[[\text{lbl}_1, \dots, \text{lbl}_n]](\cdot)$
label application	$\text{lbl}\langle e_1, \dots, e_n \rangle$	$\text{intro}[\text{lbl}](e_1, \dots, e_n)$
numerals	n	$\text{intro}[n](\cdot)$
labeled numerals	$n \text{ lbl}$	$\text{intro}[(n, \text{lbl})](\cdot)$
strings	"s"	$\text{intro}["s"](\cdot)$

$$\begin{array}{c}
\text{(s-itm-var-v)} \\
\frac{}{\triangleright(x) \text{ val}_{\mathcal{A}}}
\end{array}
\quad
\begin{array}{c}
\text{(s-itm-lam-step-1)} \\
\frac{}{\blacktriangleright(\hat{\tau}) \mapsto_{\mathcal{A}} \blacktriangleright(\hat{\tau}')}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-lam-step-2)} \\
\frac{\blacktriangleright(\hat{\tau}) \text{ val}_{\mathcal{A}} \quad \triangleright(i) \mapsto_{\mathcal{A}} \triangleright(i')}{\triangleright(\lambda[\hat{\tau}](x.i)) \mapsto_{\mathcal{A}} \triangleright(\lambda[\hat{\tau}'](x.i'))}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-lam-err-1)} \\
\frac{\blacktriangleright(\hat{\tau}) \text{ err}_{\mathcal{A}}}{\triangleright(\lambda[\hat{\tau}](x.i)) \text{ err}_{\mathcal{A}}}
\end{array}
\quad
\begin{array}{c}
\text{(s-itm-lam-err-2)} \\
\frac{}{\triangleright(i) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-lam-v)} \\
\frac{\blacktriangleright(\hat{\tau}) \text{ val}_{\mathcal{A}} \quad \triangleright(i) \text{ val}_{\mathcal{A}}}{\triangleright(\lambda[\hat{\tau}](x.i)) \text{ val}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(k-itm-unquote)} \\
\frac{\Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{ITm}}{\Delta \Gamma \vdash_{\Phi}^n \triangleright(\blacktriangleleft(\sigma)) :: \text{ITm}}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-unquote-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\triangleright(\blacktriangleleft(\sigma)) \mapsto_{\mathcal{A}} \triangleright(\blacktriangleleft(\sigma'))}
\end{array}
\quad
\begin{array}{c}
\text{(s-itm-unquote-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\triangleright(\blacktriangleleft(\sigma)) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-unquote-elim)} \\
\frac{}{\triangleright(i) \text{ val}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-ana-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\text{ana}[n](\sigma) \mapsto_{\mathcal{A}} \text{ana}[n](\sigma')}
\end{array}
\quad
\begin{array}{c}
\text{(s-ana-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\text{ana}[n](\sigma) \text{ err}_{\mathcal{A}}}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-anatrans-step)} \\
\frac{\sigma \mapsto_{\mathcal{A}} \sigma'}{\triangleright(\text{anatrans}[n](\sigma)) \mapsto_{\mathcal{A}} \triangleright(\text{anatrans}[n](\sigma'))}
\end{array}$$

$$\begin{array}{c}
\text{(s-itm-anatrans-err)} \\
\frac{\sigma \text{ err}_{\mathcal{A}}}{\triangleright(\text{anatrans}[n](\sigma)) \text{ err}_{\mathcal{A}}}
\end{array}$$

, as specified by the judgement $\mathcal{G} \rightsquigarrow \gamma : \Gamma$ defined by the following rules:

$$\begin{array}{c} \text{(ttrs-emp)} \\ \hline \emptyset \rightsquigarrow \emptyset : \emptyset \end{array} \quad \begin{array}{c} \text{(ttrs-ext)} \\ \hline \mathcal{G} \rightsquigarrow \gamma : \Gamma \\ \hline (\mathcal{G}, n : \sigma \rightsquigarrow \iota/x : \tau) \rightsquigarrow (\gamma, \iota/x) : (\Gamma, x : \tau) \end{array}$$

$$\begin{array}{c} \text{(k-itm-lam)} \\ \hline \Delta \Gamma \vdash_{\Phi}^n \blacktriangleright(\hat{\tau}) :: \text{ITy} \quad \Delta \Gamma \vdash_{\Phi}^n \triangleright(\hat{\iota}) :: \text{ITm} \\ \hline \Delta \Gamma \vdash_{\Phi}^n \triangleright(\lambda[\hat{\tau}](x.\hat{\iota})) :: \text{ITm} \end{array}$$

$$\begin{array}{c} \text{(k-raise)} \\ \hline \Delta \vdash \kappa \\ \hline \Delta \Gamma \vdash_{\Phi}^n \text{raise}[\kappa] :: \kappa \end{array} \quad \begin{array}{c} \text{(s-raise)} \\ \hline \text{raise}[\kappa] \text{ err}_{\mathcal{A}} \end{array}$$

$$\begin{array}{c} \text{(s-syn-success)} \\ \hline \text{nth}[n](\bar{e}) = e \quad \Upsilon \vdash_{\Phi} e \Rightarrow \sigma \rightsquigarrow \iota \\ \hline \text{syn}[n] \mapsto_{\bar{e}; \Upsilon; \Phi} (\sigma, \triangleright(\text{syntrans}[n])) \end{array}$$

$$\begin{array}{c} \text{(s-syn-fail)} \\ \hline \text{nth}[n](\bar{e}) = e \quad [\Upsilon \vdash_{\Phi} e \not\Rightarrow] \\ \hline \text{syn}[n] \text{ err}_{\bar{e}; \Upsilon; \Phi} \end{array}$$

$$\begin{array}{c} \text{(k-itm-syntrans)} \\ \hline n' < n \\ \hline \Delta \Gamma \vdash_{\Phi}^n \triangleright(\text{syntrans}[n']) :: \text{ITm} \end{array}$$

, e.g. for lambdas:

$$\begin{array}{c} \text{(abs-lam)} \\ \hline \hat{\tau} \parallel \mathcal{D} \rightsquigarrow_{\Phi}^{\text{TC}} \tau \parallel \mathcal{D}' \quad \hat{\iota} \parallel \mathcal{D}' \mathcal{G} \rightsquigarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} \iota \parallel \mathcal{D}'' \mathcal{G}' \\ \hline \lambda[\hat{\tau}](x.\hat{\iota}) \parallel \mathcal{D} \mathcal{G} \rightsquigarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} \lambda[\tau](x.\iota) \parallel \mathcal{D}'' \mathcal{G}' \end{array}$$

$$\begin{array}{c} \text{(abs-anatrans-stored)} \\ \hline n : \sigma \rightsquigarrow \iota/x : \tau \in \mathcal{G} \\ \hline \text{anatrans}[n](\sigma) \parallel \mathcal{G} \mathcal{D} \rightsquigarrow_{\mathcal{A}}^{\text{TC}} x \parallel \mathcal{G} \mathcal{D} \end{array}$$

$$\begin{array}{c} \text{(abs-syntrans-stored)} \\ \hline n : \sigma \rightsquigarrow \iota/x : \tau \in \mathcal{G} \\ \hline \text{syntrans}[n] \parallel \mathcal{G} \mathcal{D} \rightsquigarrow_{\mathcal{A}}^{\text{TC}} x \parallel \mathcal{G} \mathcal{D} \end{array}$$

$$\begin{array}{c} \text{(k-itm-anatrans)} \\ \hline n' < n \quad \Delta \Gamma \vdash_{\Phi}^n \sigma :: \text{Ty} \\ \hline \Delta \Gamma \vdash_{\Phi}^n \triangleright(\text{anatrans}[n'](\sigma)) :: \text{ITm} \end{array}$$

$$\begin{array}{c} \text{(abs-syntrans-new)} \\ \hline n \notin \text{dom}(\mathcal{G}) \quad \text{nth}[n](\bar{e}) = e \\ \Upsilon \vdash_{\Phi} e \Rightarrow \sigma \rightsquigarrow \iota \quad \text{trans}(\sigma) \parallel \mathcal{D} \rightsquigarrow_{\Phi}^{\text{TC}} \tau \parallel \mathcal{D}' \quad (x \text{ fresh}) \\ \hline \text{syntrans}[n] \parallel \mathcal{G} \mathcal{D} \rightsquigarrow_{\bar{e}; \Upsilon; \Phi}^{\text{TC}} x \parallel \mathcal{G}, n : \sigma \rightsquigarrow \iota/x : \tau \mathcal{D}' \end{array}$$

$$\begin{array}{c} \text{(etctx-emp)} \\ \hline \vdash_{\Phi} \emptyset \rightsquigarrow \emptyset \end{array} \quad \begin{array}{c} \text{(etctx-ext)} \\ \hline \vdash_{\Phi} \Upsilon \rightsquigarrow \Gamma \quad \sigma \text{ type}_{\Phi} \quad \vdash_{\Phi} \sigma \rightsquigarrow \tau \\ \hline \vdash_{\Phi} \Upsilon, x \Rightarrow \sigma \rightsquigarrow \Gamma, x : \tau \end{array}$$

Description	Concrete Form	Desugared Form
index projection	$e_{\text{targ}} \# n$	$\text{targ}[\text{idx}; n](e_{\text{targ}})$
label projection	$e_{\text{targ}} \# \text{lbl}$	$\text{targ}[\text{prj}; \text{lbl}](e_{\text{targ}})$
explicit invocation	$e_{\text{targ}} \cdot \text{op}[\sigma_{\text{tmidx}}](\bar{e})$ $e_{\text{targ}} \cdot \text{op}(\bar{e})$ $e_{\text{targ}} \cdot \text{op}(\text{lbl}_1 = e_1, \dots, \text{lbl}_n = e_n)$	$\text{targ}[\text{op}; \sigma_{\text{tmidx}}](e_{\text{targ}})$ $\text{targ}[\text{op}; ()](e_{\text{targ}})$ $\text{targ}[\text{op}; [\text{lbl}_1, \dots, \text{lbl}_n]](e_{\text{targ}})$
labeled case analysis	$e_{\text{targ}} \cdot \text{case} \{$ $\mid \sigma_1 \langle x_1, \dots, x_k \rangle \Rightarrow e_1$ $\mid \dots$ $\mid \sigma_n \langle x_1, \dots, x_k \rangle \Rightarrow e_n \}$	$\text{targ}[\text{case}; [\sigma_1, \dots, \sigma_n]](e_{\text{targ}})$ $\lambda(x_1 \dots \lambda(x_k) \dots)$ $\lambda(x_1 \dots \lambda(x_k) \dots)$

For example,

$$\begin{array}{c} \text{(abs-prod)} \\ \hline \hat{\tau}_1 \parallel \mathcal{D} \rightsquigarrow_{\Phi}^{\text{TC}} \tau_1 \parallel \mathcal{D}' \quad \hat{\tau}_2 \parallel \mathcal{D}' \rightsquigarrow_{\Phi}^{\text{TC}} \tau_2 \parallel \mathcal{D}'' \\ \hline \hat{\tau}_1 \times \hat{\tau}_2 \parallel \mathcal{D} \rightsquigarrow_{\Phi}^{\text{TC}} \tau_1 \times \tau_2 \parallel \mathcal{D}'' \end{array}$$

The argument interfaces that populate the list provided to opcon definitions is derived from the argument list by the judgement $\text{args}(\bar{e}) =_n \sigma_{\text{args}}$, defined as follows:

$$\begin{array}{c} \text{(args-z)} \\ \hline \text{args}(\cdot) =_0 \text{nil}[\text{Arg}] \end{array}$$

$$\begin{array}{c} \text{(args-s)} \\ \hline \text{args}(\bar{e}) =_n \sigma \\ \hline \text{args}(\bar{e}; e) =_{n+1} \text{rcons}[\text{Arg}] \sigma (\lambda \text{ty} :: \text{Ty.ana}[n](\text{ty}), \lambda \dots : 1. \text{syn}[n]) \end{array}$$

We assume that the definitions of the standard helper functions $\text{nil} :: \forall(\alpha. \text{List}[\alpha])$ and $\text{rcons} :: \forall(\alpha. \text{List}[\alpha] \rightarrow \alpha \rightarrow \text{List}[\alpha])$, which adds an item to the end of a list, have been substituted into these rules. The result is that the n th element of the argument interface list simply wraps the static terms $\text{ana}[n](\sigma)$ and $\text{syn}[n]$.