

# Ace: An Actively-Typed Compilation Environment for High-Performance Computing

Cyrus Omar, Nathan Fulton, Sharschchandra Bidargaddi and Jonathan Aldrich

*School of Computer Science*

*Carnegie Mellon University*

*Pittsburgh, PA, USA*

*<http://www.acelang.org/>*

**Abstract**—In this paper, we introduce the Ace compilation environment and demonstrate its practicality as a foundational tool in both high-performance computing research and in practice. Ace consists of a statically-typed language that shares syntax with Python and uses the Python language itself as a multi-purpose compile-time metalanguage. Unlike standard languages, however, the Ace compiler consists only of a minimal core that delegates control over most aspects of compilation to user libraries using a novel extension mechanism that we call *active type-checking and translation (AT&T)*. In particular, users can specify new primitive types and operations by equipping type definitions, which are first-class objects in the metalanguage, with special methods. The compiler then invokes these methods when type checking and translating expressions, according to a fixed type dispatch protocol associated with each syntactic form. Translation targets a user-defined backend language, and the safety of user-defined translations can be checked automatically. Using this mechanism, we implement the primitives of the C99 and OpenCL languages as simple modules within Ace. We go on show examples of other interesting primitive parallel abstractions that build on top of these basic modules, demonstrating the foundational role that an actively-typed system like Ace can play in support of research on language-based parallel programming abstractions.

In addition to AT&T, Ace supports more conventional forms of structured and unstructured metaprogramming and integrates directly with the numeric capabilities of Python itself. We use these features to design a scientific simulation framework that allows users to modularly specify and orchestrate the execution of parameterized families of simulations on clusters of GPUs. This framework is used to successfully conduct large-scale, high-performance neuroscience simulations, providing evidence that Ace can be used in practice today.

## I. INTRODUCTION

Researchers design and implement novel parallel programming abstractions on a regular basis. Each abstraction aims to make some class of developers more productive by making appropriate trade-offs between performance, portability, verifiability and ease-of-use. It can be observed that the most widely-adopted abstractions to date are those that have been implemented as libraries that can be easily deployed within existing, widely-used languages. Abstractions that require introducing new primitive constructs, changing the semantics of existing constructs in particular contexts, or controlling compilation in a fine-grained manner cannot take

this approach. This has led to a proliferation of specialized programming languages, each designed around a few privileged abstractions. Unfortunately, such languages have not been widely adopted by practitioners.

Computer-aided simulations and data analysis techniques have transformed science and engineering. Surveys show that scientists and engineers now spend up to 40% of their time writing software, generally individually or in small groups [?][?]. Most of this software targets conventional desktop hardware, but about 20% of scientists also target either local clusters or supercomputers for more numerically-intensive computations [?]. In recent years, GPUs and other kinds of coprocessors have emerged as powerful platforms for high-performance computing as well.

Many scientists and engineers prefer high-level languages like MATLAB, Python, R and Perl [?], due to their powerful domain-specific libraries and their focus on productivity and flexibility. These languages are typically interpreted rather than compiled and generally do not feature static type systems, relying instead on run-time (“dynamic”) type lookup and indirection. Although this can increase the flexibility of the language, these features also negatively impact performance. Along code paths where performance appears to be a bottleneck, developers typically turn to traditional statically-typed low-level languages like C and Fortran [?]. These languages require explicit type annotations on variables and give users explicit control over data layout and heap allocation. Although this makes writing programs more tedious and error-prone, it can also significantly improve performance, particularly if attributes of the target architecture (e.g. caching behavior) are considered.

While a sequential reimplementing of a critical code path using a low-level language will produce a modest speedup, more significant speedups on modern hardware require parallelizing over many processor cores and, on massively-parallel machines, many interconnected nodes. Despite well-known difficulties, low-level approaches that use a form of shared-memory multithreading (e.g. pthreads, OpenMP) paired with explicit message passing between nodes (typically using MPI) remain widely used [?][?].

Although researchers often propose higher-level language

features and parallel programming abstractions that aim to strike a better balance between raw performance, productivity, code portability and verifiability (see Section 2), end-users remain skeptical of new approaches. This viewpoint was perhaps most succinctly expressed by a participant interviewed in a recent study [?], who stated “I hate MPI, I hate C++. [But] if I had to choose again, I would probably choose the same.” Although this sentiment is easy to dismiss as paradoxical, we believe that it demands direct examination by those in the research community working to advance the practice of scientific and high-performance computing. We begin in Section 2 by developing a set of design and adoption criteria for new languages and abstractions based on prior empirical studies of these *professional end-user developers* as well as our observations of characteristics common to successful projects in the past.

We then introduce a new programming language named *cl.oquence* (pronounced “C eloquence”) in Section 3. *cl.oquence* initially targets code paths where low-level languages, particularly OpenCL, would be used today. By employing a collection of powerful front-end language and compiler design techniques, making pragmatic design choices, and building on top of established infrastructure, we argue that *cl.oquence* may uniquely satisfy many of the criteria established in Section 2, particularly those related to ease of use, familiarity, performance and tool support.

Unlike many existing languages, *cl.oquence* is designed to be highly extensible from inception, employing a powerful compile-time language extension and code generation mechanism based around the concept of *active libraries* [?]. We argue that this feature may make *cl.oquence* particularly suitable as a platform for researchers developing new parallel abstractions and other specialized, high-level language features, as well as development tools.

In Section 4 we briefly describe a case study where the language was used to conduct large-scale spiking neural circuit simulations on GPU hardware, demonstrating the feasibility of this approach on a non-trivial problem. Finally, conclude in Section 6 with future directions for both the *cl.oquence* language and the research community.

## II. THE CL.OQUENCE LANGUAGE

We now describe a programming language designed with many of the criteria described in the previous sections in mind, *cl.oquence*. *cl.oquence* is built around a minimal core and nearly every construct in the language is implemented as a library. Rather than beginning directly with this extension mechanism, however, let us begin with a concrete example of the language in use.

### A. Example: Higher-Order Map for OpenCL

Figure 1 shows the standard data-parallel map function written using an extension that implements the full OpenCL language as a library. Each thread, indexed by *gid*, applies

```
import clq
from clq.backends.opencl import
    OpenCL, get_global_id

@clq.fn(OpenCL)
def map(in, out, fn):
    gid = get_global_id(0)
    out[gid] = fn(in[gid])
```

Figure 1. Data-parallel higher-order map in *cl.oquence*. The *get\_global\_id* function is an OpenCL primitive function.

```
#pragma OPENCL EXTENSION cl_khr_fp64 : enable

__kernel void map_sin_dbl(__global double *in,
                          __global double *out)
{
    size_t gid = get_global_id(0);
    out[gid] = sin(in[gid]);
}
```

Figure 2. An OpenCL kernel function that maps the *sin* function over a vector of double-precision floating-point numbers in the global memory address space.

the transfer function, *fn*, to the corresponding element of the input array, *in*, writing the result into the corresponding location in the output array, *out*.

### B. Syntax

Readers familiar with the Python programming language will recognize the style of syntax used in Figure 1. In fact, *cl.oquence* uses the Python grammar and parsing facilities directly. Several factors motivated this design decision. First, Python’s syntax is widely credited as being particularly simple and readable, due to its use of significant whitespace and conventional mathematical notation. Python is one of the most widely-used languages in scientific computing, so its syntax is already familiar to much of the field. And significantly, a large ecosystem of tools already exist that work with Python files, such as code editors, syntax highlighters, style checkers and documentation generators. These can be used without modification to work with *cl.oquence* files. Therefore, by re-using an existing, widely-used grammar, we are able to satisfy many of the design criteria described in Section ?? and the adoption criteria described in Section ?? without significant development effort.

### C. Semantics

*cl.oquence* is a compiled, statically-typed language. The extension used in Figure 1 implements the full OpenCL type system, which includes features inherited from C99 like pointer arithmetic and OpenCL-specific constructs like vector types. In fact, compilation of Figure 1 will produce code semantically equivalent, and with the same performance profile, as manually-written OpenCL code. Compilation produces OpenCL source code, so *cl.oquence*, when used with

this extension, inherits OpenCL’s portability profile as well, as described in Section ??.

In OpenCL itself, however, there is no way to write a function like `map`. As described in Section ??, OpenCL lacks support for higher-order functions or polymorphism (over the types of `in` and `out` in this case.) To map a function over an array, developers must create a separate function for every combination of argument types and for every transfer function. Figure 2 shows an OpenCL kernel that implements a `map` specifically for the `sin` function applied to arrays of doubles in global memory. Note that because OpenCL does not support templates or function pointers, this is the most general solution possible without explicit code generation. However, even in CUDA, which does support templates, there is greater syntactic overhead for specializing a function with particular types than in `cl.quence`, where it happens implicitly.

To support this much more concise form for specifying kernels, `cl.quence` uses techniques pioneered by functional programming languages. We briefly describe these below.

1) *Type, Kernel and Extension Inference*: As described in Section ??, professional end-users tend to prefer languages that do not require explicit type annotations for variables. These languages accomplish this by using dynamic typing, so that types are associated with values instead. Statically-typed functional programming languages have also eliminated type annotations in many cases, however, using a technique known as *type inference* [?].

Type inference is a technique for giving types to variables by examining how these variables are being used. As a simple example, in Figure 1, the variable `gid` is being used to hold the result of calling the `get_global_id` function so the type of `gid` must be consistent with the return type of that function, `size_t` (a device-dependent integer type.) Note, however, that in C99, many possible integer types can be given to `gid` (in practice, a 32-bit integer is often used.) Picking a specific type may require further examining how `gid` is used in the remainder of the function. To deal with this non-locality, type inference is typically cast as a constraint solving problem – the type of each program variable is treated as an unknown and each use of the variable introduces a new constraint. A constraint-solving algorithm, generally a variant of unification, can be used to find a type assignment that satisfying all constraints.

`cl.quence` uses a variant of this approach (modified to admit extensions to the type system, as we will discuss shortly) to conduct whole-function type inference. This approach differs from the HM algorithm, however, which conducts whole-program (or whole-file) inference and assigns a single type to function arguments. In `cl.quence`, types are *propagated* to functions from their calling sites (to enable structural polymorphism, described below) and inference occurs only within a function body.

In addition to type annotations, OpenCL introduces addi-

tional syntactic burdens. Developers must annotate functions that meet the requirements to be called from the host with the `__kernel` attribute, and several types (notably, `double`) and specialized functions require that an OpenCL extension be enabled with a `#pragma` (as can be seen in Figure 2.) The OpenCL extension we have developed automatically infers many of these annotations as well, consistent with the principle of conciseness described in Section ??.

2) *Structural Polymorphism*: In Section ??, we discussed several strategies for achieving *polymorphism* – the ability to create functions and data structures that operate over more than a single type. In `cl.quence`, all functions are implicitly polymorphic and can be called with arguments of *any type that supports the operations used by the function*. For example, in Figure 1, `in` can be any type that supports indexing by a variable of type `size_t` to produce a value of a type that can be passed into `fn`, which must then produce a value consistent with indexing into `out`. OpenCL pointer types are consistent with these constraints, for example. Although powerful, this also demonstrates a caveat of this approach – that it is more difficult to give a function a concise signature, because arguments are constrained by capability, rather than to a single type [?].

Structural typing can be compared to the approach taken by dynamically-typed languages that rely on “duck typing”. It is more flexible than the parametric polymorphism found in many functional languages and in languages like Java (which only allow polymorphic functions that are valid for *all* possible types), but is of comparable strength to the template system found in C++. It can be helpful to think of each function as being preceded by an implicit template header that assigns each argument its own unique type parameter. At function call sites, these parameters are implicitly specialized with the types of the provided arguments. This choice is again motivated by the criteria of conciseness given in Section ??.

3) *Higher-Order Functions*: Section ?? also discusses the value of higher-order functions. The OpenCL extension supports higher-order functions, despite the absence of function pointers from OpenCL itself, by explicitly passing functions as constants at compile-time. In other words, each function *uniquely inhabits* a type corresponding to it.

Functions are not, however, *first-class* – references to functions cannot be stored in memory, for example. Most use cases that we have considered thus far have never needed to do so, however. Compared to the function pointer approach used in other C-based languages, this approach is more efficient. It can, however, result in “code bloat”, since a new copy of the higher-order function is created for each unique function argument passed to it. First-class functions can theoretically be simulated using an auxiliary lookup function in OpenCL, and we leave clean support for this in `cl.quence` as open future work.

```

from figure1 import map
from clq.backends.opencl import double, sin

map_sin_double = map.specialize(
    double.global_ptr, double.global_ptr,
    sin.clq_type)

```

Figure 3. Programmatically specializing the `map` function from Figure 1 for use by the standalone compiler to produce code equivalent to Figure 2.

#### D. Active Libraries in *cl.ouquence*

In addition to adopting Python’s grammar, *cl.ouquence* uses the Python language itself as a *compile-time met-language*. Every *cl.ouquence* file is, at the top-level, a Python script. This script is executed during compilation and has access to a number of powerful capabilities that allow developers to influence the compilation process and manipulate programs directly. Libraries that have such capabilities has been called *active libraries* in prior proposals [?]. A number of projects, such as Blitz++, have taken advantage of the C++ preprocessor and template-based metaprogramming system to implement domain-specific optimizations. In *cl.ouquence*, we replace these brittle mini-languages with a general-purpose language. This allows for several interesting uses that we discuss in the following sections.

#### E. Module System

In Figure 1, the only compile-time operations are library imports and a declaration of a *cl.ouquence* function. However, this already demonstrates an important benefit of this approach: *cl.ouquence* uses Python’s simple but flexible module system. As a consequence, libraries written using *cl.ouquence* can directly utilize the module distribution infrastructure available for Python, partially relieving the issues discussed in Section ??.

#### F. Compilation and Invocation

*cl.ouquence* functions can be invoked in one of two ways: from any host language supporting OpenCL by using the standalone compiler, or directly from Python itself.

1) *Programmatic Specialization and Compilation*: The standalone *cl.ouquence* compiler, `clqcc`, generates source code from *cl.ouquence* source files. To do so, however, users must *specialize* any *externally callable* functions with their specific types. Figure 3 shows how the specific OpenCL function given in Figure 2 can be recovered from the generic specification of `map` given in Figure 1. The `specialize` method of a generic function like `map` creates a new function that is concrete – that is, with concrete argument types.

The standalone compiler operates by first executing the module, then iterating over the externally available variables in the module’s environment to extract the OpenCL functions and any other functions and constructs that they depend on. This allows developers who wish to use the host language of their choice, rather than Python, to continue

```

from figure1 import map
import numpy as np
import clq.opencl.pyopencl as cl

ctx = cl.Context.for_device(0, 0)
d_in = ctx.to_device(np.ones(1024))
d_out = ctx.alloc(like=d_in)
map(d_in, d_out, clq.opencl.sin,
    global_size=d_in.shape, local_size=(128,))
out = ctx.from_device(d_out)

```

Figure 4. A full OpenCL program using the *cl.ouquence* Python bindings, including data transfer and kernel invocation.

writing library functions generically, in one place, only writing small stub modules to specialize them as needed.

2) *Direct Invocation from Python*: As discussed in the introduction, a common usage scenario in scientific computing involves the use of a dynamic, high-level language for workflow orchestration and exploratory visualization and analysis, paired with a low-level language for performance-critical sections. Although the extension mechanism we describe shortly will permit the development of a broad array high-level language features in the future, current extensions, such as the OpenCL extension we have been describing, have been designed to make low-level programming simpler. For this reason, the OpenCL extension includes support for directly invoking *cl.ouquence* functions from Python scripts with minimal overhead.

This feature builds on top of an established library, `pyopencl`, that wraps the OpenCL host API and integrates it with the popular `numpy` numerics library in Python. Buffers (arrays) in `pyopencl` do not retain type information, so the *cl.ouquence* extension adds type information, mirroring the API of `numpy` arrays. It also hides many of the most verbose components of the OpenCL interface by providing reasonable defaults, such as an implicit global context and a default queue associated with each context.

Using these simplifications, generic *cl.ouquence* OpenCL functions like `map` can be called directly, without specialization as in the previous section. Figure 4 shows a full OpenCL program written using these bindings. By way of comparison, the same program written using OpenCL directly is two orders of magnitude larger and correspondingly more complex. Not shown are several additional conveniences, such as delegated kernel sizing and In and Out constructs that can reduce the size and improve the clarity of this code further – due to space constraints, the reader is referred to the language documentation for additional details.

#### G. Code Generation in *cl.ouquence*

Metaprogramming refers to the practice of writing programs that manipulate other programs. There are a number of use cases for this technique, including domain-specific optimizations and code generation for programs with a

```

class PtrType(Type):
    def __init__(self, target_type, addr_space):
        self.target_type = target_type
        self.addr_space = addr_space

    def verify_Subscript(self, context, node):
        slice_type = context.validate(node.slice)
        if isinstance(slice_type, IntegerType):
            return self.target_type
        else:
            raise TypeError('<error message here>')

    def translate_Subscript(self, context, node):
        value = context.translate(node.value)
        slice = context.translate(node.slice)
        return copy_node(node,
            value = value,
            slice = slice,
            code = value.code + '['+slice.code+']')

```

Figure 5. A portion of the implementation of OpenCL pointer types implementing subscripting logic using the `cl.ouquence` extension mechanism.

repetitive structure that cannot easily be packaged using available abstractions. OpenCL in particular relies on code generation as a fundamental mechanism, partially justifying the lack of support for higher-order programming.

`cl.ouquence` supports code generation from directly within the metalanguage. A `cl.ouquence` function can be constructed from a string containing its source using the `clq.from_source` function. It can also be constructed directly from a Python abstract syntax tree, available via the standard `ast` package, using the `clq.from_ast` function. We discuss a use case for code generation for simulation orchestration in `cl.ouquence` in Section 4.

#### H. Active Language Extensions

We now turn our attention to the extension mechanism used by `cl.ouquence`, as motivated by section ???. The `cl.ouquence` type checker and compiler interacts directly with active libraries over an interface that gives substantial control over type checking and translation of an operation to a *type definition* associated with one of its operands according to a fixed dispatch protocol.

1) *Dispatch Protocol*: When the compiler encounters an expression, it must first verify that it type checks, then generate code for it. Rather than containing fixed logic for this, however, the compiler defers this responsibility to the *type* of a subexpression whenever possible. Below are example from the `cl.ouquence` dispatch protocol. Due to space constraints, we do not list the entire dispatch protocol.

- Responsibility over a **unary operation** like `-x` is handed to the type assigned to the operand, `x`.
- Responsibility over **binary operations** is first handed to the type assigned to the left operand. If it indicates that it does not understand the operation, the type assigned

to the right operand is handed responsibility, with a different method call<sup>1</sup>.

- Responsibility over **attribute access**, `obj.attr`, and **subscript access**, `obj[idx]`, is handed to the type assigned to `obj`.

2) *Verification and Translation*: A type in `cl.ouquence` is a metalanguage *instance* of a Python class deriving from `clq.Type`. The compiler hands control over an expression to a type by calling a method corresponding to the syntactic form of the expression according to the above dispatch protocol. Figure 5 gives the verification and translation logic governing subscript access for pointer types in OpenCL.

During the verification phase, the type of the primary operand, as determined by the dispatch protocol, is responsible for assigning a type to the expression as a whole. In this case, it recursively assigns a type to the slice operand. If it is assigned an integer type, the pointer’s target type is assigned to the expression as a whole. Otherwise, a type error is raised with the provided error message (an ability that can be used by extension authors to satisfy the error message criteria described in Section ??.)

If verification succeeds, the compiler must subsequently translate the `cl.ouquence` expression into an expression in the output language, here OpenCL. It does so by again applying the dispatch protocol to call a method prefixed with `translate_`. This method is responsible for returning a copy of the expression’s ast node with an additional attribute, `code`, containing the source code of the translation. In this case, it is simply a direct translation to the corresponding OpenCL attribute access, using the recursively-determined translations of the operands. More sophisticated abstractions may insert arbitrarily complex statements and expressions. The context also provides some support for non-local effects, such as new top-level declarations (not shown.)

3) *Modular Backends*: Thus far, we have discussed using OpenCL as a backend with `cl.ouquence`. The OpenCL extension is the most mature as of this writing. However, `cl.ouquence` supports the introduction of new backends in a manner similar to the introduction of new types, by extending the `clq.Backend` base class. Backends are provided as the first argument to the `@clq.fn` decorator, as can be seen in Figure 1. Backends are responsible for some aspects of the grammar that do not admit simple dispatch to the type of a subterm, such as number and string literals or basic statements like `while`.

In addition to the OpenCL backend, preliminary C99 and CUDA backends are available (with the caveat that they have not been as fully developed or tested as of this writing.) Backends not based on the C family are also possible, but we leave such developments for future work.

<sup>1</sup>Note that this operates similarly to Python’s operator overloading protocol.

4) *Use Cases*: The development of the full OpenCL language using only the extension mechanisms described above provides evidence of the power of this approach. Nothing about the core language was designed specifically for OpenCL. However, to be truly useful, as described in Sections ?? and ??, the language must be able to support a wide array of primitive abstractions. We briefly describe a number of other abstractions that may be possible using this mechanism. Many of these are currently available either via inconvenient libraries or in standalone languages. With the `cl.ouquence` extension mechanism, we hope to achieve robust, natural implementations of many of these mechanisms within the same language.

*Partitioned Global Address Spaces*: A number of recent languages in high-performance computing have been centered around a partitioned global address space model, including UPC, Chapel, X10 and others. These languages provide first-class support for accessing data transparently across a massively parallel cluster, which is verbose and poorly supported by standard C. The extension mechanism of `cl.ouquence` allows inelegant library-based approaches such as the Global Arrays library to be hidden behind natural wrappers that can use compile-time information to optimize performance and verify correctness. We have developed a prototype of this approach using the C backend and hope to expand upon it in future work.

*Other Parallel Abstractions*: A number of other parallel abstractions, some of which are listed in ??, also suffer from inelegant C-based implementations that spurred the creation of standalone languages. A study comparing a language-based concurrency solution for Java with an equivalent, though less clean, library-based solution found that language support is preferable but leads to many of the issues we have described [?]. The extension mechanism is designed to enable library-based solutions that operate as first-class language-based solutions, barring the need for particularly exotic syntactic extensions.

*Domain-Specific Type Systems*: `cl.ouquence` is a statically-typed language, so a number of domain-specific abstractions that promise to improve verifiability using types, as discussed in Section ??, can be implemented using the extension mechanism. We hope that this will allow advances from the functional programming community to make their way into the professional end-user community more quickly, particularly those focused on scientific domains (e.g. [?]).

*Specialized Optimizations*: In many cases, code optimization requires domain-specific knowledge or sophisticated, parametrizable heuristics. Existing compilers make implementing and distribution such optimizations difficult. With active libraries in `cl.ouquence`, optimizations can be distributed directly with the libraries that they work with. For instance, we have implemented substantial portions of the NVidia GPU-specific optimizations described in [?] as a library that uses the extension mechanism to track affine

transformations of the thread index used to access arrays, in order to construct a summary of the memory access patterns of the kernel, which can be used both for single-kernel optimization (as in [?]) and for future research on cross-kernel fusion and other optimizations.

*Instrumentation*: Several sophisticated feedback-directed optimizations and adaptive run-time protocols require instrumenting code in other ways. The extension mechanism enables granular instrumentation based on the form of an operation as well as its constituent types, easing the implementation of such tools. This ability could also be used to collect data useful for more rigorous usability and usage studies of languages and abstractions, and we plan on following up on this line of research going forward.

### I. Availability

`cl.ouquence` is available under the LGPL license and is developed openly and collaboratively using the popular Github platform<sup>2</sup>. Documentation and other learning material is being made available at <http://cl.ouquence.org/>.

## III. CASE STUDY: NEUROBIOLOGICAL CIRCUIT SIMULATION

An important criteria that practitioners use to evaluate a language or abstraction, as discussed in Section ??, is whether significant case studies have been conducted with it. In this section, we briefly (due to space limitations) discuss an application of the `cl.ouquence` OpenCL library, Python host bindings and code generation features for developing a modular, high-performance scientific simulation library used to simulate thousands of parallel realizations of a spiking neurobiological circuit on a GPU.

### A. Background

A neural circuit can be modeled as a network of coupled differential equations, where each node corresponds to a single neuron. Each neuron is modeled using one or more ordinary differential equations. These equations capture the dynamics of physically important quantities like the cell's membrane potential or the conductance across various kinds of ion channels and can take many forms [?]. Single simulations can contain from hundreds to tens of millions of neurons each, depending on the specific problem being studied. In some cases, such as when studying the effects of noise on network dynamics or to sweep a parameter space, hundreds or thousands of realizations must be generated. In these cases, care must be taken to only probe the simulation for relevant data and process portions of it as the simulation progresses, because the amount of data generated is often too large to store in its entirety for later analysis.

The research group we discuss here (of which the first author was a member) was studying a problem that required running up to 1,000 realizations of a network of between

<sup>2</sup><https://github.com/cyrus-cl.ouquence>

4,000 and 10,000 neurons each. An initial solution to this problem used the Brian framework, written in Python, to conduct these simulations on a CPU. Brian was selected because it allowed the structure of the simulation to be specified in modular and straightforward manner. This solution required between 60 and 70 minutes to conduct the simulations and up to 8 hours to analyze the data each time a parameter of the simulation was modified.

Unsatisfied with the performance of this approach, the group developed an accelerated variant of the simulation using C++ and CUDA. Although this produced significant speedups, reducing the time for a simulation by a factor of 40 and the runtime of the slowest analyses by a factor of 200, the overall workflow was also significantly disrupted. In order to support the many variants of models, parameter sets, and probing protocols, C preprocessor flags were necessary to selectively include or exclude code snippets. This quickly led to an incomprehensible and difficult to maintain file structure. Moreover, much of the simpler data analysis and visualization was conducted using Python, so marshalling the relevant data between processes also became an issue.

### B. The cl.egans Simulation Library

In order to eliminate these issues while retaining the performance profile of the GPU-accelerated code, the project was ported to cl.oquence. Rather than using preprocessor directives to control the code contained in the final GPU kernels used to execute the simulation and data analyses, the group was able to develop a more modular library called cl.egans<sup>3</sup> based on the language’s compile-time code generation mechanism and Python and OpenCL bindings.

cl.egans leverages Python’s object-oriented features to enable modular, hierarchical simulation specifications. For example, Figure 6 shows an example where a neuron model (ReducedLIF) is added to the root of the simulation, a synapse model (ExponentialSynapse) is then added to it, and its conductance is probed in the same way, by adding a probe model as a child of the synapse model. If interleaved analysis needed to be conducted as well, it would be specified in the same way.

Implementations of these classes do not evaluate the simulation logic directly, but rather contain methods that generate cl.oquence source code for insertion at various points, called *hooks*, in the final simulation kernel. The hook that code is inserted into is determined by the method name, and code can be inserted into any hook defined anywhere upstream in the simulation tree. New hooks can also be defined in these methods and these become available for use by child nodes. Figure 7 shows an example of a class that inserts code in the `model_code` hook and defines several new hooks. This protocol is closely related to the notion of *frame-oriented programming*. Although highly modular, this

```
# Create the root node of the simulation
sim = Simulation(ctx, n_timesteps=10000)
neurons = ReducedLIF(sim, count=N, tau=20.0)
e_synapse = ExponentialSynapse(neurons, 'ge',
                                tau=5.0, reversal=60.0)
probe = StateVariableProbeCopyback(e_synapse)
```

Figure 6. An example of a nested simulation tree, showing that specifying a simulation is both simple and modular.

```
class SpikingModel(Model):
    """Base class for spiking neuron models."""
    def in_model_code(self, g):
        """
        idx_state = idx_model + (realization_num -
                                realization_start)*count
        """ << g # << appends to code generator, g
        self.insert_cg_hook("read_incoming", g)
        self.insert_cg_hook("read_state", g)
        self.insert_cg_hook("calculate_inputs", g)
        self.insert_cg_hook("state_calculations", g)
        self.insert_cg_hook("spike_processing", g)
        # ...
```

Figure 7. An example of a hook that inserts code and also inserts new, nested hooks for downstream simulation nodes below that.

strategy avoids the performance penalties associated with standard object-oriented methodologies via code generation.

Compared to a similar protocol targeting OpenCL directly, the required code generation logic is significantly simpler because it enables classes like `StateVariable` to be written generically for all types of state variables, without carrying extra parameters and *ad hoc* logic to extract and compute the result types of generated expressions. Moreover, because types are first-class objects in the metalanguage, they can be examined during the memory allocation step to enable features like fully-automatic parallelization of multiple realizations across one or more devices, a major feature of cl.egans that competing frameworks cannot easily offer.

Once the kernel has been generated and memory has been allocated, the simulation can be executed directly from Python using the bindings described in Section II-F2. The results of this simulation are immediately available to the Python code following the simulation and can be visualized and further analyzed using standard tools. Once the computations are complete, the Python garbage collector is able to handle deallocation of GPU memory automatically (a feature of the underlying `pyopencl` library [?].)

Using this cl.oquence-based framework, the benefits of the Brian-based workflow were recovered without the corresponding decrease in performance relative to the previous CUDA-based solution, leading ultimately to a satisfying solution for the group conducting this research.

## IV. CONCLUSION

Professional end-users demand much from new languages and abstractions. In this paper, we began by generating a

<sup>3</sup>...after *c. elegans*, a model organism in neuroscience

concrete, detailed set of design and adoption criteria that we hope will be of broad interest and utility to the research community. Based on these constraints, we designed a new language, `cl.oquence`, making several pragmatic design decisions and utilizing advanced techniques, including type inference, structural typing, compile-time metaprogramming and active libraries, to uniquely satisfy many of the criteria we discuss, particularly those related to extensibility. We validated the extension mechanism with a mature implementation of the entirety of the OpenCL type system, as well as preliminary implementations of some other features. Finally, we demonstrated that this language was useful in practice, drastically improving performance without negatively impacting the high-level scientific workflow of a large-scale neurobiological circuit simulation project. Going forward, we hope that `cl.oquence` (or simply the key techniques it proposes, by some other vehicle) will be developed further by the community to strengthen the foundations upon which new abstractions are implemented and deployed into professional end-user development communities.