

Statically Typed String Sanitation Inside a Python

Nathan Fulton

Cyrus Omar

Jonathan Aldrich

Carnegie Mellon University
Pittsburgh, PA

{nathanfu, comar, aldrich}@cs.cmu.edu

ABSTRACT

Web applications must ultimately command systems like web browsers and database engines using strings. Strings derived from improperly sanitized user input can thus be a vector for command injection attacks. In this paper, we introduce *regular string types*, which classify strings known statically to be in a specified regular language. These types come equipped with common operations like concatenation, substitution and coercion, so they can be used to implement, in essentially a conventional manner, the portions of a web application or web application framework that must directly construct command strings. Simple type annotations at key interfaces can be used to statically verify that sanitization has been performed correctly without introducing redundant run-time checks. We specify this type system as a minimal typed lambda calculus, λ_{RS} .

To be practical, adopting a specialized type system like this should not require the adoption of a new programming language. Instead, we advocate for extensible type systems: new type system fragments like this should be implemented as libraries atop a mechanism that guarantees that they can be safely composed. We support this by 1) specifying a translation from λ_{RS} to a language containing only strings and regular expressions, then, taking Python as such a language, 2) implementing the type system together with the translation as a library using *atlang*, an extensible static type system for Python (being developed by the authors).

1. INTRODUCTION

Command injection vulnerabilities are among the most common and severe security vulnerabilities in modern web applications [14]. They arise because web applications, at their boundaries, control external systems using commands represented as strings. For example, web browsers are controlled using HTML and Javascript sent from a server as a string, and database engines execute SQL queries also sent as strings. When these commands include data derived from user input, care must be taken to ensure that the user cannot subvert the intended command by carefully crafting the

data they send. For example, a SQL query constructed using string concatenation exposes a SQL injection vulnerability:

```
'SELECT * FROM users WHERE name=' + name + ''
```

If a malicious user enters the name `""; DROP TABLE users --`, the entire database could be erased.

To avoid this problem, the program must *sanitize* user input. For example, in this case, the developer (or, more often, a framework) might define a function *sanitize* that escapes double quotes and existing backslashes with a backslash, which SQL treats safely. Guaranteeing that user input has already been sanitized like this before it is used to construct a command is challenging. Note that this function is not idempotent, so it should only be called once.

We observe that many such sanitization techniques can be understood using *regular languages* [8]. For example, *name* must be a string in the language described by the regular expression $([^\backslash"]|(\backslash")|(\backslash\backslash))^*$ – a sequence of characters other than quotation marks and backslashes; these can only appear escaped. This concrete syntax for regular expression patterns can be understood to desugar, in a standard way, to the syntax for regular expressions shown in Figure 1, where $r \cdot r$ is sequencing and $r + r$ is disjunction. We will work with this “core” for simplicity.

In this paper, we present a static type system that tracks the regular language a string belongs to. For example, the output of *sanitize* will be a string in the regular language described by the regular expression above (we describe such a regular language, following convention, as $\mathcal{L}\{r\}$). By leveraging closure and decidability properties of regular languages, the language of a string is tracked through uses of a number of operations, including *replacement* of substrings matching a given regular expression. This makes it simple to implement sanitation procedures like the one just described in an essentially conventional manner. The result is a system where the fact that a string has been correctly sanitized is manifest in its type. Missing calls to sanitization functions are detected statically, and, importantly, so are *incorrectly implemented sanitization functions* (i.e. these functions need not be trusted). These guarantees require run-time checks only when going from less precise to more precise types (e.g. at the edges of the system, where user input has not yet been validated).

We will begin in Sec. 2 by specifying this type system minimally, as a conservative extension of the simply typed lambda calculus called λ_{RS} . This allows us to specify the guarantees that the type system provides precisely. We also formally specify a translation from this calculus to a typed calculus with only standard strings and regular expressions,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

intending it as a guide to language implementors interested in building this feature into their own languages. This also demonstrates that no additional space overhead is required.

Waiting for a language designer to build this feature in is unsatisfying in practice. Moreover, we also face a “chicken-and-egg problem”: justifying its inclusion into a commonly used language benefits from empirical demonstrations that it is useful, but this is difficult to do if developers have no way to use it in practice. As such, we take the position that a better path forward for the community is to work within a programming language where such type system fragments can be introduced modularly and orthogonally, as libraries.

In Sec. 3, we show how to implement the type system fragment we have specified using **atlang**, an extensible static type system implemented as a library inside Python. **atlang** leverages local type inference to control the semantics of literal forms, so regular string types can be introduced using string literals without any run-time overhead. Coercions that are known to be safe due to a sublanguage relationship are performed implicitly, also without run-time overhead. This results in a *usably secure* system: working with regular strings differs little from working with standard strings.

We conclude after discussing related work in Sec. 4.

2. REGULAR STRING TYPES, MINIMALLY

This section is organized as follows:

- Section 2.1 describes λ_{RS} . We also give proof outlines of type safety and correctness theorems and relevant propositions about regular languages.
- Section 2.2 describes a simple target language, λ_P , with a minimal regular expression library. In Section 3, we will take Python to be such a language.
- Section 2.3 describes the translation from λ_{RS} to λ_P and ensures the correctness result for 2.1 is preserved under translation.

2.1 The Language of Regular Strings

In this section, we define a minimal typed lambda calculus with regular string types called λ_{RS} . The syntax of λ_{RS} is specified in Figure 2, its static semantics in Figure 3 and its evaluation semantics in Figure 4.¹

There are two type constructors in λ_{RS} , \rightarrow and **stringin**. Arrow types classify functions, which are introduced via lambda abstraction, $\lambda x.e$, and can be applied, written $e(e)$, in the usual way [7]. Regular string types are of the form **stringin** $[r]$, where r is a regular expression. Values of such regular string types take the form **rstr** $[s]$, where s is a string (i.e. $s \in \Sigma^*$, defined in the usual way). The rule S-T-STRINGIN-I statically guarantees that $s \in \mathcal{L}\{r\}$.

λ_{RS} provides several familiar operations on strings. The type system handles these operations by relating operations over strings to corresponding operations over the regular languages they belong to. Since these operations over regular languages are known to be closed and decidable, we can use the corresponding operations over regular strings as a basis for static analysis of sanitation protocols.

¹For convenience, a single sheet containing Figures 2-8 is available at <http://nfulton.org/strings/printout.pdf>.

$$r ::= \epsilon \mid \cdot \mid a \mid r \cdot r \mid r + r \mid r^* \quad a \in \Sigma$$

Figure 1: Regular expressions over the alphabet Σ .

$$\begin{array}{ll} \sigma ::= \sigma \rightarrow \sigma & \text{source types} \\ \quad \mid \text{stringin}[r] & \\ v ::= \lambda x.e & \text{source values} \\ \quad \mid \text{rstr}[s] & s \in \Sigma^* \langle \$e\$ \rangle ::= x \quad \text{source terms} \\ \quad \mid v & \\ \quad \mid e(e) & \\ \quad \mid \text{rconcat}(e; e) & \\ \quad \mid \text{rstrcase}(e; e; x.y.e) & \\ \quad \mid \text{rreplace}[r](e; e) & \\ \quad \mid \text{rcoerce}[r](e) & \\ \quad \mid \text{rcheck}[r](e; x.e; e) & \end{array}$$

Figure 2: Syntax of λ_{RS} .

2.1.1 Concatenation and String Decomposition

The S-T-CONCAT rule is the simplest example of our approach. The rule is sound because the result of concatenating two strings will always be in the sequential composition of two regular expressions matching the respective strings. The rule therefore relates string concatenation to sequential composition of regular expressions.

Whereas concatenation allows the construction of large strings from smaller strings, S-T-CASE allows the decomposition, or elimination, of large strings into smaller strings. Intuitively, this rule “peels off” the first character, and then performs a specified operation on the first character and the remaining string. In terms of strings, this is essentially the ubiquitous substring operation.

The definition of S-T-CASE is subtle because even simple string operations sometimes require complicated operations on regular expressions. Peeling off the first character of a string is deterministic – the first character is always one, concrete value. However, a regular expression describing the first character of any matching string might recognize multiple concrete values. For instance, the first character of “CMU” is always “C”; however, if the expression associated with this string is CMU+KIT, then the expression describing the first character is C+K. Therefore, an operation on languages corresponding to case analysis on strings must consider one-character prefixes (a head) and their corresponding suffixes (a tail).

Fortunately, regular expression derivatives [2] conveniently capture this intuition. The regular expression recognizing any one-character prefix of the strings in $\mathcal{L}\{r\}$ is easily defined. The one-character prefix definition combined with derivatives gives a concise definition of the remaining string.

Definition 1 (Definition of $\text{lhead}(r)$). The function $\text{lhead}(r)$ is defined in terms of the structure of r :

- $\text{lhead}(ar') = a$ where $a \in \Sigma$
- $\text{lhead}(r_1 + r_2) = \text{lhead}(r_1) + \text{lhead}(r_2)$
- $\text{lhead}(r^*) = \text{lhead}(r)$
- $\text{lhead}(\cdot) = a_1 + a_2 + \dots + a_n$ for all $a_i \in \Sigma$ where $|\Sigma| = n$
- $\text{lhead}(\epsilon) = \epsilon$.

| | |
|--|--|
| $\boxed{\Psi \vdash e : \sigma}$ | $\Psi ::= \emptyset \mid \Psi, x : \sigma$ |
| $\frac{\text{S-T-VAR}}{x : \sigma \in \Psi} \quad \frac{\text{S-T-ABS}}{\Psi, x : \sigma_1 \vdash e : \sigma_2}$ $\Psi \vdash x : \sigma \quad \Psi \vdash \lambda x. e : \sigma_1 \rightarrow \sigma_2$ | |
| $\frac{\text{S-T-APP}}{\Psi \vdash e_1 : \sigma_2 \rightarrow \sigma \quad \Psi \vdash e_2 : \sigma_2} \quad \frac{\text{S-T-STRINGIN-I}}{s \in \mathcal{L}\{r\}}$ $\Psi \vdash e_1(e_2) : \sigma \quad \Psi \vdash \text{rstr}[s] : \text{stringin}[r]$ | |
| $\frac{\text{S-T-CONCAT}}{\Psi \vdash e_1 : \text{stringin}[r_1] \quad \Psi \vdash e_2 : \text{stringin}[r_2]}$ $\Psi \vdash \text{rconcat}(e_1; e_2) : \text{stringin}[r_1 \cdot r_2]$ | |
| $\frac{\text{S-T-CASE}}{\Psi \vdash e_1 : \text{stringin}[r] \quad \Psi \vdash e_2 : \sigma}$ $\Psi, x : \text{stringin}[\text{lhead}(r)], y : \text{stringin}[\text{ltail}(r)] \vdash e_3 : \sigma$ $\Psi \vdash \text{rstrcase}(e_1; e_2; x.y.e_3) : \sigma$ | |
| $\frac{\text{S-T-REPLACE}}{\Psi \vdash e_1 : \text{stringin}[r_1] \quad \Psi \vdash e_2 : \text{stringin}[r_2]}$ $\text{lreplace}(r; r_1; r_2) = r'$ $\Psi \vdash \text{rreplace}[r](e_1; e_2) : \text{stringin}[r']$ | |
| $\frac{\text{S-T-SAFECOERCE}}{\Psi \vdash e : \text{stringin}[r'] \quad \mathcal{L}\{r'\} \subseteq \mathcal{L}\{r\}}$ $\Psi \vdash \text{rcoerce}[r](e) : \text{stringin}[r]$ | |
| $\frac{\text{S-T-CHECK}}{\Psi \vdash e_0 : \text{stringin}[r_0] \quad \Psi, x : \text{stringin}[r] \vdash e_1 : \sigma \quad \Psi \vdash e_2 : \sigma}$ $\Psi \vdash \text{rcheck}[r](e_0; x.e_1; e_2) : \sigma$ | |

Figure 3: Typing rules for λ_{RS} . The typing context Ψ is standard.

Given this definition of $\text{lhead}(r)$, regular expression derivatives provide a useful tool for defining $\text{ltail}(r)$.

Definition 2 (Definition of $\text{ltail}(r)$). The function $\text{ltail}(r)$ is defined in terms of $\text{lhead}(r)$. Note that $\text{lhead}(r) = a_1 + a_2 + \dots + a_i$. We define $\text{ltail}(r) = \delta_{a_1}(r) + \delta_{a_2}(r) + \dots + \delta_{a_i}(r)$ where $\delta_s(r)$ is the Brzozowski derivative of r with respect to s [2].

The S-T-CONCAT rule, which is defined in terms of these operations, relates the result of “peeling off” the first character of a string to regular expression derivatives.

2.1.2 Coercion

The λ_{RS} language supports two forms of coercion. Safe coercion only allows coercion between strings types which are guaranteed to be safe. This form of coercion is useful because the replacement operation (introduced below in Section 2.1.3) is often more conservative than absolutely necessary. Conversely, unsafe coercion – which we refer to as a checked coercion – allows for potential unsafe type casts. Our semantics for check ensures that only safe values are used. Whenever a value is determined to be unsafe at runtime, a default value e_2 is used instead.

Summarily, S-T-SAFECOERCE allows only safe coercions between string types by exploiting the decidability of language inclusion, while S-T-CHECK allows casts between strings that cannot be guaranteed at compile time, but inserts a runtime check a runtime check.

2.1.3 Replacement

| | |
|---|--|
| $\boxed{e \Downarrow e}$ | |
| $\frac{\text{S-E-ABS}}{\lambda x. e \Downarrow \lambda x. e}$ | $\frac{\text{S-E-APP}}{e_1 \Downarrow \lambda x. e_3 \quad e_2 \Downarrow v_2 \quad [v_2/x]e_3 \Downarrow e'}$ $e_1(e_2) \Downarrow e'$ |
| $\frac{\text{S-E-RSTR}}{\text{rstr}[s] \Downarrow \text{rstr}[s]}$ | $\frac{\text{S-E-CONCAT}}{e_1 \Downarrow \text{rstr}[s_1] \quad e_2 \Downarrow \text{rstr}[s_2]}$ $\text{rconcat}(e_1; e_2) \Downarrow \text{rstr}[s_1 s_2]$ |
| | $\frac{\text{S-E-CASE-}\epsilon}{e_1 \Downarrow \text{rstr}[\epsilon] \quad e_2 \Downarrow v_2}$ $\text{rstrcase}(e_1; e_2; e_3) \Downarrow v_2$ |
| | $\frac{\text{S-E-CASE-CONCAT}}{e_1 \Downarrow \text{rstr}[ps] \quad e_3 \Downarrow x.y.e_4 \quad [p/x][s/y]e_4 \Downarrow e'}$ $\text{rstrcase}(e_1; e_2; e_3) \Downarrow e'$ |
| $\frac{\text{S-E-REPLACE}}{e_1 \Downarrow \text{rstr}[e_1] \quad e_2 \Downarrow \text{rstr}[e_2] \quad \text{subst}(r; s_1; s_2) = s}$ $\text{rreplace}[r](e_1; e_2) \Downarrow \text{rstr}[s]$ | |
| | $\frac{\text{S-E-SAFECOERCE}}{e \Downarrow \text{rstr}[s]}$ $\text{rcoerce}[r](e) \Downarrow \text{rstr}[s]$ |
| | $\frac{\text{S-E-CHECK-OK}}{e \Downarrow \text{rstr}[s] \quad s \in \mathcal{L}\{r\} \quad [\text{rstr}[s]/x]e_1 \Downarrow e'}$ $\text{rcheck}[r](e; x.e_1; e_2) \Downarrow e'$ |
| | $\frac{\text{S-E-CHECK-NOTOK}}{e \Downarrow \text{rstr}[s] \quad s \notin \mathcal{L}\{r\}}$ $\text{rcheck}[r](e; x.e_1; e_2) \Downarrow e_2$ |

Figure 4: Big step semantics for λ_{RS}

The premier operation for manipulating strings in λ_{RS} is string substitution. On strings, the string substitution operator is familiar, and roughly analogous to `str_replace` in PHP and `String.replace` in Java. The definition of S-T-CASE requires a definition of replacement for regular expressions (or languages) corresponding to string replacement. In this section, we define this operation.

Both string and language replacement are defined extralinguistically. The system λ_{RS} is defined in terms of these functions. The function $\text{subst}(r; s_1; s_2)$ replaces every substring of s_1 matching r with s_2 . In an analogous manner, the function $\text{lreplace}(r; r_1; r_2)$ replaces every sublanguages of r_1 matching r with r_2 . If an intuition for this operation is not clear, it may be helpful to think in terms replacing sub-automata.

Throughout this section, we fix an alphabet Σ over which strings s and regular expressions r are defined. We use $\mathcal{L}\{r\}$ to refer to the language recognized by the regular expression r .

Lemma 3 (Properties of Regular Languages and Expressions.). *The following are properties of regular expressions which are necessary for our proofs: If $s_1 \in \mathcal{L}\{r_1\}$ and $s_2 \in \mathcal{L}\{r_2\}$ then $s_1 s_2 \in \mathcal{L}\{r_1 r_2\}$. For all strings s and regular expressions r , either $s \in \mathcal{L}\{r\}$ or $s \notin \mathcal{L}\{r\}$. Regular languages are closed under reversal.*

If any of these properties are unfamiliar, the reader may refer to a standard text on the subject [8].

Definition 4 (subst). The relation $\text{subst}(r; s_1; s_2) = s$ produces a string s in which all substrings of s_1 matching r are replaced with s_2 .

A proper definition of lreplace would give an rewrite system with correctness and termination proofs, which is beyond the scope of this paper on security-motivated type system extensions. Instead, we provide an abstract definition of the relation and state necessary propositions.

Definition 5 (lreplace). The relation $\text{lreplace}(r; r_1; r_2) = r'$ relates r, r_1 , and r_2 to a language r' containing all strings of r_1 except that any substring $s_{pre}ss_{post} \in \mathcal{L}\{r_1\}$ where $s \in \mathcal{L}\{r\}$ is replaced by the set of strings $s_{pre}s_2s_{post}$ for all $s_2 \in \mathcal{L}\{r_2\}$ (the prefix and postfix positions may be empty).

Proposition 6 (Closure.). If $\mathcal{L}\{r\}, \mathcal{L}\{r_1\}$ and $\mathcal{L}\{r_2\}$ are regular expressions, then $\mathcal{L}\{\text{lreplace}(r; r_1; r_2)\}$ is also a regular language.

Proof Sketch. Algorithms for the inclusion problem may be adopted to identify any sublanguage $x \subseteq r$ of r_1 . The language x_{pre} can be computed by taking total derivatives until the remaining language equals x ; x_{post} can be computed in a similar way after reversal. Then r' is $x_{pre}r_2x_{post}$. \square

Proposition 7 (Substitution Correspondence.). If $s_1 \in \mathcal{L}\{r_1\}$ and $s_2 \in \mathcal{L}\{r_2\}$ then $\text{subst}(r; s_1; s_2) \in \mathcal{L}\{\text{lreplace}(r; s_1; s_2)\}$.

Proof Sketch. The proposition follows from the definitions of subst and lreplace ; note that language substitutions over-approximate string substitutions. \square

2.1.4 Safety

In this section, we establish type soundness for λ_{RS} . The theorem relies upon the lemmas and propositions established above. We also rely on additional lemmas which establish the preservation of well-formedness of regular expressions.

Lemma 8. If $\Psi \vdash e : \text{stringin}[r]$ then r is a well-formed regular expression.

Proof Sketch. The only non-trivial case is S-T-Replace, which follows from lemma 6. \square

Lemma 9. If $\Theta \vdash \iota : \text{regex}$ then $\iota \Downarrow \text{rx}[r]$ such that r is a well-formed regular expression.

Safety for the string fragment of λ_{RS} requires validating that the type system's definition is justified by our theorems about regular languages. We avoid the most significant issues inherent to big-step semantics by avoiding non-termination. The simply typed lambda calculus terminates, and our conservative, compositional extension clearly terminates modulo termination of subst and lreplace (but even if these were non-terminating, this still does not pose a problem for the big-step semantics itself). However, a more careful treatment or treatment for a language with non-termination might proceed by either coinduction [1, 11] or step-indexing [4].

Theorem 10 (Canonical Forms for String Fragment of λ_{RS} .). If $\Psi \vdash e : \text{stringin}[r]$ then $e \Downarrow \text{rstr}[s]$

Theorem 11 (Type Safety.). If $\Psi \vdash e : \sigma$ then $e \Downarrow e'$ and $\Psi \vdash e' : \text{stringin}[r]$.

Proof Sketch. By induction on the typing relation. The S-T-Concat case requires Lemma 3 and the S-T-Replace case appeals to Lemma 7. \square

2.1.5 The Security Theorem

The chief benefit of λ_{RS} is its safety theorem, which states that any value of a regular expression type is recognized by the regular expression corresponding to its type. Our main technical result, stated later in this section, establishes that this property is preserved under translation into λ_P .

Theorem 12 (Correctness of Input Sanitation for λ_{RS} .). If $\Psi \vdash e : \text{stringin}[r]$ and $e \Downarrow \text{rstr}[s]$ then $s \in \mathcal{L}\{r\}$.

Proof Sketch. The theorem follows directly from type safety, canonical forms for λ_{RS} , and inversion of the typing relation for λ_{RS} . \square

2.2 A Target Language with a Regular Expression Library

The system λ_P is a straight-forward extension of a simply typed lambda calculus with a `string` type and a regular expression type, as well as some operations – such as concatenation and replacement – found in many the standard libraries of many programming languages. The operations of λ_P correspond precisely to the operations of λ_{RS} in a way made precise by the translation rules described in the next section. Unlike λ_{RS} , λ_P does not statically track the effects of string operations. The language λ_P is so-called because it is reminiscent of popular web programming languages, such as Python or PHP.

The grammar of λ_P is defined in Figure 5. The typing rules P-T- are defined in Figure 6 and a big-step semantics is defined by the rules P-E- in Figure 7.

| | |
|---|---------------|
| $\theta ::= \theta \rightarrow \theta$ | target types |
| <code>string</code> | |
| <code>regex</code> | |
| $\dot{v} ::= \lambda x. \iota$ | target values |
| <code>str</code> [s] | |
| $\iota ::= x$ | target terms |
| v | |
| ι | |
| <code>rx</code> [r] | |
| <code>concat</code> ($\iota; \iota$) | |
| <code>pstrcase</code> ($\iota; \iota; x.y.\iota$) | |
| <code>preplace</code> ($\iota; \iota; \iota$) | |
| <code>check</code> (<code>rx</code> [r]; $\iota; \iota; \iota$) | |

Figure 5: Syntax for the target language, λ_P , containing strings and statically constructed regular expressions.

2.2.1 Safety

Type safety for λ_P is straight-forward, but is necessary in order to establish the correctness of our translation. Again, we elide a more careful treatment by treating a special case where our language terminates but note the multiple approaches to proving soundness for non-terminating languages with natural semantics [1, 4, 11].

$$\boxed{\Theta \vdash \iota : \tau} \quad \Theta ::= \emptyset \mid \Theta, x : \tau$$

$$\begin{array}{c}
\text{P-T-VAR} \\
\frac{x : \tau \in \Theta}{\Theta \vdash x : \tau} \\
\\
\text{P-T-ABS} \\
\frac{\Theta, x : \tau_1 \vdash \iota_2 : \tau_2}{\Theta \vdash \lambda x. \iota_2 : \tau_1 \rightarrow \tau_2} \\
\\
\text{P-T-APP} \\
\frac{\Theta \vdash \iota_1 : \tau_2 \rightarrow \tau \quad \Theta \vdash \iota_2 : \tau_2}{\Theta \vdash \iota_1(\iota_2) : \tau} \\
\\
\text{P-T-STRING} \\
\frac{}{\Theta \vdash \text{str}[s] : \text{string}} \\
\\
\text{P-T-REGEX} \\
\frac{}{\Theta \vdash \text{rx}[r] : \text{regex}} \\
\\
\text{P-T-CONCAT} \\
\frac{\Theta \vdash \iota_1 : \text{string} \quad \Theta \vdash \iota_2 : \text{string}}{\Theta \vdash \text{concat}(\iota_1; \iota_2) : \text{string}} \\
\\
\text{P-T-CASE} \\
\frac{\Theta \vdash \iota_1 : \text{string} \quad \Theta \vdash \iota_2 : \tau \quad \Theta, x : \text{string}, y : \text{string} \vdash \iota_3 : \tau}{\Theta \vdash \text{pstrcase}(\iota_1; \iota_2; \iota_3) : \tau} \\
\\
\text{P-T-REPLACE} \\
\frac{\Theta \vdash \iota_1 : \text{regex} \quad \Theta \vdash \iota_2 : \text{string} \quad \Theta \vdash \iota_3 : \text{string}}{\Theta \vdash \text{preplace}(\iota_1; \iota_2; \iota_3) : \text{string}} \\
\\
\text{P-T-CHECK} \\
\frac{\Theta \vdash \iota_r : \text{regex} \quad \Theta \vdash \iota_r : \text{regex} \quad \Theta, x : \text{string} \vdash \iota_2 : \sigma \quad \Theta \vdash \iota_3 : \sigma}{\Theta \vdash \text{check}(\iota_r; \iota_1; x. \iota_2; \iota_3) : \sigma}
\end{array}$$

Figure 6: Typing rules for λ_P . The typing context Θ is standard.

Theorem 13. Let ι be a term in the target language. If $\Theta \vdash \iota : \tau$ then $\iota \Downarrow \iota'$ and $\Theta \vdash \iota' : \tau$.

2.3 Translation from λ_{RS} to λ_P

The translation from λ_{RS} to λ_P is defined in figure 5. The coercion cases are most interesting. If the safety of coercion is manifest in the types of the expressions, then no runtime check is inserted. If the safety of coercion is not manifest in the types, then a check is inserted.

In practice, the type of a replacement rarely matches a specification. Therefore, it is convenient in an implementation to always insert the appropriate safe coercion. This way, violations of the sanitation protocol manifest as type errors, but the definition of an adequate replacement is not arduous. Alternatively, this policy may be codified in the type system itself using subtyping [5].

The translation is defined by the rules TR- in Figure 8. This section ultimately establishes that the security theorem for λ_{RS} is preserved under compilation.

Theorem 14 (Translation Correctness). If $\Theta \vdash e : \text{stringin}[r]$ then there exists an ι such that $\llbracket e \rrbracket = \iota$, $\iota \Downarrow \text{str}[s]$, and $e \Downarrow \text{rstr}[s]$.

Proof Sketch. The proof proceeds by induction on the typing relation for e and an appropriate choice of ι ; in each case, the choice is the syntactic form in λ_P corresponding to the form under consideration (e.g. choose `preplace` when considering `sreplace`). After the correct choice, the proof proceeds by our type safety theorems and an appeal to the induction hypothesis. \square

2.3.1 Preservation of Security

Finally, our main result establishes that correctness of λ_{RS} is preserved under the translation into λ_P .

$$\boxed{\iota \Downarrow \iota}$$

$$\begin{array}{c}
\text{P-E-ABS} \\
\frac{}{\lambda x. e \Downarrow \lambda x. e} \\
\\
\text{P-E-APP} \\
\frac{\iota_1 \Downarrow \lambda x. \iota_3 \quad \iota_2 \Downarrow v \quad [v/x]\iota_3 \Downarrow \iota'}{\iota_1(\iota_2) \Downarrow \iota'} \\
\\
\text{P-E-STR} \\
\frac{}{\text{str}[s] \Downarrow \text{str}[s]} \\
\\
\text{P-E-RX} \\
\frac{}{\text{rx}[r] \Downarrow \text{rx}[r]} \\
\\
\text{P-E-CONCAT} \\
\frac{\iota_1 \Downarrow \text{str}[s_1] \quad \iota_2 \Downarrow \text{str}[s_2]}{\text{concat}(\iota_1; \iota_2) \Downarrow \text{str}[s_1 s_2]} \\
\\
\text{P-E-CASE-}\epsilon \\
\frac{\iota_1 \Downarrow \text{str}[\] \quad \iota_2 \Downarrow \iota'_2}{\text{pstrcase}(\iota_1; \iota_2; \iota_3) \Downarrow \iota'_1} \\
\\
\text{P-E-CASE-CONCAT} \\
\frac{\iota_1 \Downarrow \text{str}[ps] \quad x.y.\iota_3 \Downarrow \iota_4 \quad [p/x][s/y]\iota_4 \Downarrow \iota'}{\text{pstrcase}(\iota_1; \iota_2; \iota_3) \Downarrow \iota'} \\
\\
\text{P-E-REPLACE} \\
\frac{\iota_1 \Downarrow \text{rx}[r] \quad \iota_2 \Downarrow \text{str}[s_2] \quad \iota_3 \Downarrow \text{str}[s_3] \quad \text{subst}(r; s_2; s_3) = s}{\text{preplace}(\iota_1; \iota_2; \iota_3) \Downarrow \text{str}[s]} \\
\\
\text{P-E-CHECK-OK} \\
\frac{\iota \Downarrow \text{str}[s] \quad s \in \mathcal{L}\{r\} \quad [\text{str}[s]/x]\iota_1 \Downarrow \iota'}{\text{check}(\text{rx}[r]; \iota; x.\iota_1; \iota_2) \Downarrow \iota'} \\
\\
\text{P-E-CHECK-NOTOK} \\
\frac{\iota \Downarrow \text{str}[s] \quad s \notin \mathcal{L}\{r\}}{\text{check}(\text{rx}[r]; \iota; x.\iota_1; \iota_2) \Downarrow \iota_2}
\end{array}$$

Figure 7: Big step semantics for λ_P

Theorem 15 (Correctness of Input Sanitation for Translated Terms.). If $\llbracket e \rrbracket = \iota$ and $\Psi \vdash e : \text{stringin}[r]$ then $\iota \Downarrow \text{str}[s]$ for $s \in \mathcal{L}\{r\}$.

Proof Sketch. By theorem 14, $\iota \Downarrow \text{str}[s]$ implies that $e \Downarrow \text{rstr}[s]$. By theorem 12, the above property together with the assumption that e is well-typed implies that $s \in \mathcal{L}\{r\}$. \square

3. IMPLEMENTATION IN ATLANG

In the previous section, we specified a type system and a translation semantics to a language containing only strings and regular expressions. In this section, we take Python to be such a target language. Python does not have a static type system, however, so to implement these semantics, we will leverage `atlang`, an extensible type system for Python (being developed by the authors). By using `atlang`, which builds on Python's quotations and reflection facilities, we can implement these semantics as a library, rather than as a new dialect of Python.

3.1 Example Usage

Figure 9 demonstrates the use of two type constructors, `fn` and `string_in`, both of which we have included in `atlib`, the standard library for `atlang`. The `fn` type constructor can be used to annotate functions that should be statically checked by `atlang`. The function `sanitize` on lines 3-7 specifies one argument, `s`, of type `string_in[r'.*'].fix sizes`.

The `sanitize` function takes an arbitrary string and returns a string without double quotes or left and right brackets. In this example, we use HTML escape sequences.

The main function receives user input and passes this input to a `sanitize` function, which replaces all double quotes

$$\boxed{\llbracket e \rrbracket = \iota}$$

$$\begin{array}{c}
\text{TR-VAR} \quad \frac{}{\llbracket x \rrbracket = x} \quad \text{TR-ABS} \quad \frac{\llbracket e \rrbracket = \iota}{\llbracket \lambda x. e \rrbracket = \lambda x. \iota} \quad \text{TR-APP} \quad \frac{\llbracket e_1 \rrbracket = \iota_1 \quad \llbracket e_2 \rrbracket = \iota_2}{\llbracket e_1(e_2) \rrbracket = \iota_1(\iota_2)} \\
\\
\text{TR-CASE} \quad \frac{\llbracket e_1 \rrbracket = \iota_1 \quad \llbracket e_2 \rrbracket = \iota_2 \quad \llbracket e_3 \rrbracket = \iota_3}{\llbracket \text{rstrcase}(e_1; e_2; e_3) \rrbracket = \text{pstrcase}(\iota_1; \iota_2; \iota_3)} \\
\\
\text{TR-STRING} \quad \frac{}{\llbracket \text{rstr}[s] \rrbracket = \text{str}[s]} \quad \text{TR-CONCAT} \quad \frac{\llbracket e_1 \rrbracket = \iota_1 \quad \llbracket e_2 \rrbracket = \iota_2}{\llbracket \text{rconcat}(e_1; e_2) \rrbracket = \text{concat}(\iota_1; \iota_2)} \\
\\
\text{TR-SUBST} \quad \frac{\llbracket e_1 \rrbracket = \iota_1 \quad \llbracket e_2 \rrbracket = \iota_2}{\llbracket \text{rreplace}[r](e_1; e_2) \rrbracket = \text{replace}(\text{rx}[r]; \iota_1; \iota_2)} \\
\\
\text{TR-SAFECOERCE} \quad \frac{\llbracket e \rrbracket = \iota}{\llbracket \text{rcoerce}[r'](e) \rrbracket = \iota} \\
\\
\text{TR-CHECK} \quad \frac{\llbracket e \rrbracket = \iota_1 \quad \llbracket e_1 \rrbracket = \iota_1 \quad \llbracket e_2 \rrbracket = \iota_2}{\llbracket \text{rcheck}[r](e; x.e_1; e_2) \rrbracket = \text{check}(\text{rx}[r]; \iota; x.\iota_1; \iota_2)}
\end{array}$$

Figure 8: Translation from source terms (e) to target terms (ι). The translation is type-directed.

```

1  from atlib import fn, string_in
2
3  @fn
4  def sanitize(s : string_in[r'.*']):
5      return (s.replace(r'\"', '"')
6              .replace(r'<', '&lt;')
7              .replace(r'>', '&gt;'))
8
9  @fn
10 def results_query(s : string_in[r'[""]*']):
11     return 'SELECT * FROM users WHERE name=' + s + ''
12
13 @fn
14 def results_div(s : string_in[r'^<>*']):
15     return '<div>Results for ' + s + '</div>'
16
17 @fn
18 def main(db):
19     input = sanitize(user_input())
20     results = db.execute(results_query(input))
21     return results_div(input) + format(results)

```

Figure 9: Regular string types in atlang, a library that enables static type checking for Python.

and brackets with HTML escape sequences.

The result of applying `sanitize` to input is appended in two functions which construct a safe query (avoiding command injection) and safe HTML output (avoiding Cross-Site Scripting (XSS) attacks). The arguments to the result and output construction functions constitute *specifications*. In the case of `results_query`, this specification ensures that user input is always interpreted as a string literal by the SQL server. In the case of `results_div`, this specification ensures that user input does not contain any HTML tags, which is a conservative but effective policy for preventing XSS attacks.

Note that `input` does not actually meet these specifications without additional machinery. The type of `input` is quite

large and does not actually equal the specified domains of the query or output construction methods. This mismatch is common – in fact, nearly universal. Therefore, our implementation includes a simple subtyping relation between regular expression types.

This subtyping relation is justified theoretically by the fact that language inclusion is decidable; see [5] for a formal definition of the subtyping relation. Additionally, our extension remains composable because subtyping is defined on a type-by-type basis; see [6] for a discussion of subtyping in Atlang (referred to there as Ace).

3.2 Implementation of the Regular Expression Type

We implemented a variation on the type system presented in this paper with two significant differences. first, we only support replacements where s_2 is the empty string. Therefore, our implementation respects the system presented in section 2 only modulo the definition of `lreplace`. Second, we use subtyping instead of `S*-SAFE-COERCE`, which decreases inessential verbosity of the language.

Atlang translates programs using type definitions, which may extend both the static and dynamic semantics of the language. New types are defined as Python classes; figure 10 contains the source code of our implementation.

The `string_in` type has an indexing regular expression `idx`. Our translation is defined by the `trans_` methods while the `syn_` methods define our type checker. Atlang defers type checking and translation to these methods whenever an expression of type `string_in` is encountered.

4. RELATED WORK AND ALTERNATIVE APPROACHES

The input sanitation problem is well-understood. There exist a large number of techniques and technologies, proposed by both practitioners and researchers, for preventing injection-style attacks. In this section, we explain how our approach to the input sanitation problem differs from each of these approaches. More important than these differences, however, is our more general assertion that language extensibility is a promising approach toward consideration of security goals in programming language design.

Unlike *frameworks and libraries* provided by languages such as Haskell and Ruby, our type system provides a *static* guarantee that input is always properly sanitized before use. Doing so requires reasoning about the operations on regular languages corresponding to standard operations on strings; we are unaware of any production system which contains this form of reasoning. Therefore, even where frameworks and libraries provide a viable interface or wrapper around input sanitation, our approach is complementary because it ensures the correctness of the framework or library itself. Furthermore, our approach is more general than database abstraction layers because our mechanism is applicable to all forms of command injection (e.g. shell injection or remote file inclusion).

A number of research languages provide static guarantees that a program is free of input sanitation vulnerabilities [3]. Unlike this work, our solution to the input sanitation problem has a very low barrier to adoption; for instance, our implementation conservatively extends Python – a popular language among web developers. We also believe our gen-


```

1 class string_in(atlang.Type):
2     def __init__(self, rx):
3         rx = rx_normalize(rx)
4         atlang.Type.__init__(idx=rx)
5
6     def ana_Str(self, ctx, node):
7         if not in_lang(node.s, self.idx):
8             raise atlang.TypeError("...", node)
9
10    def trans_Str(self, ctx, node):
11        return astx.copy(node)
12
13    def syn_BinOp_Add(self, ctx, node):
14        left_t = ctx.syn(node.left)
15        right_t = ctx.syn(node.right)
16        if isinstance(left_t, string_in):
17            left_rx = left_t.idx
18            if isinstance(right_t, string_in):
19                right_rx = right_t.idx
20                return string_in[lconcat(left_rx, right_rx)]
21            raise atlang.TypeError("...", node)
22
23    def trans_BinOp_Add(self, ctx, node):
24        return astx.copy(node)
25
26    def syn_Method_replace(self, ctx, node):
27        [rx, exp] = node.args
28        if not isinstance(rx, ast.Str):
29            raise atlang.TypeError("...", node)
30        rx = rx.s
31        exp_t = ctx.syn(exp)
32        if not isinstance(exp_t, string_in):
33            raise atlang.TypeError("...", node)
34        exp_rx = exp_t.idx
35        return string_in[lreplace(self.idx, rx, exp_rx)]
36
37    def trans_Method_replace(self, ctx, node):
38        return astx.quote(
39            """__import__(re); re.sub(%0, %1, %2)""",
40            astx.Str(s=node.args[0]),
41            astx.copy(node.func.value),
42            astx.copy(node.args[1]))
43
44    def syn_Method_check(self, ctx, node):
45        [rx] = node.args
46        if not isinstance(rx, ast.Str):
47            raise atlang.TypeError("...", node)
48        return string_in[rx.s]
49
50    def trans_Method_check(self, ctx, node):
51        return astx.quote(
52            """__import__(string_in_helper);
53            string_in_helper.coerce(%0, %1)""",
54            astx.Str(s=other_t.idx),
55            astx.copy(node))
56
57    def check_Coerce(self, ctx, node, other_t):
58        # coercions can only be defined between
59        # types with the same type constructor,
60        if rx_sublang(other_t.idx, self.idx):
61            return other_t
62        else: raise atlang.TypeError("...", node)

```

Figure 10: Implementation of the `string_in` type constructor in `atlang`.

eral approach is better-positioned for security, where continuously evolving threats might require frequent addition of new analyses; in these cases, the composability and generality of our approach is a substantial advantage.

The Wyvern programming language provides a general framework for composing language extensions [13][12]. Our work identifies one particular extension, and is therefore complementary to Wyvern and related work on extensible programming languages. We are also unaware of any extensible programming languages which emphasize applications to security concerns.

Incorporating regular expressions into the type system is

not novel. The XDuce system [10, 9] checks XML documents against schema using regular expressions. Similarly, XHaskell [15] focuses on XML documents. We differ from this and related work in at least three ways:

- Our system is defined within an extensible type system.
- We demonstrate that regular expression types are applicable to the web security domain, whereas previous work on regular expression types focused on XML schema.
- Although our static replacement operation is definable in some languages with regular expression types, we are the first to expose this operation and connect the semantics of regular language replacement with the semantics of string substitution via a type safety and compilation correctness argument.

In conclusion, our contribution is a type system, implemented within an extensible type system, for checking the correctness of input sanitation algorithms.

5. FUTURE WORK

We believe that this type system extension serves as a useful basis for web-oriented static analysis. The use sites of arbitrary strings are usually constrained to a few sections of the application. Assuming regular expressions are used for input sanitation, annotations at these use sites could be sufficient input to a sound and complete static analysis.

We also believe that extensible programming languages are a promising approach toward incorporating security and privacy into programming languages. The Software Industry moves quickly, and developing new languages for the particular security and privacy concerns of each new domain is likely a losing strategy. Conversely, extensible languages might provide developers with a way to incorporate domain-specific security and privacy analyses into their daily development activities. In future work, we hope to bolster this hypothesis by extending our approach in this paper to more realistic settings and by developing new type extensions addressing security and privacy concerns.

6. CONCLUSION

Composable analyses which complement existing approaches constitute a promising approach toward the integration of security concerns into programming languages. In this paper, we presented a system with both of these properties and defined a security-preserving transformation. Unlike other approaches, our solution complements existing, familiar solutions while providing a strong guarantee that traditional library and framework-based approaches are implemented and utilized correctly.

7. REFERENCES

- [1] D. Ancona. How to prove type soundness of java-like languages without forgoing big-step semantics. In *Proceedings of 16th Workshop on Formal Techniques for Java-like Programs, FTfJP'14*, pages 1:1–1:6, New York, NY, USA, 2014. ACM.
- [2] J. A. Brzozowski. Derivatives of regular expressions. *J. ACM*, 11(4):481–494, Oct. 1964.

- [3] A. Chlipala. Static checking of dynamically-varying security policies in database-backed applications. In *OSDI'10: Proceedings of the 9th USENIX Symposium on Operating Systems Design and Implementation*, Oct. 2010.
- [4] E. Ernst, K. Ostermann, and W. R. Cook. A virtual class calculus. In *Conference Record of the 33rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '06, pages 270–282, New York, NY, USA, 2006. ACM.
- [5] N. Fulton. Security through extensible type systems. In *Proceedings of the 3rd Annual Conference on Systems, Programming, and Applications: Software for Humanity*, SPLASH '12, pages 107–108, New York, NY, USA, 2012. ACM.
- [6] N. Fulton. A typed lambda calculus for input sanitation. Undergraduate thesis in mathematics, Carthage College, 2013.
- [7] R. Harper. *Practical Foundations for Programming Languages*. Cambridge University Press, 2012.
- [8] J. E. Hopcroft and J. D. Ullman. *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley, 1979.
- [9] H. Hosoya and B. C. Pierce. XDuce: A statically typed XML processing language. *ACM Transactions on Internet Technology*, 3(2):117–148, May 2003.
- [10] H. Hosoya, J. Vouillon, and B. C. Pierce. Regular Expression Types for XML. In *ICFP '00*, 2000.
- [11] X. Leroy. Coinductive big-step operational semantics. In P. Sestoft, editor, *Programming Languages and Systems*, volume 3924 of *Lecture Notes in Computer Science*, pages 54–68. Springer Berlin Heidelberg, 2006.
- [12] L. Nistor, D. Kurilova, S. Balzer, B. Chung, A. Potanin, and J. Aldrich. Wyvern: A simple, typed, and pure object-oriented language. In *Proceedings of the 5th Workshop on Mechanisms for Specialization, Generalization and Inheritance*, MASPEGHI '13, pages 9–16, New York, NY, USA, 2013. ACM.
- [13] C. Omar, D. Kurilova, L. Nistor, B. Chung, A. Potanin, and J. Aldrich. Safely composable type-specific languages. In R. Jones, editor, *ECOOP 2014 – Object-Oriented Programming*, volume 8586 of *Lecture Notes in Computer Science*, pages 105–130. Springer Berlin Heidelberg, 2014.
- [14] OWASP. Open web application security project top 10.
- [15] M. Sulzmann and K. Lu. Xhaskell – adding regular expression types to haskell. In O. Chitil, Z. Horváth, and V. Zsák, editors, *Implementation and Application of Functional Languages*, volume 5083 of *Lecture Notes in Computer Science*, pages 75–92. Springer Berlin Heidelberg, 2008.