

Collaborative Test-Driven Scientific Model Validation

Cyrus Omar, Jonathan Aldrich
Carnegie Mellon University
{comar,aldrich}@cs.cmu.edu

Richard C. Gerkin
Arizona State University
rgerkin@asu.edu

ABSTRACT

One of the pillars of the modern scientific method is *model validation*: comparing a scientific model’s predictions against empirical observations. Today, a scientist demonstrates the validity of a model by making an argument in a paper and submitting it for peer review, a process comparable to *code review* in software engineering. While human review helps to ensure that contributions meet high-level goals, software engineers typically supplement this with *unit testing* to get a more complete picture of the status of a software project, particularly for complex projects involving many developers.

We argue that a similar test-driven methodology would be valuable to scientific communities as they seek to validate increasingly complex models against growing collections of empirical data. The dynamics of scientific communities and software communities differ in several key ways, however. In this paper, we introduce *SciUnit*, a framework for test-driven scientific model validation. We describe how SciUnit, supported by new and existing collaborative infrastructure, can integrate into the modern scientific process.

1. INTRODUCTION

Scientific theories often take the form of a *quantitative model*: a formal structure that can generate predictions about observable quantities. Such a model is characterized by its *scope*: the set of observable quantities that the model can predict, and by its *validity*: the extent to which its predictions agree with experimental observations of these quantities.

Quantitative models are described and validated using a social process: the scientific publication system. For a model to be accepted by the scientific community, scientists must write a paper describing it and providing evidence that it predicts some quantity of interest more accurately than previous models, or that the tradeoff it makes between accuracy and complexity may be useful to the community [?]. During the *peer review* process, other members of the relevant community are tasked with ensuring that the validation criteria

used (e.g. a *p*-value) are strong enough and that all relevant data and competing models were fully considered, drawing on knowledge of statistical methods and the prior literature. Publishing a paper is one of the primary motivators for most scientists [1].

Quantitative scientific modeling and software development have much in common. Indeed, quantitative models are increasingly implemented as computer programs. In some cases, the program *is* the model (e.g. complex simulations). The peer review process is similar in many ways to the *code review* process used in many development teams. During code review, team members look for errors, enforce style and architectural guidelines and check that the code is *valid* (i.e. that achieves its intended goal [?]) before it can be committed to the primary source code repository [?].

Code review is reasonably effective [?, ?]. However, to be most effective reviewers must expend considerable effort [?], and code review is most effective for resolving issues related to software architecture [?]. For these reasons, most software development teams supplement code reviews with a number of more automated approaches to verification and validation. One of the most widely-adopted approaches is *unit testing*. In brief, unit tests are functions that check a single component of a program against a single well-defined correctness criterion. A suite of such tests can be seen as a partial specification of a program or component. Test suites complement code review procedures in several ways:

1. They help ensure that a code modification does not cause a *regression* in another part of a program.
2. They help measure development progress. Developers can use a test suite to determine which functionality is not yet adequately implemented.
3. They serve as a form of documentation. By inspecting the test suite itself, a developer can see which modes of usage the development team collectively considers most important.

However, in many areas, the number of publications being generated every year can overwhelm even the most conscientious scientists [?].

Unfortunately, there are few alternatives to a comprehensive literature review available when scientists need to answer questions like these:

- Which models are capable of predicting the quantities I am interested in?

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

- Which metrics should be used to evaluate the goodness-of-fit between these models and data?
- How well do these models perform, as judged to these metrics, given currently available data?
- What other quantities can and can't these models predict?
- What observations are not adequately explained by any available model?

Professional software developers face similar issues. They must understand the scope of each component of a complex software project and validate it by measuring how well each component achieves its specified input/output behavior. But software developers do not validate components by simply choosing a few interesting inputs and presenting the outputs to reviewers. Rather, they typically follow a *test-driven development* methodology by creating a suite of executable *unit tests* that serve to specify each component's scope and validate its implementation as it is being developed and modified [?]. Each test individually checks that a small portion of the program meets a single correctness criterion. For example, a unit test might verify that one function within the program correctly handles malformed inputs. Collectively, the test results serve as a summary of the validity of the project as it progresses through its development cycle. Developers can determine which features are unimplemented or buggy by examining the set of failed tests, and progress can be measured in terms of how many tests the program passes over time. This methodology is widely adopted in practice [?].

Test-driven methodologies have started to see success in neuroscience as well. Modeling competitions in neuroscience, for example, are typically organized around a collection of simple validation criteria, implemented as executable tests. These competitions continue to drive important advances and improve scientists' understanding of the relative merits of different models. For example, the quantitative single neuron modeling competition (QSNMC) [?] investigates the complexity-accuracy tradeoff among reduced models of excitable membranes; the "Hopfield" challenge [?] tested techniques for generating neuronal network form given its function; the Neural Prediction Challenge sought the best stimulus reconstructions, given neuronal activity (<http://neuralprediction.berkeley.edu>); the Diadem challenge is advancing the art of neurite reconstruction (<http://www.diademchallenge.org>); and examples from other subfields of biology abound (<http://www.the-dream-project.org>).

Each of these examples has leveraged *ad hoc* infrastructure to support test generation. While the specific criteria used to evaluate models varies widely between disciplines in neuroscience, the underlying test-driven methodology has many common features that could be implemented once. Recognizing this, we developed a discipline-agnostic framework for developing scientific validation test suites called *SciUnit* (<http://www.sciunit.org>). Here we describe *NeuronUnit*, which builds upon *SciUnit*, allowing neuroscientists to build *SciUnit* tests that validate neurophysiology models against electrophysiological data. We provide a concrete example pipeline, showing how models described using NeuroML and provided freely by the *Open Source Brain Project* (OSB, [?], <http://www.opensourcebrain.org>) can be tested in fully automated fashion using published, curated data available

```

1 class SpikeCountTest(sciunit.Test):
2     """Tests spike counts produced in response to
3         several current stimuli against observed means
4         and standard deviations.
5
6     goodness of fit metric: Computes p-values based on a
7         chi-squared test statistic, and pools them
8         using Fisher's method.
9     parameters:
10         inputs: list of numpy arrays containing input
11             currents (nA)
12         means, stds: lists of observed means and standard
13             deviations, one per input
14
15     """
16     def __init__(self, inputs, means, stds):
17         self.inputs, self.means, self.stds = inputs, means
18         , stds
19
20     required_capabilities = [SpikeCountFromCurrent]
21
22     def _judge(self, model):
23         inputs, means, stds = self.inputs, self.means,
24             self.stds
25         n = len(inputs)
26         counts = numpy.empty((n,))
27         for i in xrange(n):
28             counts[i] = model.spike_count_from_current(
29                 inputs[i])
30         chisquared = sum((counts-means)**2 / means) # An
31             array of chi-squared values.
32         p = scipy.stats.chi2.cdf(chisquared,n-1) # An
33             array of p-values.
34         pooled_p = sciunit.utils.fisher(p_array) # A
35             pooled p-value.
36         return sciunit.PValue(pooled_p, related_data={
37             "inputs": inputs, "counts": counts, "obs_means":
38                 means, "obs_stds": stds
39         })

```

Figure 2: An example single neuron spike count test class implemented using *SciUnit*. Because this implementation contains logic common to many different systems, *NeuronUnit* was developed to provide a simpler means to deliver it (see Sec. ??).

through the *NeuroElectro Project* (Neuroelectro, [?], <http://neuroelectro.org>), leveraging facilities from the *NeuroTools* library (<http://neuralensemble.org/NeuroTools>) to extract relevant features of model output. This is summarized in Figure 1, which shows the relationships between the layers described here.

2. VALIDATION TESTING WITH *SCIUNIT*

2.1 Example: The Quantitative Single Neuron Modeling Competition

We first illustrate the form of a generic example *SciUnit* test suite that could be used in neurophysiology. Suppose we have collected data from an experiment where current stimuli (measured in pA) are delivered to neurons in some brain region, while the somatic membrane potential of each stimulated cell (in mV) is recorded and stored. A model claiming to capture this cell type's membrane potential dynamics must be able to accurately predict a variety of features observed in these data.

One simple validation test would ask candidate models to predict the number of action potentials (a.k.a. spikes) generated in response to a stimulus (e.g. white noise), and compare these *spike count* predictions to the distribution observed in repeated experimental trials using the same stimulus. For data of this type, goodness-of-fit can be measured by first calculating a p-value from a chi-squared statistic




Figure 1: NeuronUnit overview. NeuronUnit is set of testing tools built upon the discipline-agnostic SciUnit framework. NeuronUnit can in principle test arbitrary neurophysiology models using arbitrary data but we provide here an example using models described in NeuroML as part of the *Open Source Brain Project* (OSB, [?], <http://www.opensourcebrain.org>), and single neuron electrophysiology data available as part of the NeuroElectro Project (Neuroelectro, [?], <http://neuroelectro.org>). Records of test results for various model/test combinations are accessible via SciDash, which indexes GitHub repositories of these records, models, and tests so they can be searched and filtered by the community.

for each prediction and then combining these p-values using Fisher’s method [?].

Alongside this *spike count test*, we might also specify a number of other tests capturing different features of the data to produce a more comprehensive suite. For data of this sort, the QSNMC defined 17 other validation criteria in addition to one based on the overall spike count, capturing features like spike latencies (SL), mean subthreshold voltage (SV), interspike intervals (ISI) and interspike minima (ISM) that can be extracted from the data [?]. They then defined a combined metric favoring models that broadly succeeded at meeting these criteria, to produce an overall ranking. Such combined criteria are simply validation tests that invoke other tests to produce a result.

2.2 Implementing a Validation Test in SciUnit

Fig. 2 shows how a scientist can implement spike count tests such as the one described above using *SciUnit*. A *SciUnit* validation test is an instance (i.e. an object) of a Python class implementing the `sciunit.Test` interface (cf. line 1). Here, we show a class `SpikeCountTest` taking three *parameters* in its constructor (constructors are named `__init__` in Python, lines 9-10). The meaning of each parameter along with a description of the goodness-of-fit metric used by the test is documented on lines 4-7. To create a *particular* spike count test, we instantiate this class with particular experimental observations. For example, given observations from hippocampal CA1 cells (not shown), we can instantiate a test as follows:

```
1 CA1_sc_test = SpikeCountTest(CA1_inputs, CA1_means,
                              CA1_std)
```

We emphasize the crucial distinction between the *class* `SpikeCountTest`, which defines a *parameterized family* of validation tests, and the particular *instance* `CA1_sc_test`, which is an individual validation test because the necessary parameters, derived from data, have been provided. As we will describe below, we expect communities to build repositories of such families capturing the criteria used in their subfields of neuroscience so that test generation for a particular system of interest will often require simply instantiating a previously-developed family with particular experimental parameters and data. For single-neuron test families like `SpikeCountTest`, we have developed such a library, called *NeuronUnit* (<http://github.com/scidash/neuronunit>) (Sec. ??).

Classes that implement the `sciunit.Test` interface must contain a `_judge` method that receives a candidate *model* as input and produces a *score* as output. To specify the interface between the test and the model (that is, to specify an appropriate scope), the test author provides a list of *capabilities* in the `required_capabilities` attribute, seen on line 12 of Fig. 2. Capabilities are simply collections of methods that a test will need to invoke in order to receive relevant data, and are analogous to *interfaces* in e.g. Java (<http://docs.oracle.com/javase/tutorial/java/concepts/interface.html>). In Python, capabilities are written as classes with unimplemented members. The capability required by the test in Fig. 2 is shown in Fig. 3. In *SciUnit*, classes defining capabilities are tagged as such by inheriting from `sciunit.Capability`. The test in Figure 2 uses this capability on line 19 to produce a spike count prediction for each input current.

The remainder of the `_judge` method implements the goodness-of-fit metric described above, returning an instance of `sciunit`

```

1 class SpikeCountFromCurrent(sciunit.Capability):
2     def spike_count_from_current(self, input):
3         """Takes a numpy array containing current stimulus
4           (in nA) and
5           produces an integer spike count. Can be called
6           multiple times."""
7         raise NotImplementedError("Model does not
8           implement capability.")

```

Figure 3: An example capability specifying a single required method (used by the test in Figure 2).

```

1 class TrainSpikeCountFromCurrent(sciunit.Capability):
2     def train_with_currents(self, currents, counts):
3         """Takes a list of numpy arrays containing current
4           stimulus (in nA) and
5           observed spike counts. Model parameters should be
6           adjusted based on this
7           training data."""
8         raise NotImplementedError("Model does not
9           implement capability.")

```

Figure 4: Another capability specifying a training protocol (not used by the test in Figure 2).

.PValue, a subclass of `sciunit.Score` that is included with *SciUnit*. In addition to the *p*-value itself, the returned score object also contains metadata, via the `related_data` parameter, for scientists who may wish to examine the result in more detail later. In this case we save the input currents, the model outputs and the observed means and standard deviations (line 24).

2.3 Models

Capabilities are *implemented* by models. In *SciUnit*, models are instances of Python classes that inherit from `sciunit.Model`. Like tests, the class itself represents a family of models, parameterized by the arguments of the constructor. A particular model is an instance of such a class.

Figure 5 shows how to write a simple family of models, `LinearModel`, that implement the capability in Fig. 3 as well as another capability shown in Fig. 4, which we will discuss below. Models in this family generate a spike count by applying a linear transformation to the mean of the provided input current. The family is parameterized by the scale factor and the offset of the transformation, both scalars. To create a *particular* linear model, a modeler can provide particular parameter values, just as with test families:

```

1 CA1_linear_model_heuristic = LinearModel(3.0, 1.0)

```

Here, the parameters to the model were picked by the modeler heuristically, or based on externally-available knowledge. An alternative test design would add a training phase where these parameters were fit to data using the capability shown in Fig. 4. This test could thus only be used for those models for which parameters can be adjusted without human involvement. Whether to build a training phase into the test protocol is a choice left to each test development community.

Fig. 2 does not include a training phase. If training data is externally available, models that nevertheless do implement a training capability (like `LinearModel`) can simply be trained explicitly by calling the capability method just like any other Python method:

```

1 CA1_linear_model_fit = LinearModel()
2 CA1_linear_model_fit.train_with_currents(
3     CA1_training_in, CA1_training_out)

```

```

1 class LinearModel(sciunit.Model, SpikeCountFromCurrent
2     , TrainSpikeCountFromCurrent):
3     def __init__(self, scale=None, offset=None):
4         self.scale, self.offset = scale, offset
5
6     def spike_count_from_current(self, input):
7         return int(self.scale*numpy.mean(input) + self.
8             offset)
9
10    def train_with_currents(self, currents, counts):
11        means = [numpy.mean(c) for c in currents]
12        [self.offset, self.scale] = numpy.polyfit(means,
13            counts, deg=1)

```

Figure 5: A model that returns a spike count by applying a linear transformation to the mean input current. The parameters can be provided manually or learned from data provided by a test or user (see text).

2.4 Executing Tests

A test is executed against a model using the `judge` method:

```

1 score = CA1_sc_test.judge(CA1_linear_model_heuristic)

```

This method proceeds by first checking that the provided model implements all required capabilities. It then calls the test’s `_judge` method to produce a score. A reference to the test and model are added to the score for convenience (accessible via the test and model attributes, respectively), before it is returned.

2.5 Test Suites and Score Matrices

A collection of tests intended to be run on the same model can be put together to form a test suite. The following is a test suite that could be used for a simplified version of the QSNMC, as described above:

```

1 CA1_suite = sciunit.TestSuite([CA1_sc_test,
2     CA1_sl_test, CA1_sv_test, CA1_isi_test,
3     CA1_ism_test])

```

Like a single test, a test suite is capable of judging one or more models. The result is a score matrix much like the one diagrammed in Fig. 1.

```

1 CA1_matrix = CA1_suite.judge([
2     CA1_linear_model_heuristic, CA1_linear_model_fit
3 ])

```

A simple summary of the scores in a score matrix can be printed to the console or visualized by other tools, such as the web application *SciDash* described in Sec. ??.

3. REFERENCES

- [1] J. Howison and J. Herbsleb, “Scientific software production: incentives and collaboration,” in *Proceedings of the ACM 2011 conference on Computer supported cooperative work*. ACM, 2011, pp. 513–522.
- [2] J. Hannay, C. MacLeod, J. Singer, H. Langtangen, D. Pfahl, and G. Wilson, “How do scientists develop and use scientific software?” in *Proceedings of the 2009 ICSE Workshop on Software Engineering for Computational Science and Engineering*. IEEE Computer Society, 2009, pp. 1–8.